

ECSE 446/546

# REALISTIC & ADVANCED IMAGE SYNTHESIS



Prof. Derek Nowrouzezahrai  
[derek@cim.mcgill.ca](mailto:derek@cim.mcgill.ca)

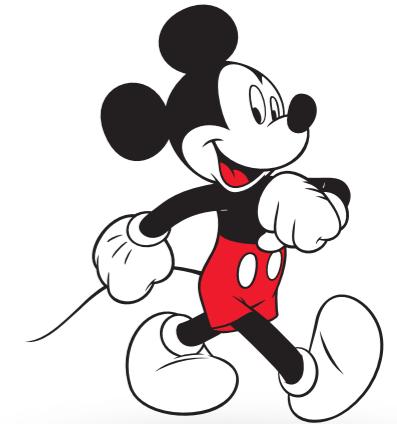
# About me



Derek Nowrouzezahrai  
**Associate Professor**

derek@cim.mcgill.ca

- “Those who can’t do, teach.”
- Previously worked films, games, parks, and products at:
  - Disney Research, Electronic Arts, Microsoft & Microsoft Research
- McGill’s Graphics & Imaging Lab
- NSERC/Ubisoft Industrial Research Chair



# Course Details

## Lecture Time and Location

- Wednesdays & Fridays
  - 2:35pm – 3:55pm in Trottier 1100

## Tutorials:

- Mondays, Wednesdays, Fridays
  - 9:35pm – 11:25pm in Trottier 4160

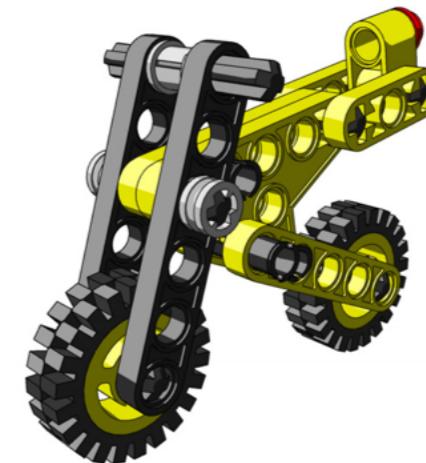
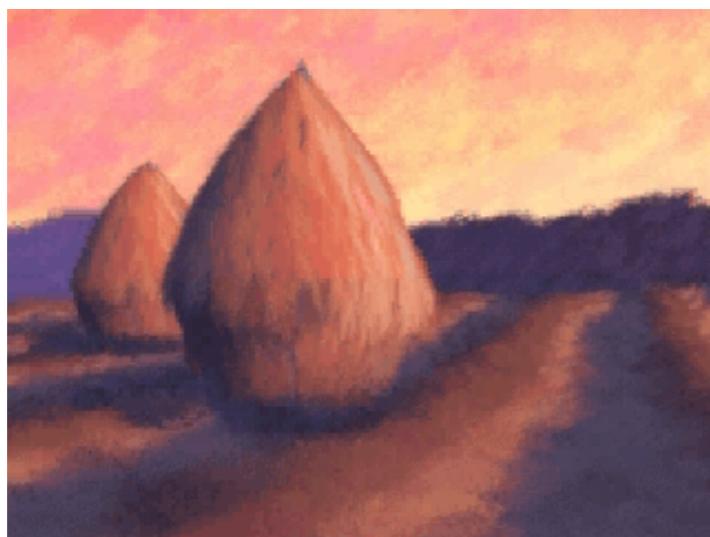
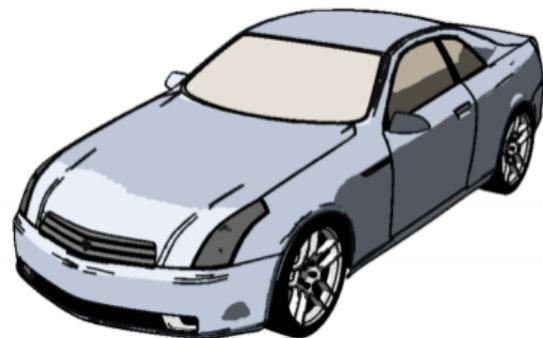
## Website

- <http://www.cim.mcgill.ca/~derek/ecse446.html>
- Will also use **myCourses** for discussion, slides and assignment hand-in

# Rendering

The process of converting a description of a 3D scene into an image

- **Non-Photorealistic Rendering (NPR)**



# Rendering

The process of converting a description of a 3D scene into an image

- **Photorealistic Rendering:** create an image of a 3D scene that is indistinguishable from a photo
- **Physically-based Rendering:** simulate the physical behaviour of light as closely as possible to *predict* what would be observed



# Curriculum

## Graphics & Imaging

**ECSE532: Computer Graphics**

**COMP557: Computer Graphics**

**ECSE446/546: Image Synthesis**

**COMP559: Computer Animation**

**ECSE544: Computational Photography**

**COMP521: Modern Computer Games**

## Vision/Imaging

**ECSE529: Computer & Biological Vision**

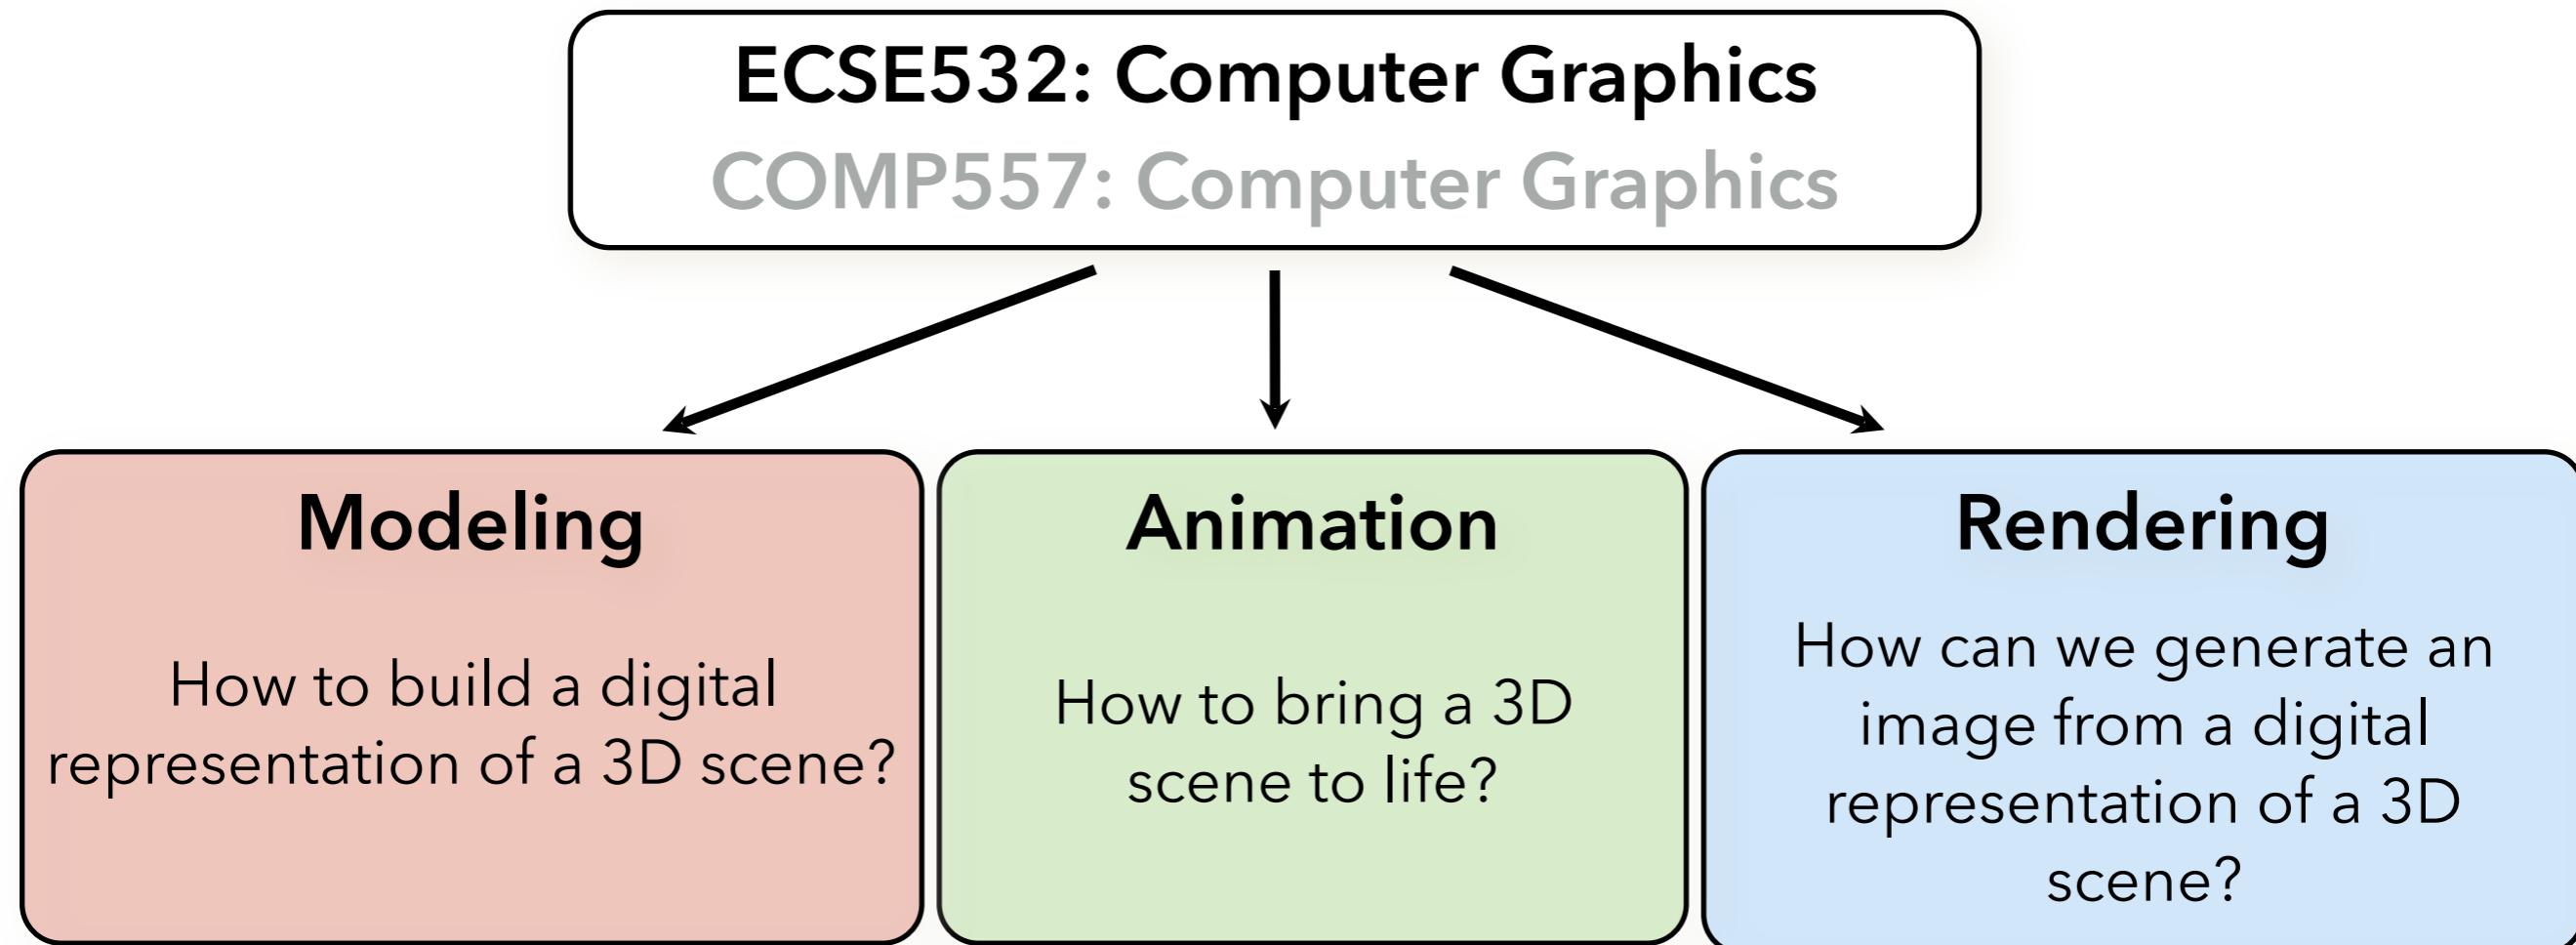
**ECSE626: Statistical Computer Vision**

**COMP558: Computer Vision**

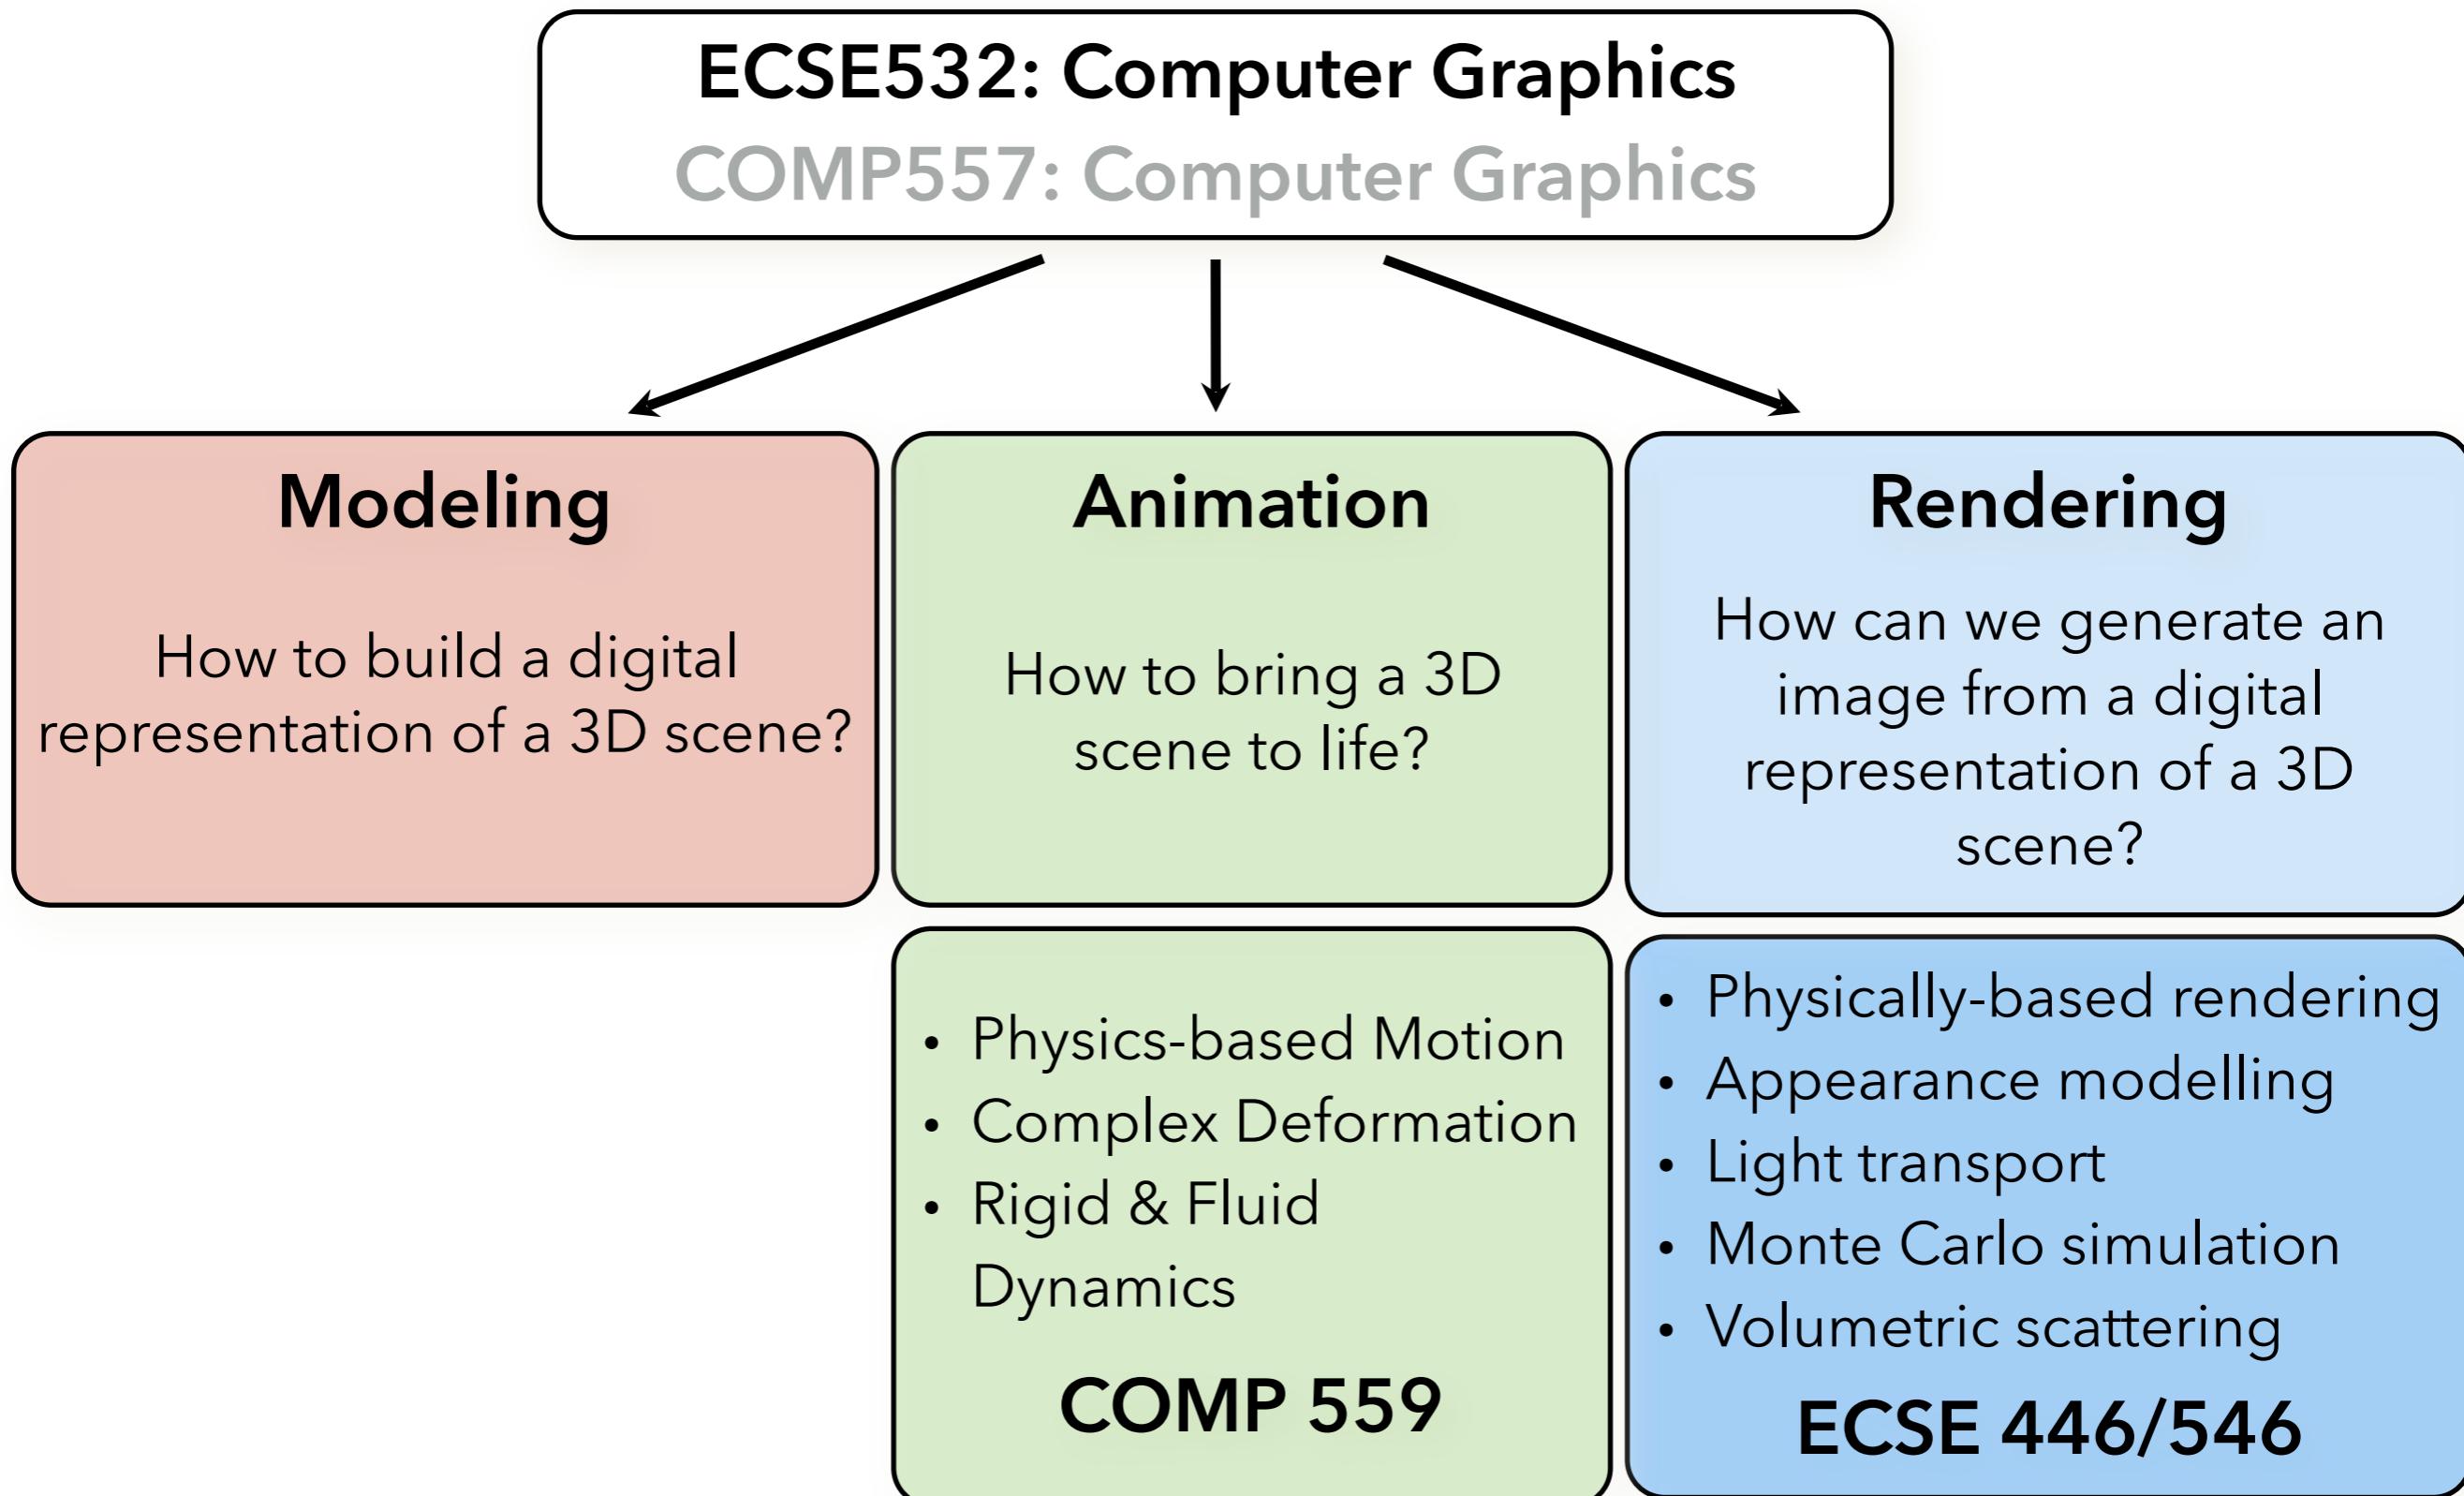
**Probability & Statistical Methods**

**Applied Numerical Methods**

# Relation to curriculum



# Relation to curriculum



# Motivating questions

---

Why do things look the way they do?

How can we generate realistic images?

# Motivation



wikipedia

# Motivation



<http://www.math.msu.edu/~jech/>

# Motivation

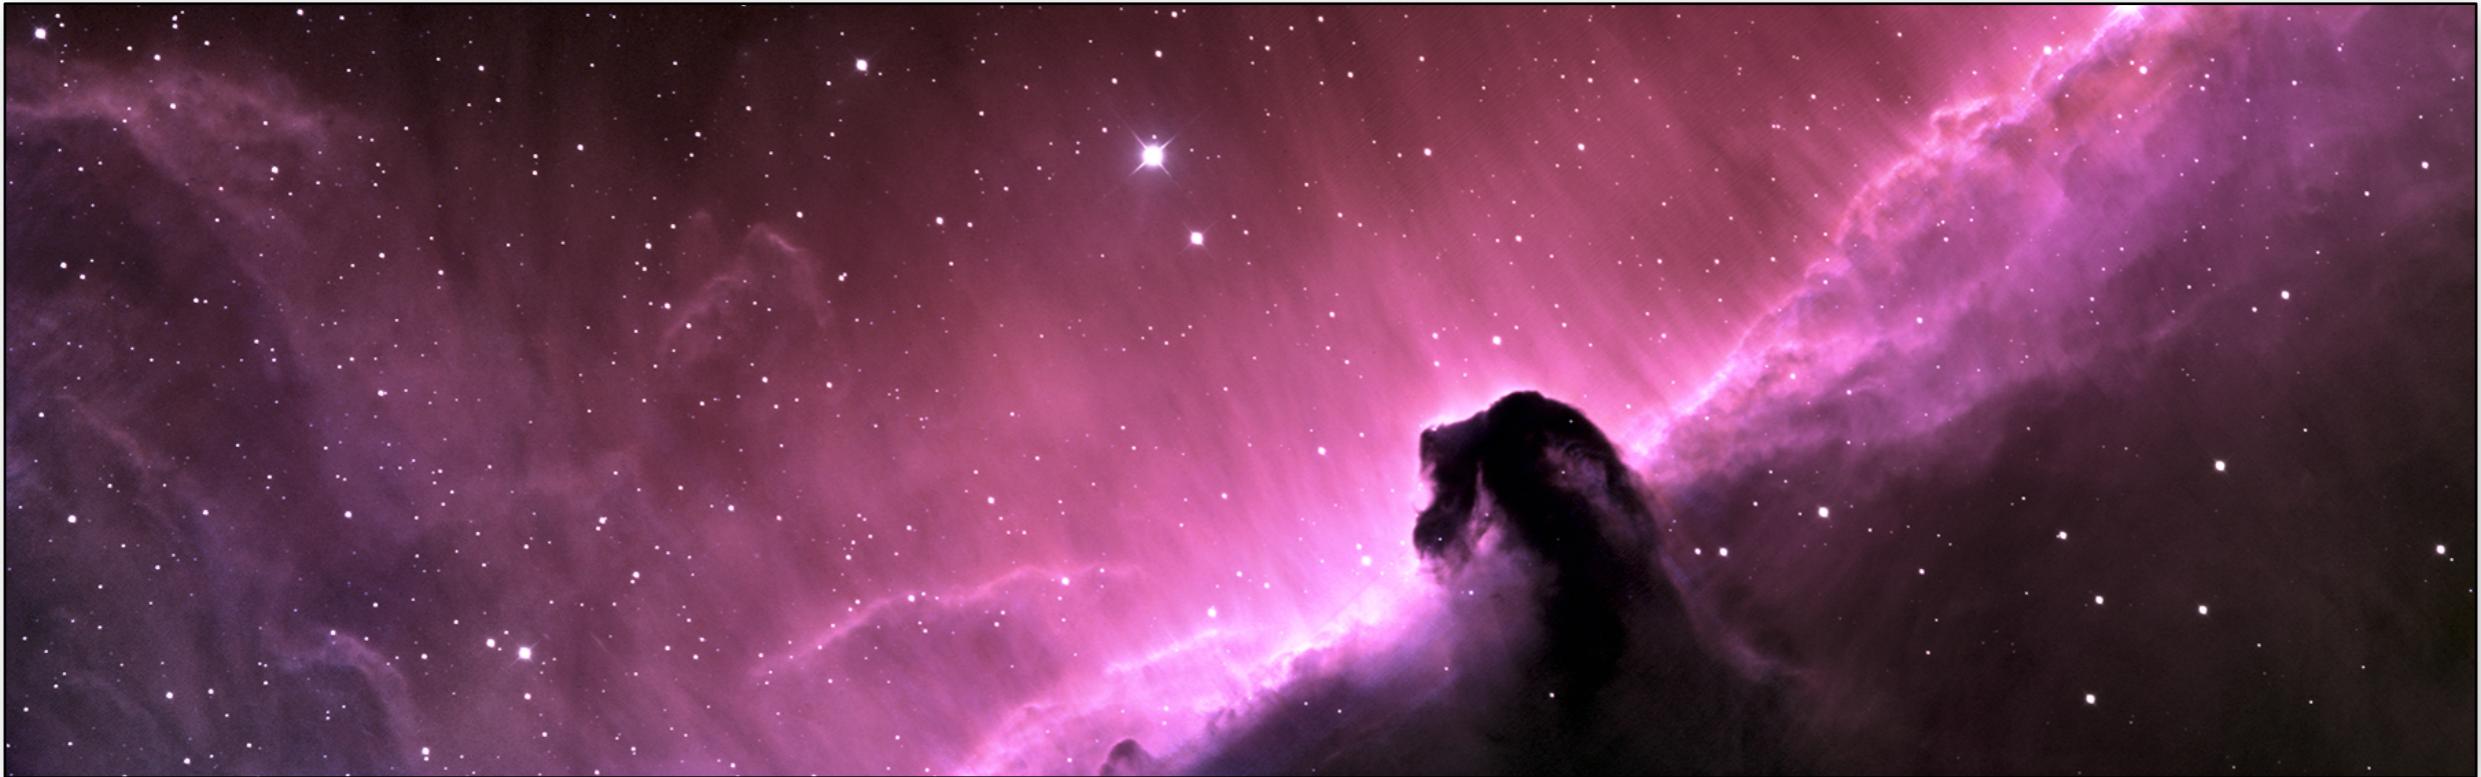


# Motivation



<http://mev.fopf.mipt.ru>

# Motivation



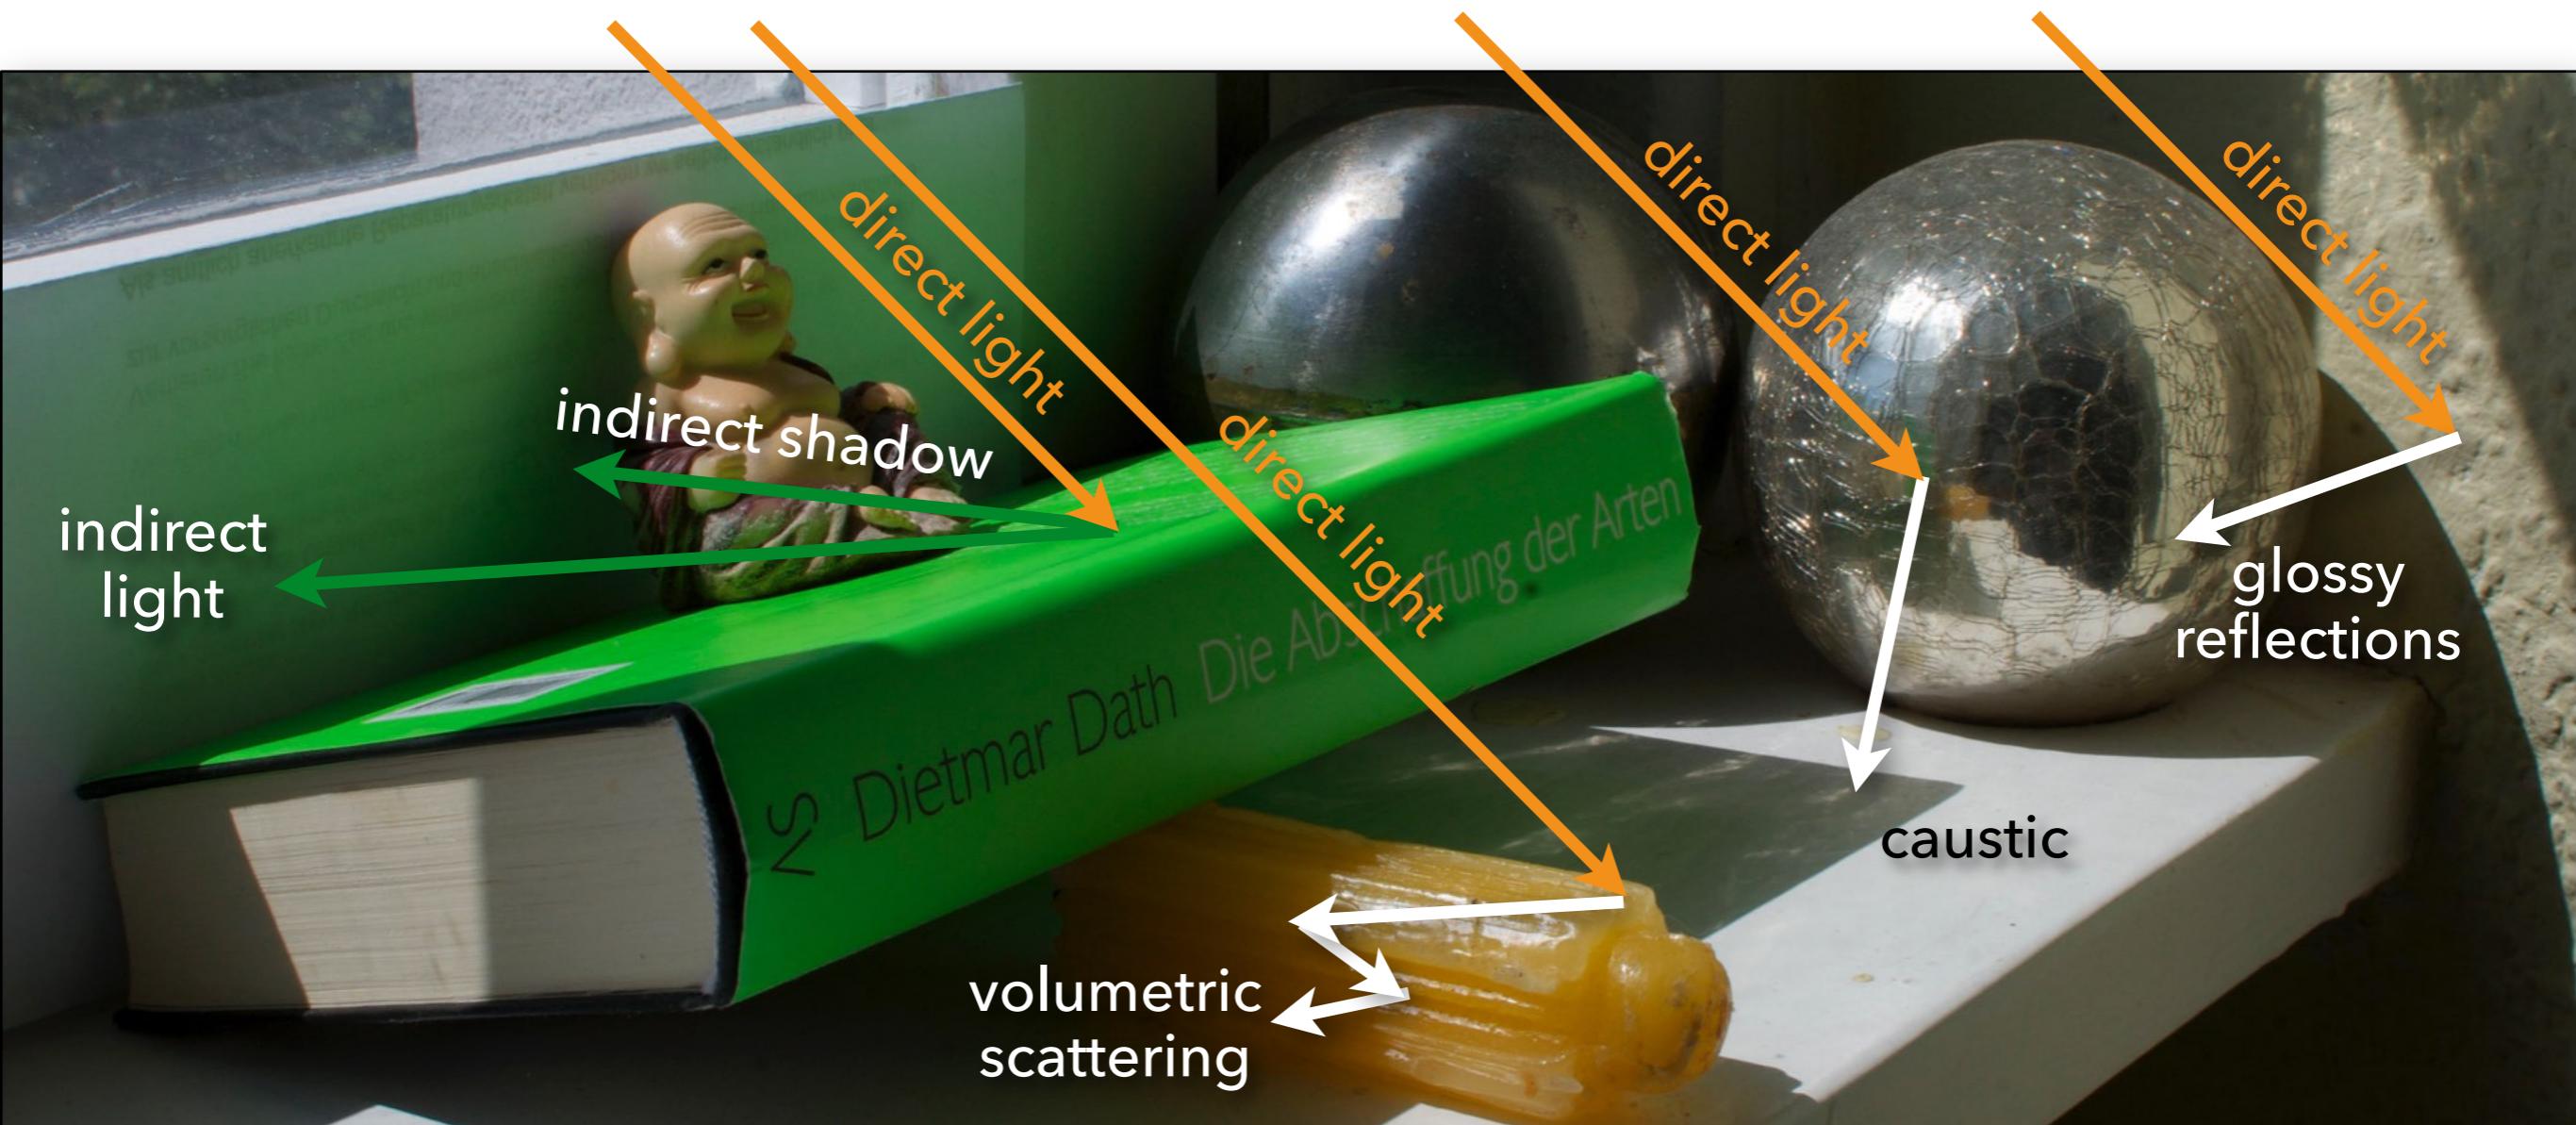
T.A.Rector (NOAO/AURA/NSF) and the Hubble Heritage Team (STScI/AURA/NASA)

# Motivation



Wojciech Jarosz

# Light transport in the real world



After [Ritschel et al. 2011]

# Visual effects



# Animated films



# Video games



# Architectural visualization



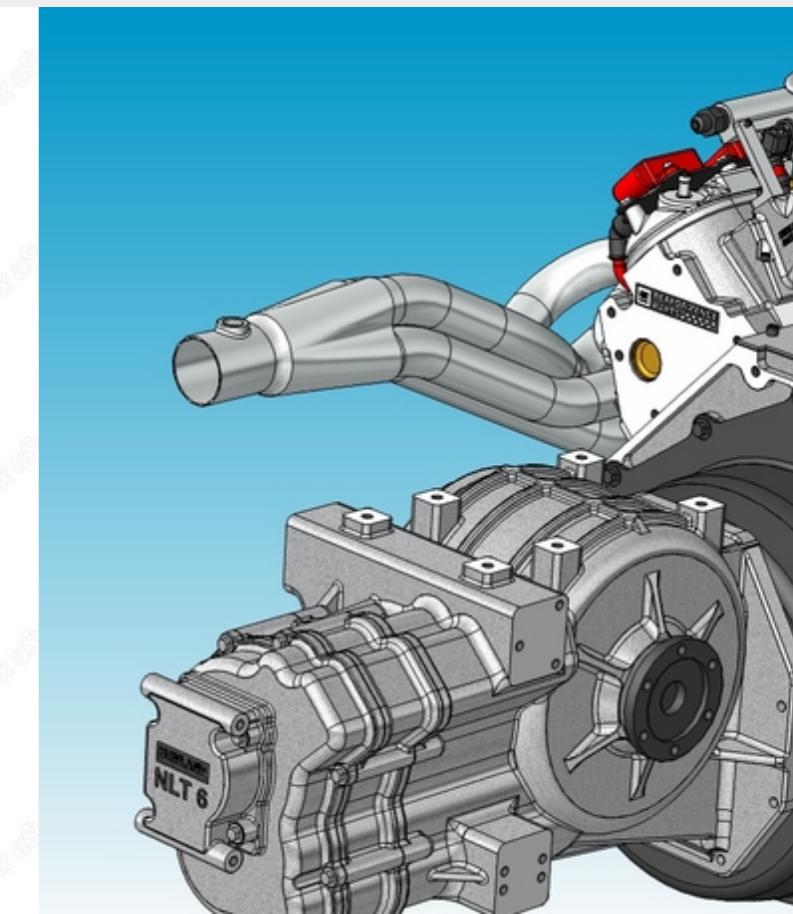
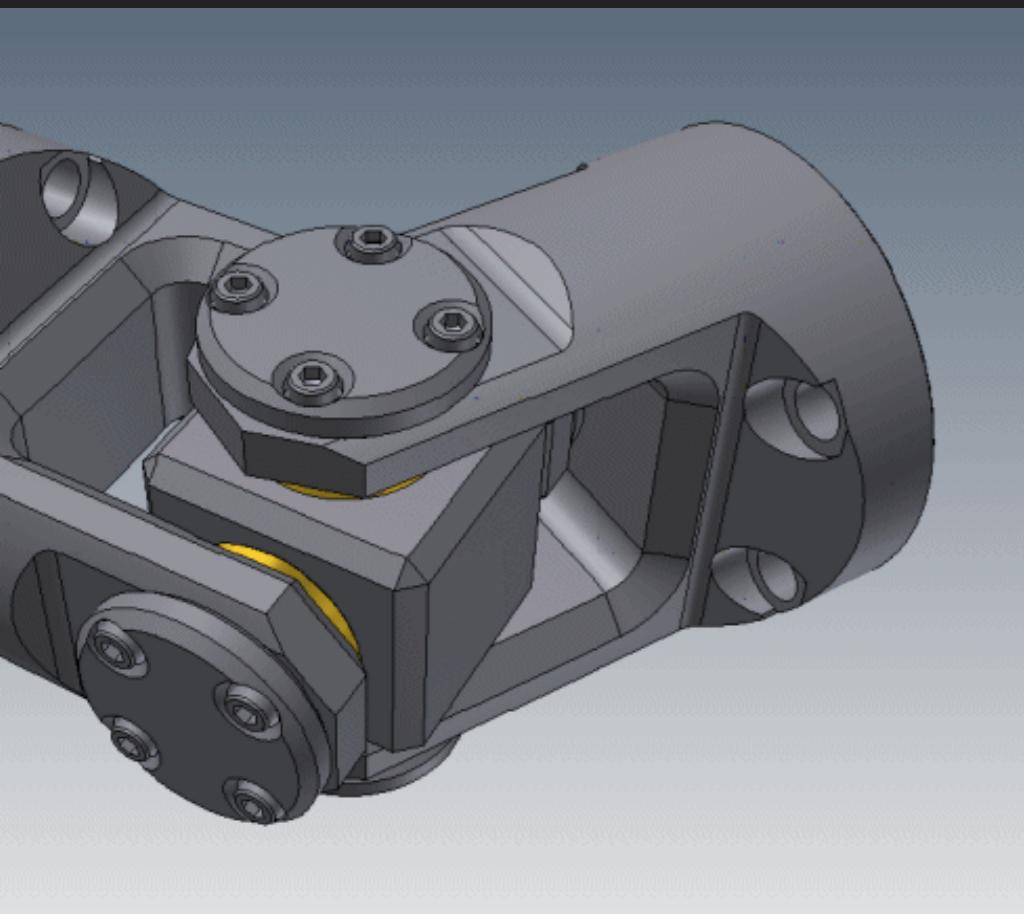
# Architectural visualization



# Architectural visualization



# Computer aided design/visualization



# Product visualization



# Product visualization



# Product visualization



# Product visualization



# Advertising & E-commerce



A screenshot of a product page for the Moto X smartphone. The page features a blue header with navigation links like 'ACCESSORIES', 'SHOP ALL', and 'GET HELP'. Below the header, there's a search bar and a sign-in button. The main content area shows the phone from both front and back views. The front view displays a pink floral wallpaper and the time '11:35'. The back view shows the wood-grain pattern. To the right, the text 'moto x' is displayed, followed by 'From \$124.99 w/ a 2-yr contract. From \$524.99 off contract'. A 'Buy' button is prominent, along with 'Compare', 'Share', and 'Help' links. At the bottom, there are '3D' and search icons, and a 'Discover More' link.

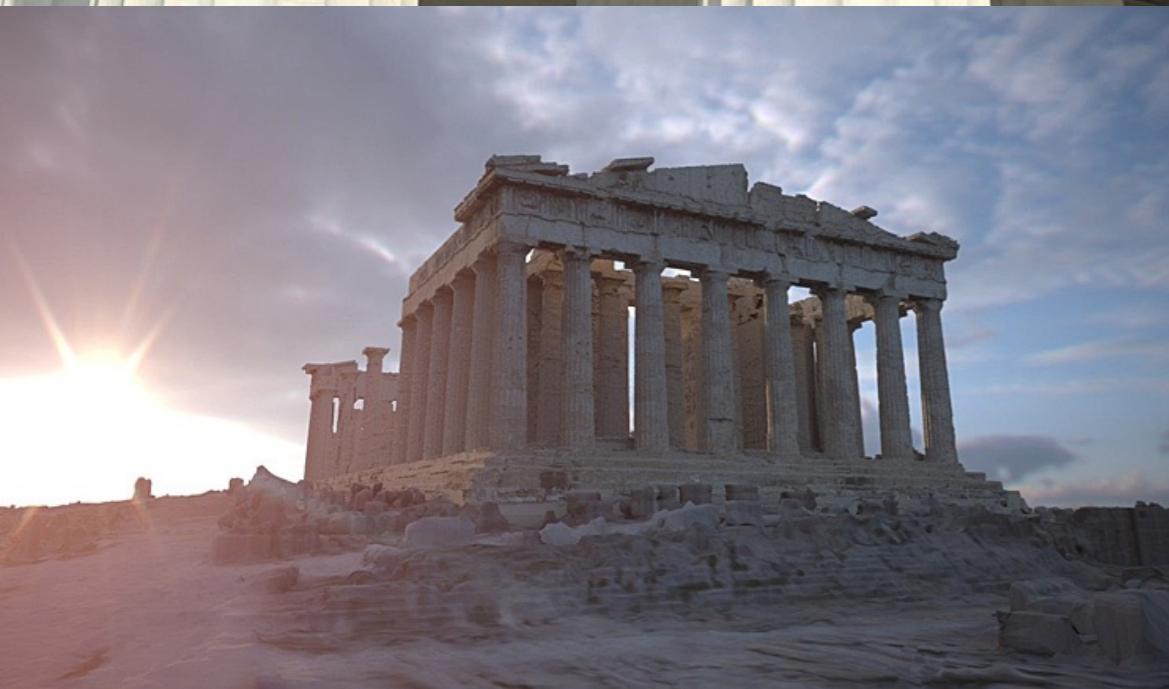
Start with one of these



# Advertising & E-commerce



# Cultural heritage



# Scientific visualization (& Visual Effects)



# Digital Fabrication



# Graphic arts



# What is this class about?

Algorithms for realistic image synthesis

Simulating light transport (global illumination)

Modelling appearance (simulating materials)

Understanding why things look the way they do:

- Why is the sky blue?
- Why is the grass green?
- Why does metal look different than marble?

# What you should already know

Basic concepts of algorithms and data structures

Calculus

- you will need to compute some derivatives & integrals

Linear algebra

- should know how to perform basic vector/matrix ops.

**Programming**

Basic Concepts of Computer Graphics\*

- basics of 3D spaces + objects\*
- basics of ray tracing + shading/lighting concepts\*

# What you will learn

By the end of course you should:

- understand advanced concepts in rendering, light transport, and appearance modelling
- be able to program both **interactive** and **offline** graphics applications
- have a better sense of:
  - good software engineering practices
  - applied numerical methods and analysis
  - hardware-accelerated graphics

# Literature

## Reference Textbook:

- Pharr & Humphreys.

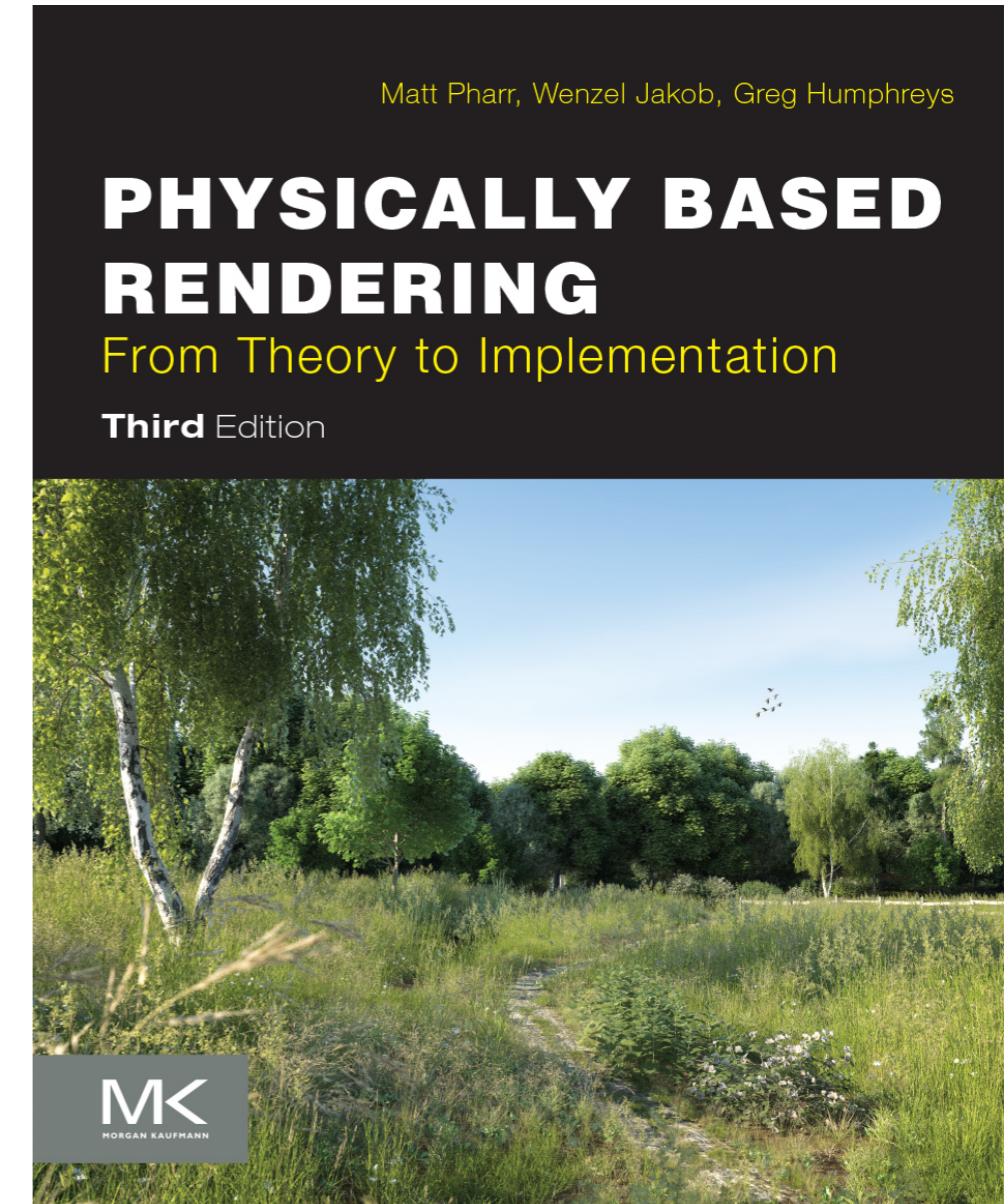
***Physically Based  
Rendering:***

***From Theory to  
Implementation,***

3rd edition (2016).

- digital e-book version available from McGill IPs

- see link on course website



# Literature (optional)

Get up to speed on ray tracing basics quickly:

- 3 kindle mini-book series by Pete Shirley



**Ray Tracing: in one weekend**

**Ray Tracing: the next week**

**Ray Tracing the rest of your life**

# Tentative Schedule

Week	Lecture Topics	Checkpoints
1	Rendering Systems and Radiometry	
2	Appearance Modelling and the Reflection Equation	
3	Monte Carlo Methods	<b>Assignment 1</b>
4	Direct Illumination	<b>Assignment 2</b>
5	Advanced Monte Carlo for Variance Reduction	
6	<b>[THANKSGIVING]</b> Path-space Light Transport Formulation	<b>Assignment 3</b>
7	Density Estimation	

# Tentative Schedule

Week	Lecture Topics	Checkpoints
8	Many Light Methods and Hybrid Rendering Systems	
9	Basis-space Representations for Interactive Shading	<b>Assignment 4</b>
10	Volumetric Participating Media	
11	Advanced Monte Carlo and Density Estimation in Volumes	<b>Assignment 5</b>
12	Markov Chain Monte Carlo for Path-space Exploration	
13	Subsurface Scattering Approximations	<b>Project*</b>

# Schedule – Coding Viewpoint

Week	Lecture Topics	Checkpoints
1	Rendering Systems and Radiometry	
2	Appearance Modelling and the Reflection Equation	
3	Monte Carlo Methods	<b>Assignment 1</b>
4	Direct Illumination	<b>Assignment 2</b>
5	Advanced Monte Carlo for Variance Reduction	
6	<b>[THANKSGIVING]</b> Path-space Light Transport Formulation	<b>Assignment 3</b>
7	Density Estimation*	

# Schedule – Coding Viewpoint

Week	Lecture Topics	Checkpoints
8	Many Light Methods and Hybrid Rendering Systems	
9	Basis-space Representations for Interactive Shading	<b>Assignment 4</b>
10	Volumetric Participating Media	
11	Advanced Monte Carlo and Density Estimation in Volumes	<b>Assignment 5</b>
12	Markov Chain Monte Carlo for Path-space Exploration	
13	Subsurface Scattering Approximations	<b>Project*</b>

# Grading Scheme

## 0%: Participation

- I want an *interactive course experience*
  - Don't be shy! Ask questions. Start discussions.
  - Technical questions are fine, too. But not "how do I compile the code"-type questions.

Without organic participation, I will move to a "volunteered" participation model...

# Grading Scheme – ECSE 446

75%: Programming Assignments

- A1 = 10%
- A2 = 10%
- A3 = 17%
- A4 = 18%
- A5 = 20%

(up to) 15%: Bonus/Hacker Points

# Grading Scheme – ECSE 546

50%: Programming Assignments

- A1 = 5%
- A2 = 5%
- A3 = 12%
- A4 = 13%
- A5 = 15%

25%: Project\*

(up to) 15%: Bonus/Hacker Points

# Programming assignments

## Late policy:

- Submit on time
  - **-100%** for late submissions
- Follow submission instructions exactly
  - **-100%** for invalid submissions

**Zero tolerance for cheating!**

All tasks will be submitted using  
**myCourses**

# Exam

---

25%: one (1) final exam

- multiple choice
- technical questions based on assignments
- conceptual questions based on remaining topics (inclusive)

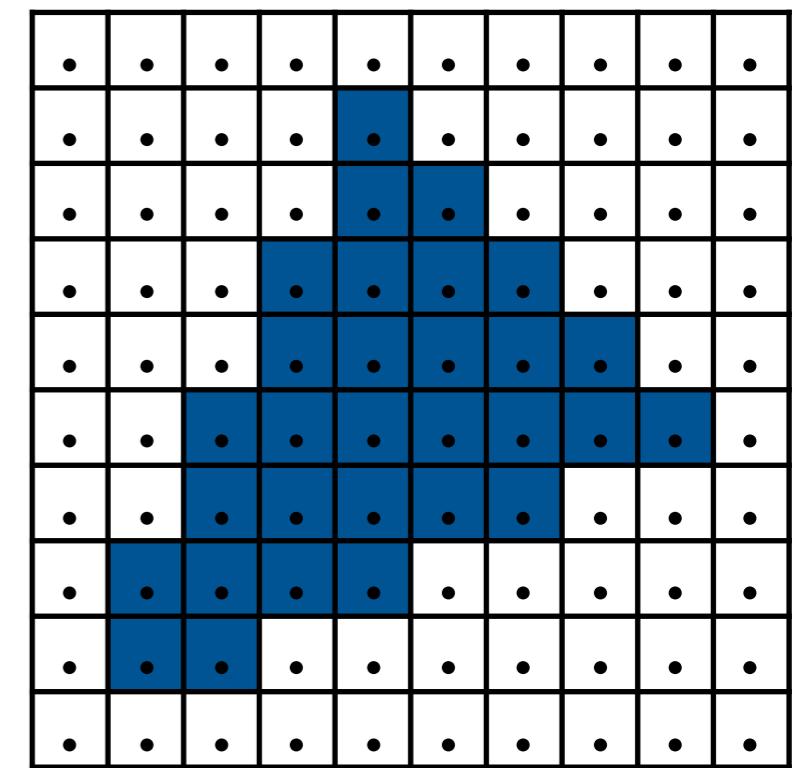
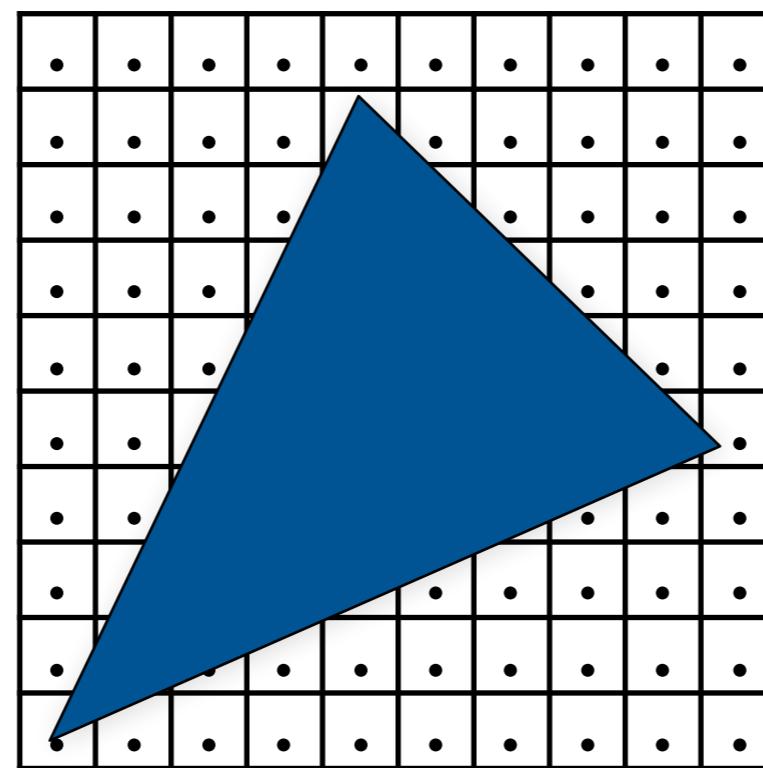
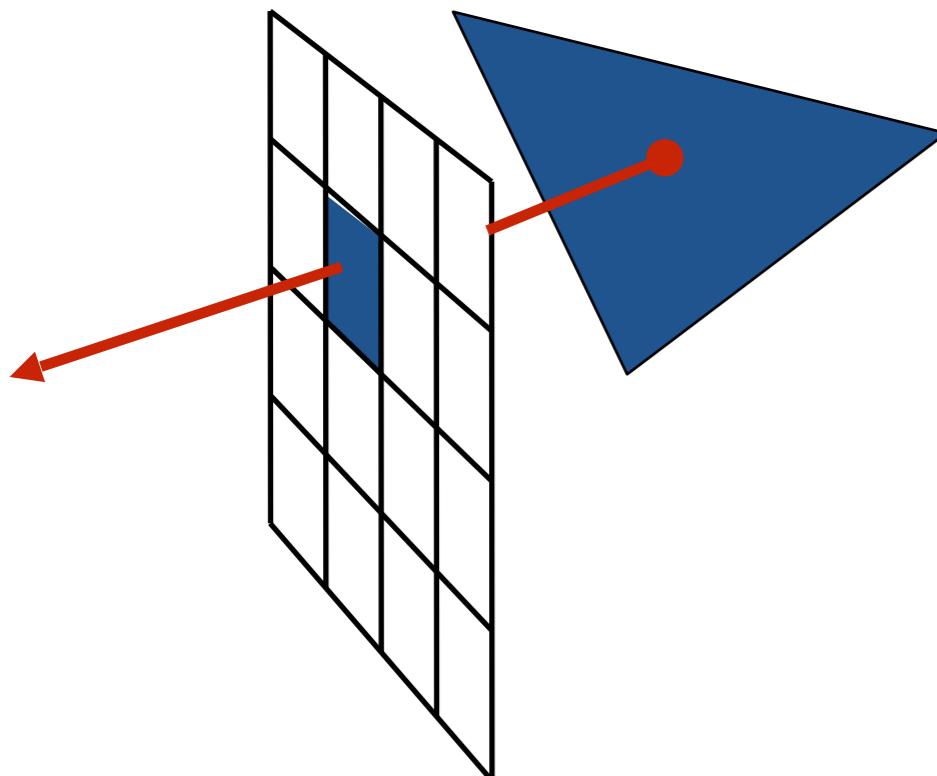
**Takeaway:** do the assignments yourself, attend class, and the exam will be easy

# Questions?

---

ECSE 446/546

# IMAGE SYNTHESIS



# RENDERING SYSTEMS

Prof. Derek Nowrouzezahrai  
[derek@cim.mcgill.ca](mailto:derek@cim.mcgill.ca)

# Rendering Systems

We provide you with a medium-sized codebase that includes basic plumbing:

- loading and representing 3D objects
- generating and saving 2D images
- rendering systems skeleton & support API

The code allows you to abstract **two** fundamental rendering system paradigms

- rasterization-based rendering, and
- ray tracing-based rendering

# Rendering Systems

Rasterization and ray-tracing are two ways of answering the following two questions:

1. what objects do I see from my point of view?
  - the eye- or *primary*- or *camera*-visibility problem
2. how do I compute their colors?
  - the *shading* or *lighting* problem

You'll build atop scaffolding we provide for both these paradigms, in the base code

- it's still a good idea to understand how they operate and their differences

# Rendering Systems

Rasterization and ray-tracing are two ways of answering the following two questions:

1. what objects do I see from my point of view?
  - the eye- or *primary*- or *camera*-*visibility problem*
2. how do I compute their colors?
  - the *shading* or *lighting problem*

You'll build atop scaffolding we provide for both these paradigms, in the base code

- it's still a good idea to understand how they operate and their differences

# Rendering Systems

Rasterization and ray-tracing are two ways of answering the following two questions:

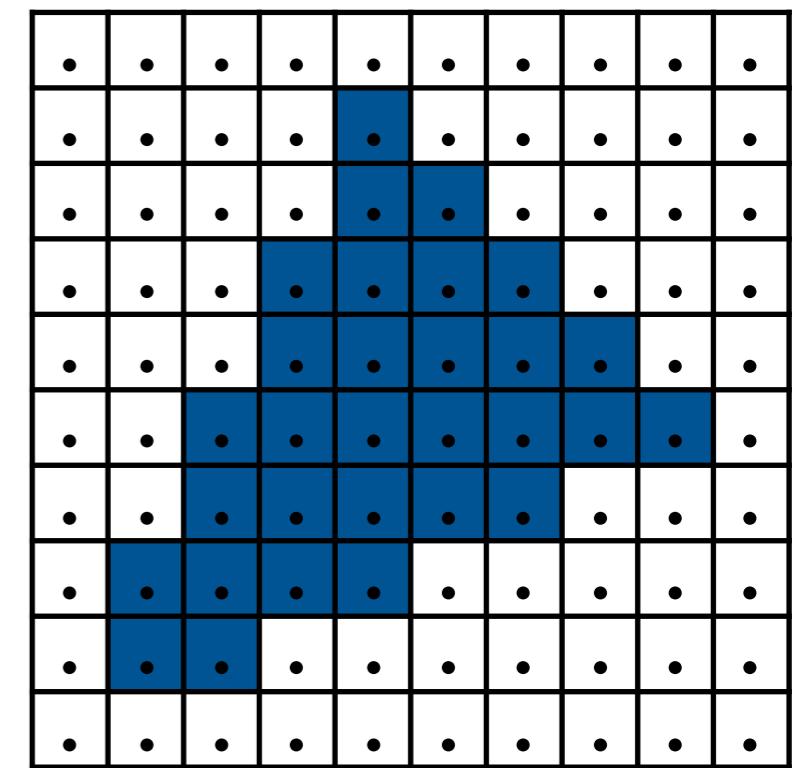
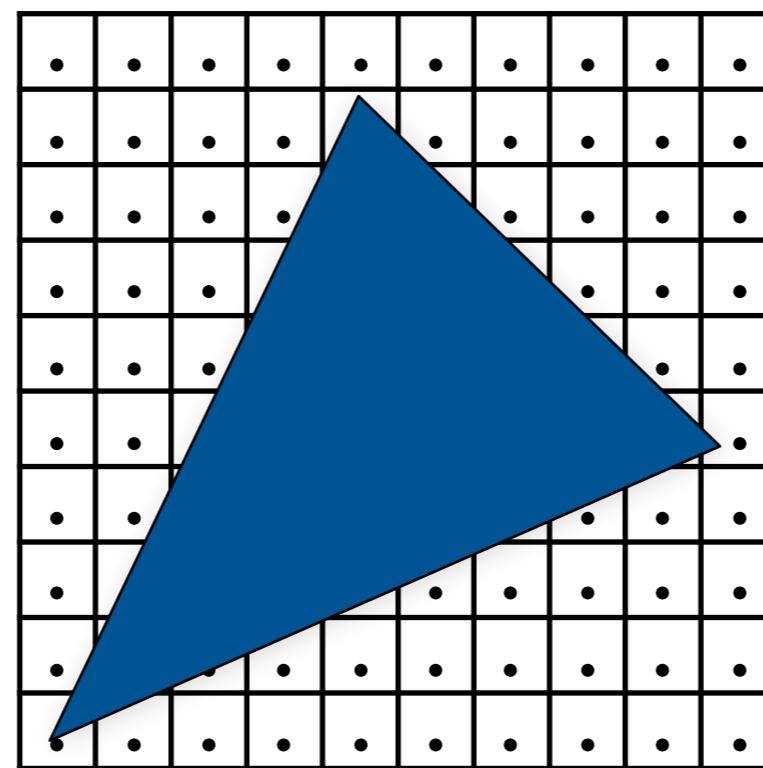
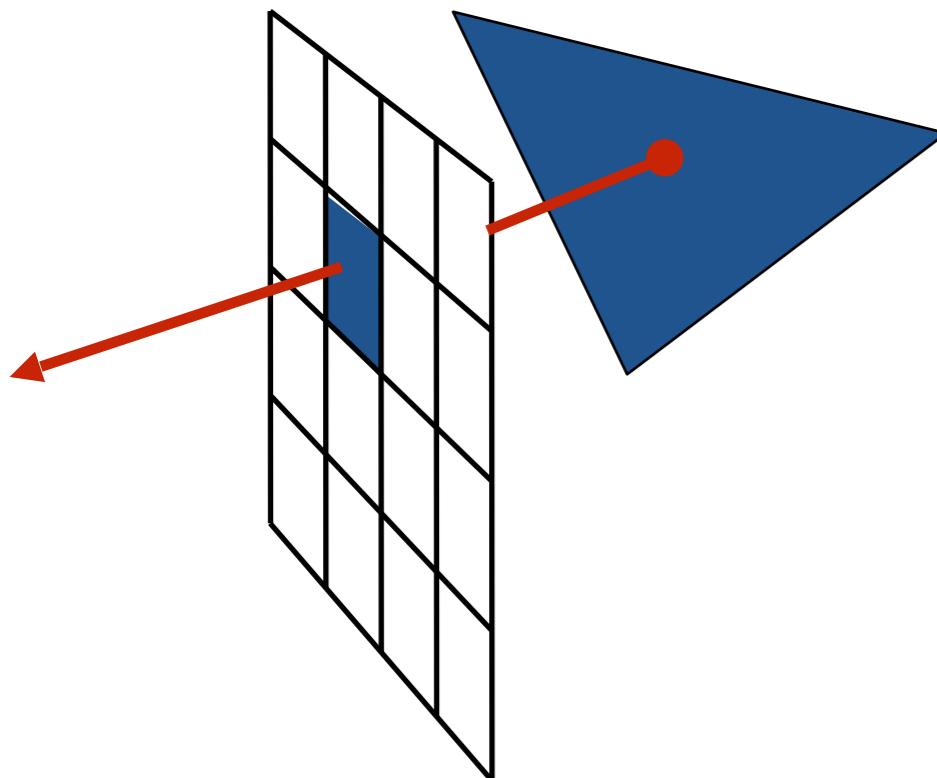
1. what objects do I see from my point of view?
  - the eye- or *primary*- or *camera*-*visibility problem*
2. how do I compute their colors?
  - the *shading* or *lighting problem*

You'll build atop scaffolding we provide for both these paradigms, in the base code

- it's still a good idea to understand how they operate and their differences

ECSE 446/546

# IMAGE SYNTHESIS

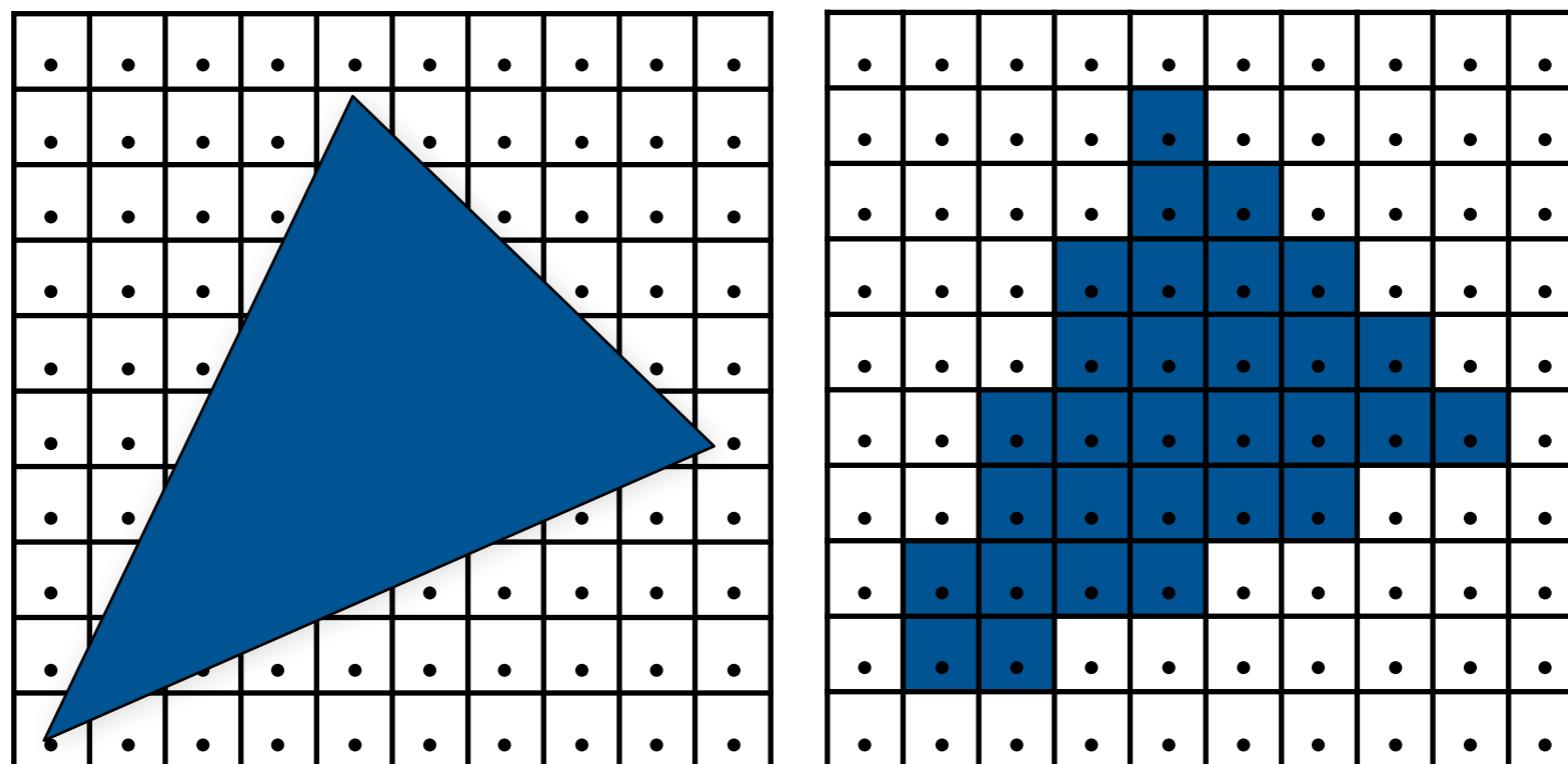


# RENDERING SYSTEMS

Prof. Derek Nowrouzezahrai  
[derek@cim.mcgill.ca](mailto:derek@cim.mcgill.ca)

ECSE 446/546

# IMAGE SYNTHESIS



## SYSTEMS 1 – RASTERIZATION

Prof. Derek Nowrouzezahrai  
[derek@cim.mcgill.ca](mailto:derek@cim.mcgill.ca)

# Rasterization Pipeline – GPUs

GPUs have dedicated hardware\* that implements a programmable rasterization pipeline

This HW is exposed through various APIs:

- OpenGL
- DirectX
- Vulkan
- Metal

# Rasterization Pipeline

We'll outline the different steps that the GPU can be directed to perform in this pipeline...

... but first, a simple example...

# Activity: Modelling & Drawing a Cube

Goal: realistically draw a cube

Key questions:

- Modelling: how do we describe the cube?
- Rendering: how do we visualize this model?



# Modelling the Cube

Suppose our cube is...

- centered at the origin  $(0,0,0)$
- has dimensions  $2 \times 2 \times 2$
- edges are aligned with the x/y/z-axes

**Q:** What are the cube vertex coordinates?

## VERTICES

A:  $(1, 1, 1)$

B:  $(-1, 1, 1)$

C:  $(1, -1, 1)$

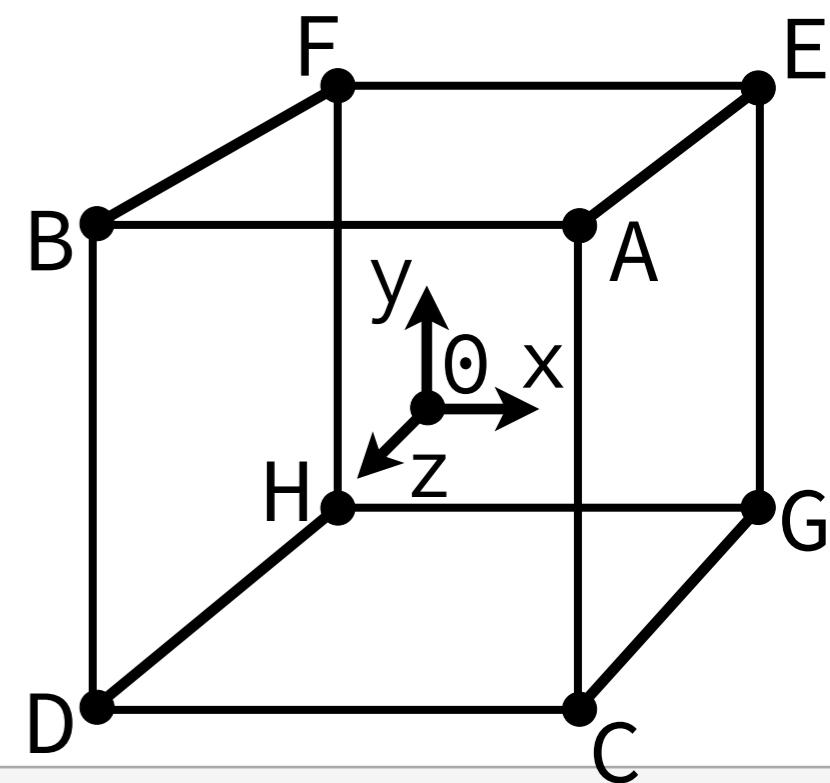
D:  $(-1, -1, 1)$

E:  $(1, 1, -1)$

F:  $(-1, 1, -1)$

G:  $(1, -1, -1)$

H:  $(-1, -1, -1)$



# From Modelling to Drawing

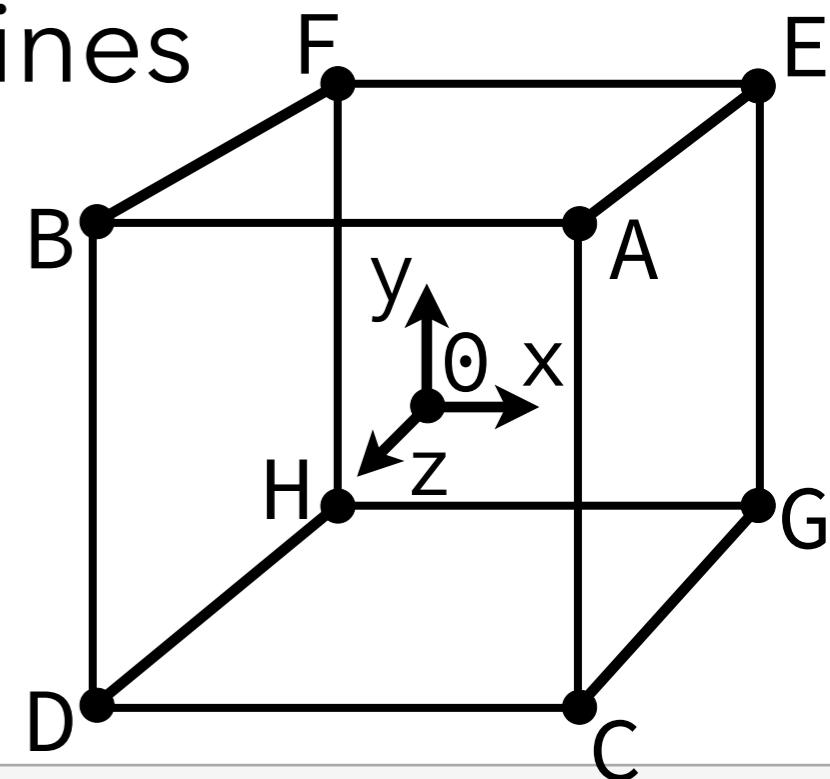
Now with a digital description of the cube

- How do we draw the 3D cube as a 2D image?

Basic strategy:

- map 3D vertices to 2D points in the image
- connect 2D points with straight lines

OK... but how?



# Perspective

After a slide by Alyosha Efros

MENU the ONION® Study: People Far Away From You Not Actually Smaller SEARCH TOP HEADLINES ▾

## Study: People Far Away From You Not Actually Smaller

NEWS  
August 22, 2013  
VOL 49 ISSUE 34  
Science & Technology

f t e



*Researchers say that, contrary to prior assertions, the subject above stands at equal height at left and at right, and does not grow smaller as he walks away from the camera.*

# Perspective

After a slide by Alyosha Efros

≡ MENU    Study: People Far Away From You Not Actually Smaller   SEARCH Q   TOP HEADLINES ▾

## Study: People Far Away From You Not Actually Smaller

NEWS  
August 22, 2013  
VOL 49 ISSUE 34  
Science & Technology

f   t   e



*Researchers say that, contrary to prior assertions, the subject above stands at equal height at left and at right, and does not grow smaller as he walks away from the camera.*

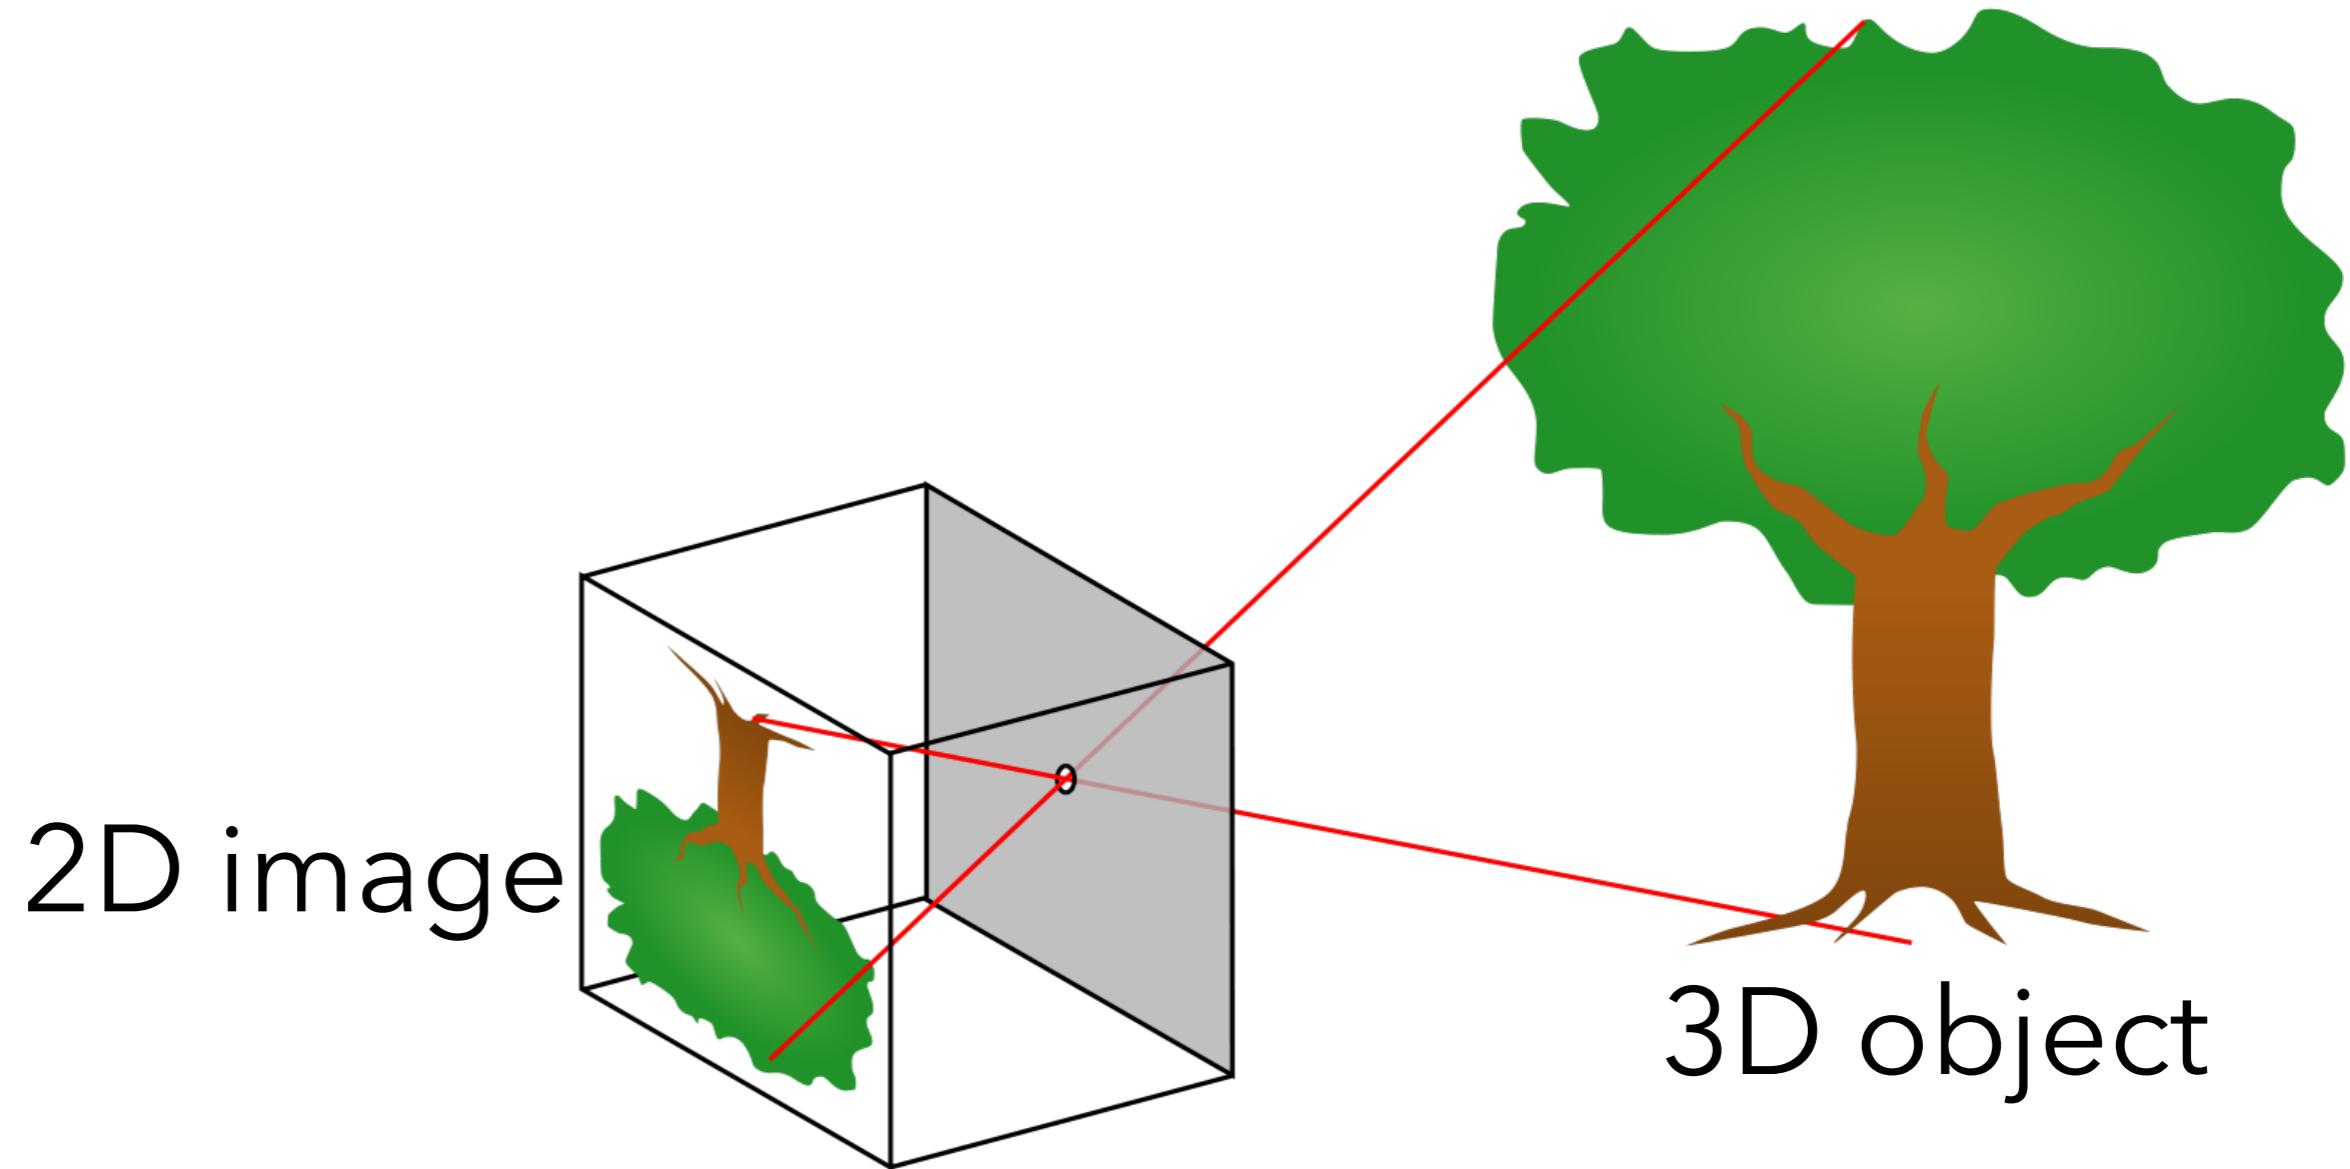
Objects appear smaller as they get further away (“perspective”)

Why does this happen?

# Perspective Projection

Consider simple (“pinhole”) camera:

After a slide by Keenan Crane

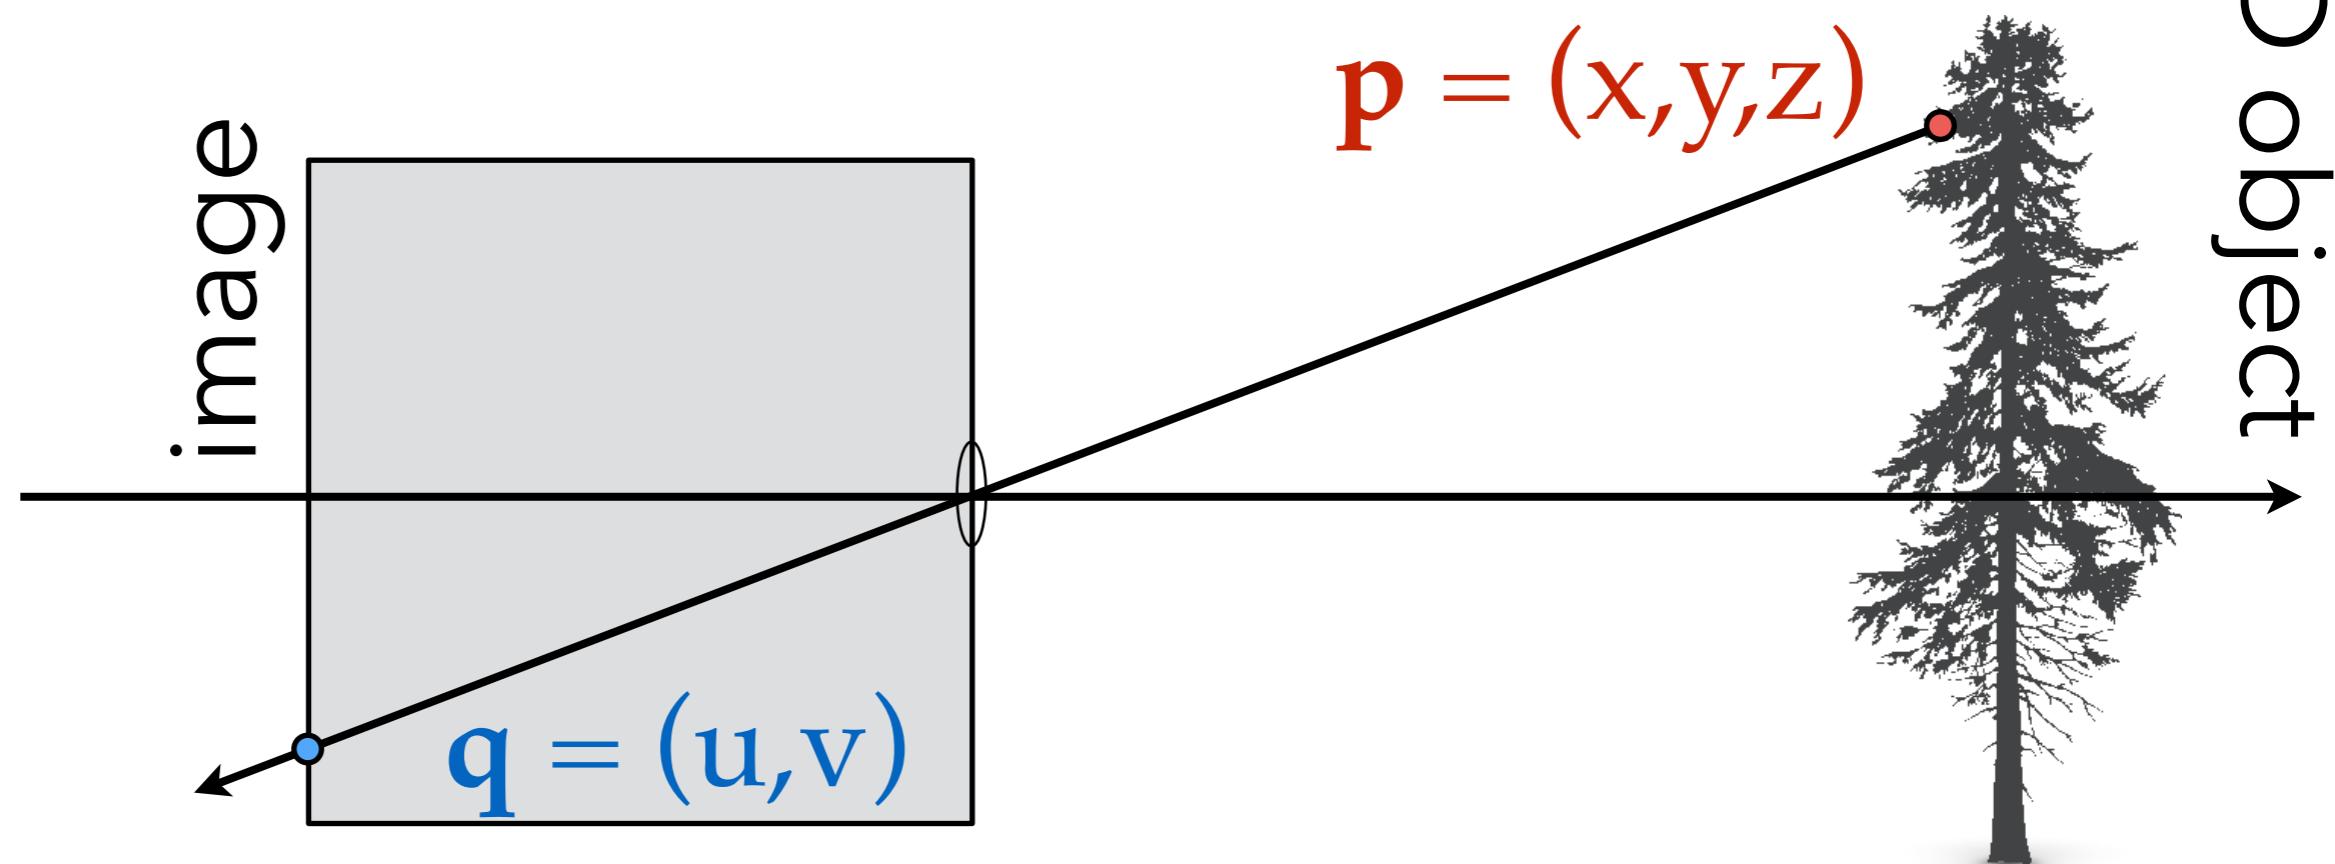


# Perspective Projection

Where does a point  $p = (x, y, z)$  end up on the image?

Let's call the image point  $q = (u, v)$

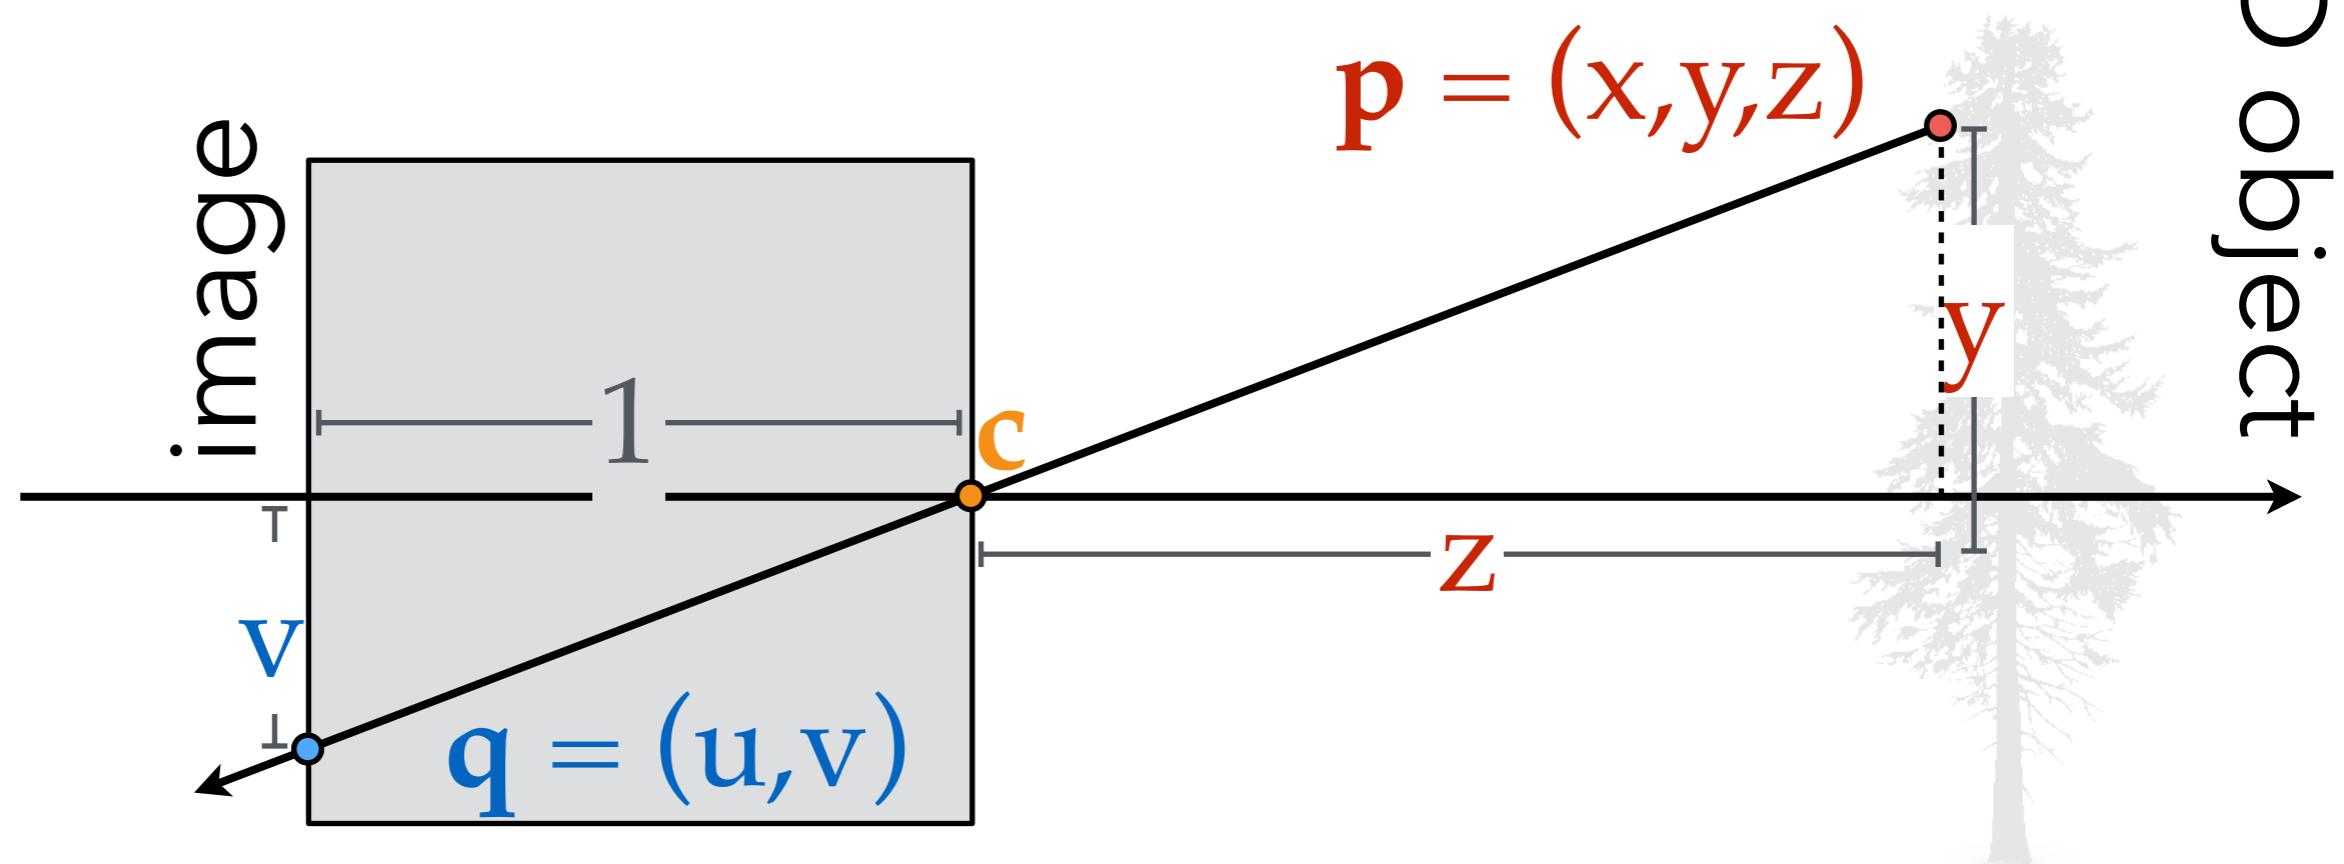
After a slide by Keenan Crane



# Perspective Projection

Assume camera at  $C$ , two similar triangles:

After a slide by Keenan Crane



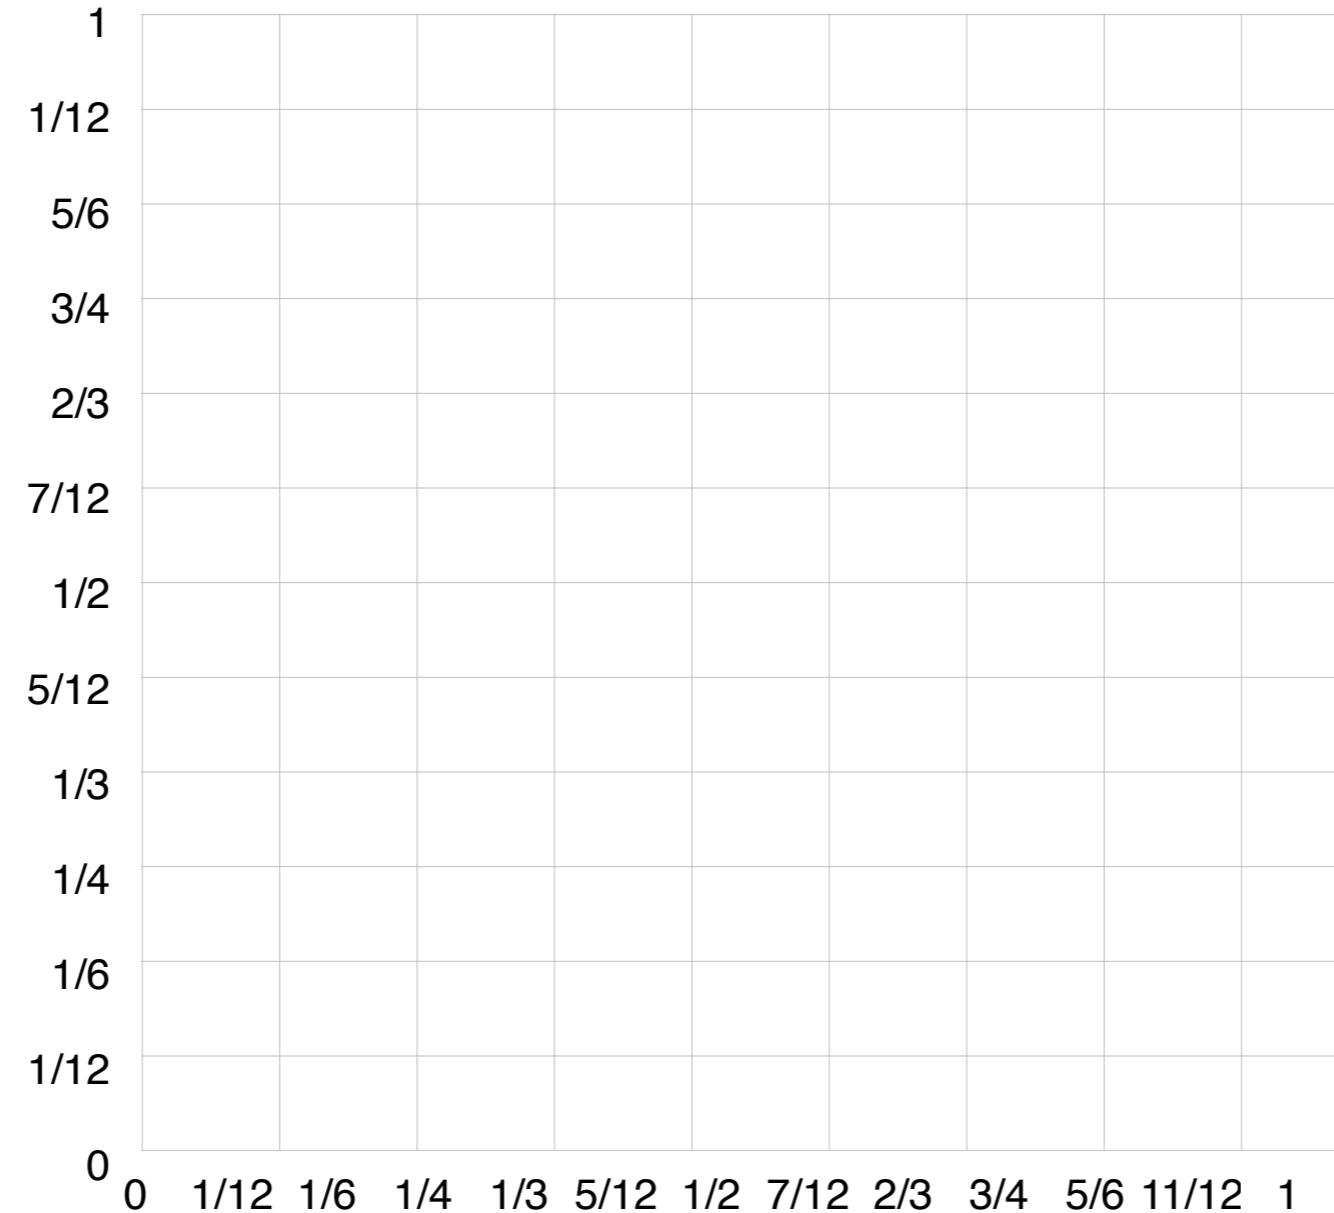
# Draw the Projected Cube

## Specification & Algorithm:

- assume camera is at  $\mathbf{c} = (2,3,5)$
- for each edge
  - convert  $(x,y,z)$  of both endpoints to  $(u,v)$ :
    1. subtract camera  $\mathbf{c}$  from vertex  $(x,y,z)$  to get  $(x',y',z')$
    2. divide  $(x',y')$  by  $z'$  to get  $(u,v)$  – write as a fraction
  - draw line between  $(u_1,v_1)$  and  $(u_2,v_2)$

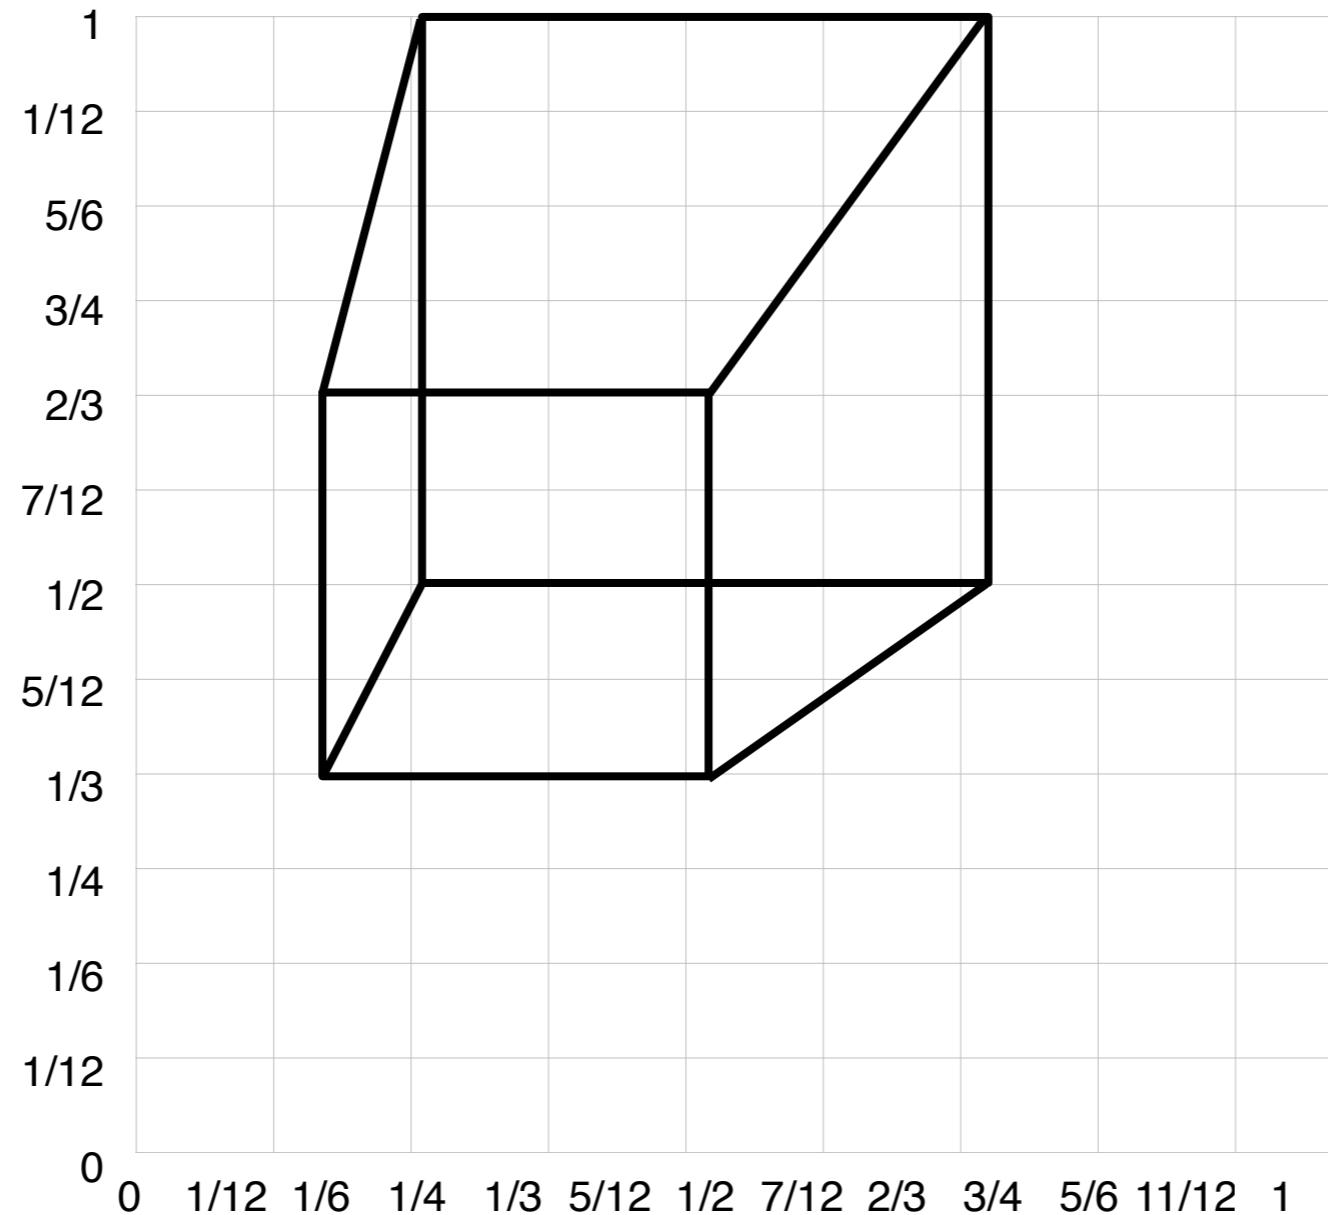
# Draw the Projected Cube

After a slide by Keenan Crane



# Draw the Projected Cube

After a slide by Keenan Crane



**2D coordinates:**

A: ( $?$ ,  $?$ )

B: ( $?$ ,  $?$ )

C: ( $?$ ,  $?$ )

D: ( $?$ ,  $?$ )

E: ( $?$ ,  $?$ )

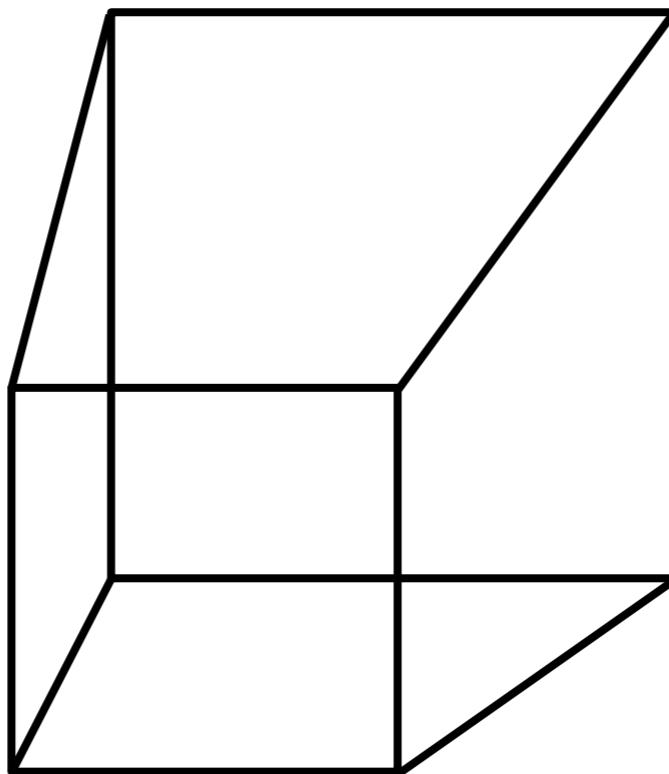
F: ( $?$ ,  $?$ )

G: ( $?$ ,  $?$ )

H: ( $?$ ,  $?$ )

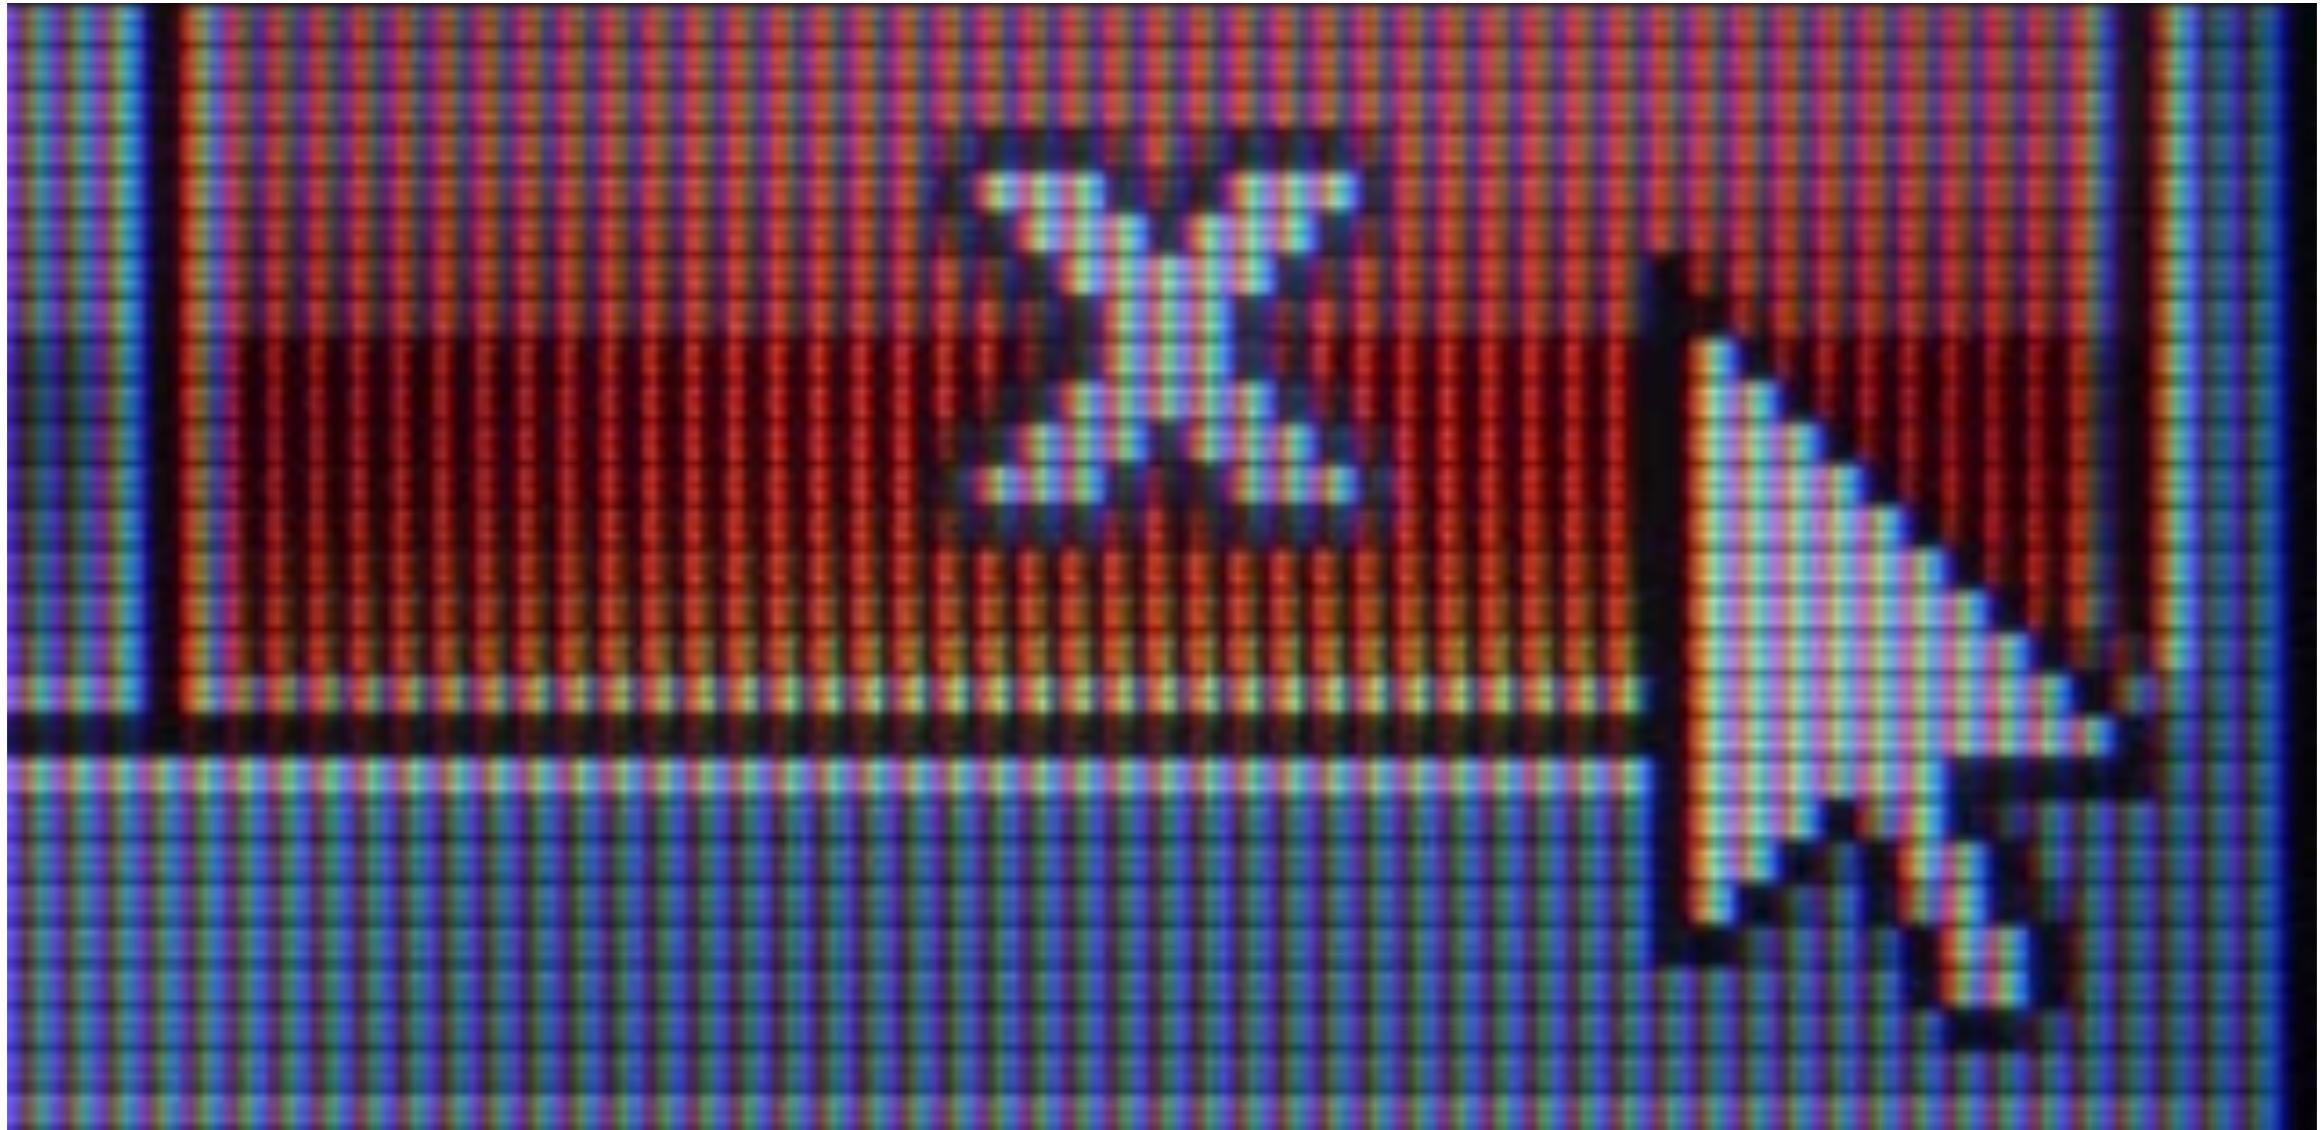
# But wait...

## How do we actually draw lines on a computer?



# Raster Graphics

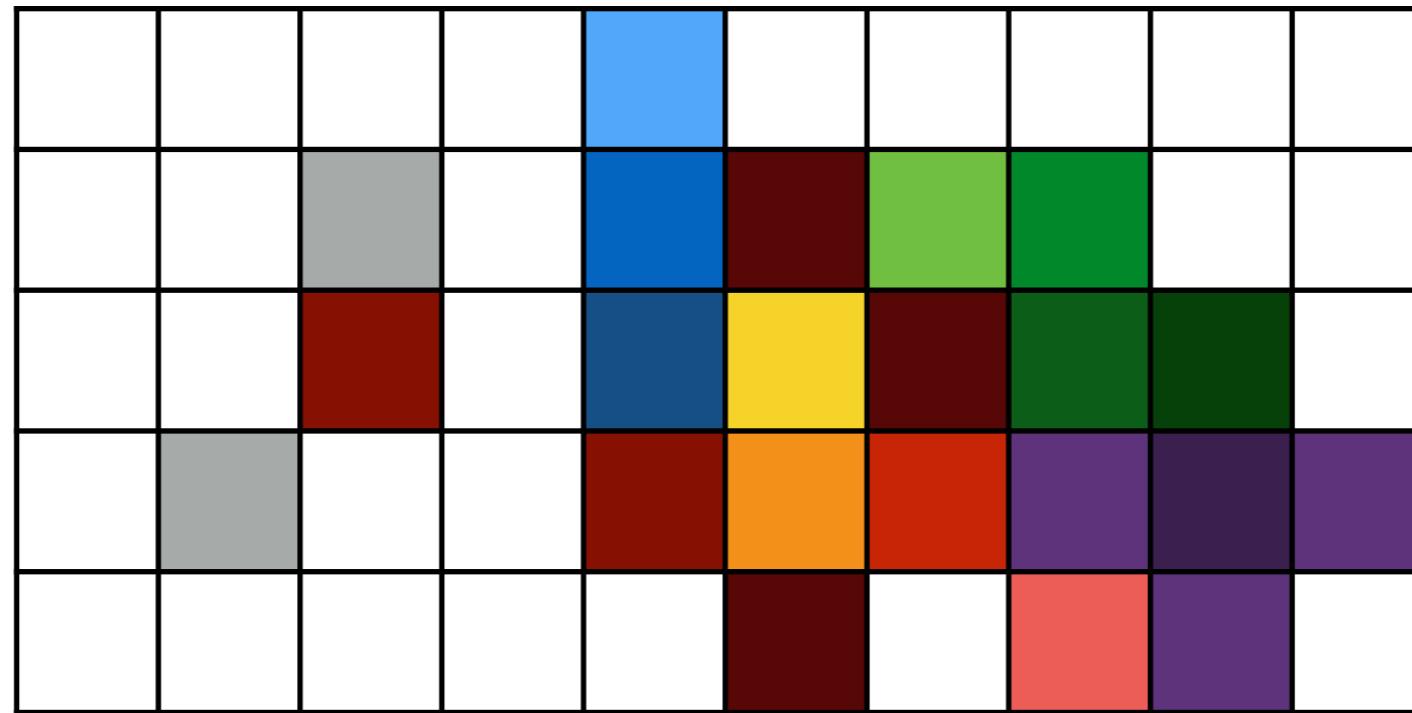
After a slide by Keenan Crane



# Raster Graphics

Common abstraction of a raster display:

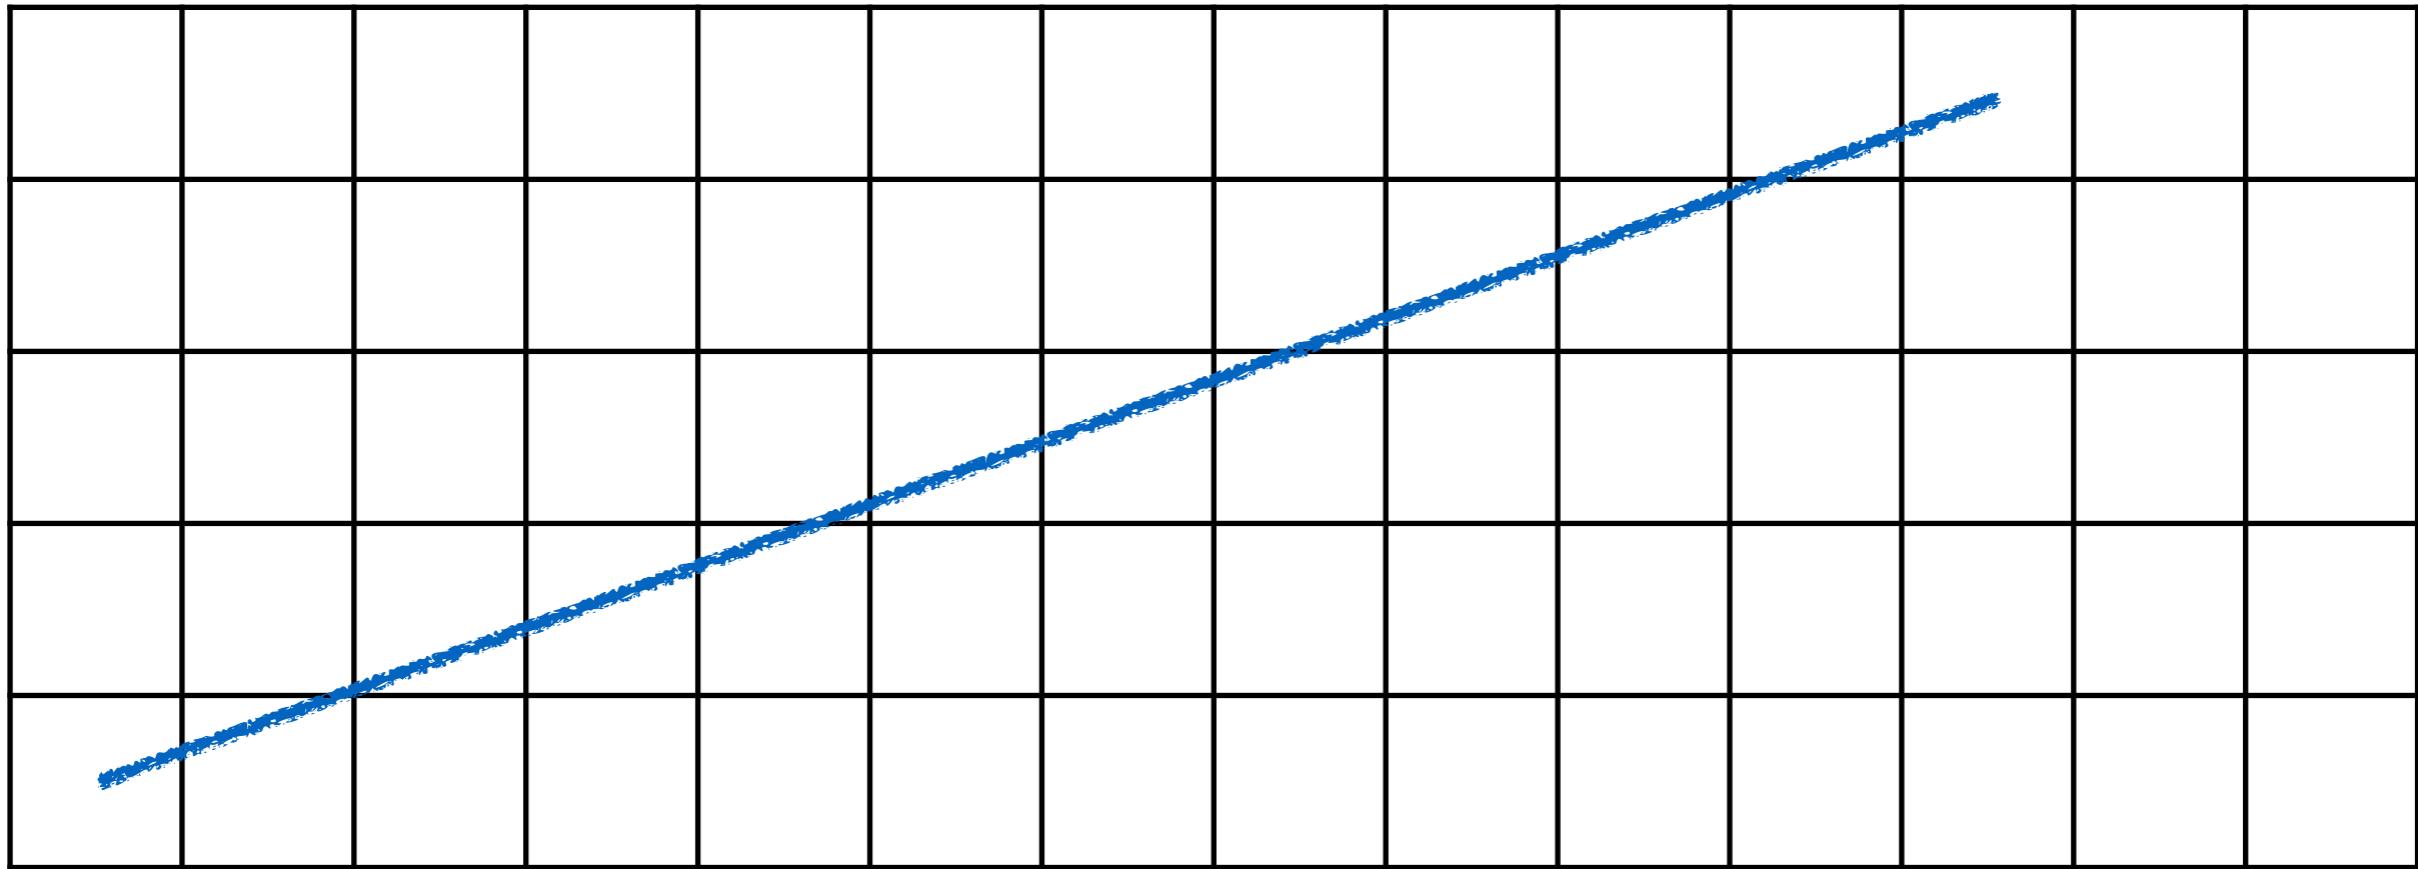
- Image represented as a 2D grid of “pixels” (picture elements)
- Each pixel can take on a unique color value



After a slide by Keenan Crane

# Which Pixels to Fill to Draw the Line?

**Rasterization** is the process of converting a continuous object to a discrete representation on a raster (pixel) grid

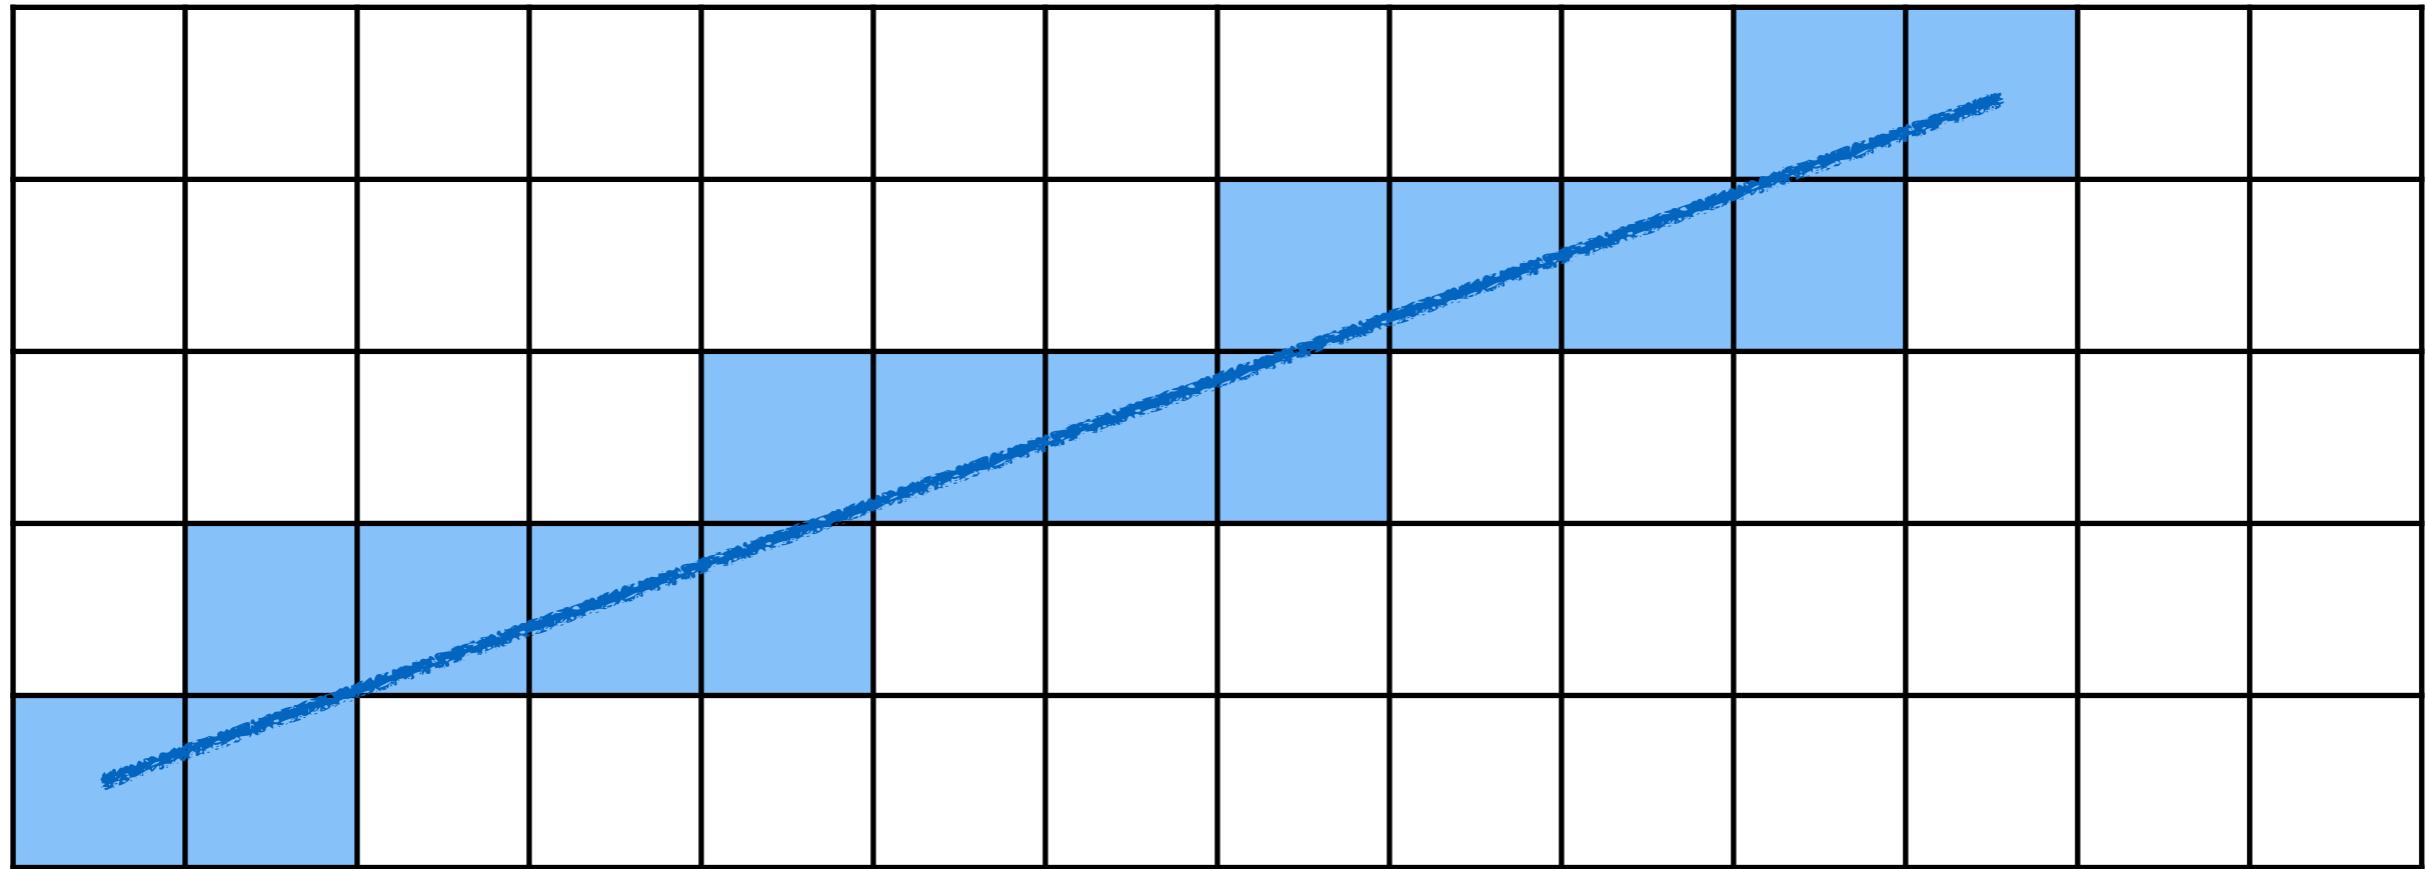


After a slide by Keenan Crane

# Rasterization Rules

Light up all pixels intersected by the line?

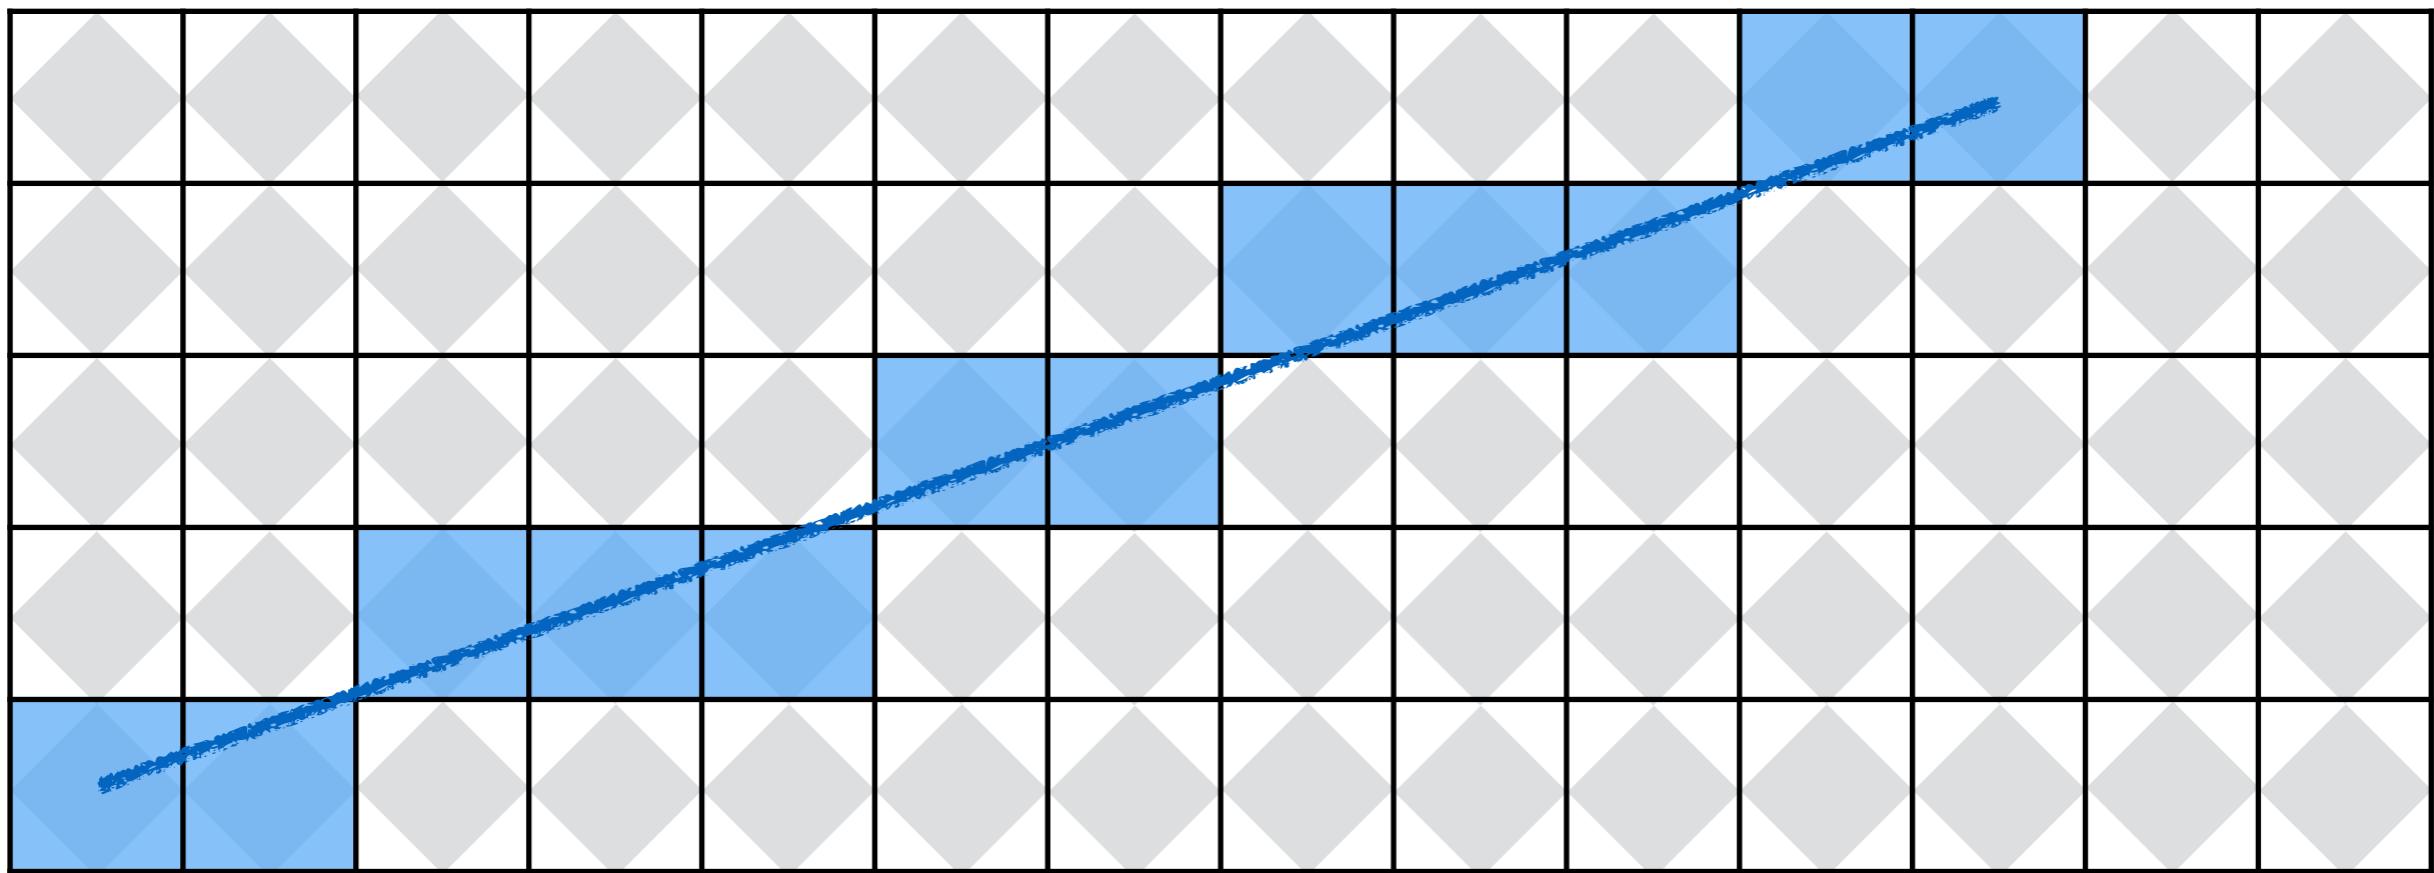
After a slide by Keenan Crane



# Rasterization Rules

Diamond rule (used by modern GPUs):  
light up pixel if line passes through  
associated diamond

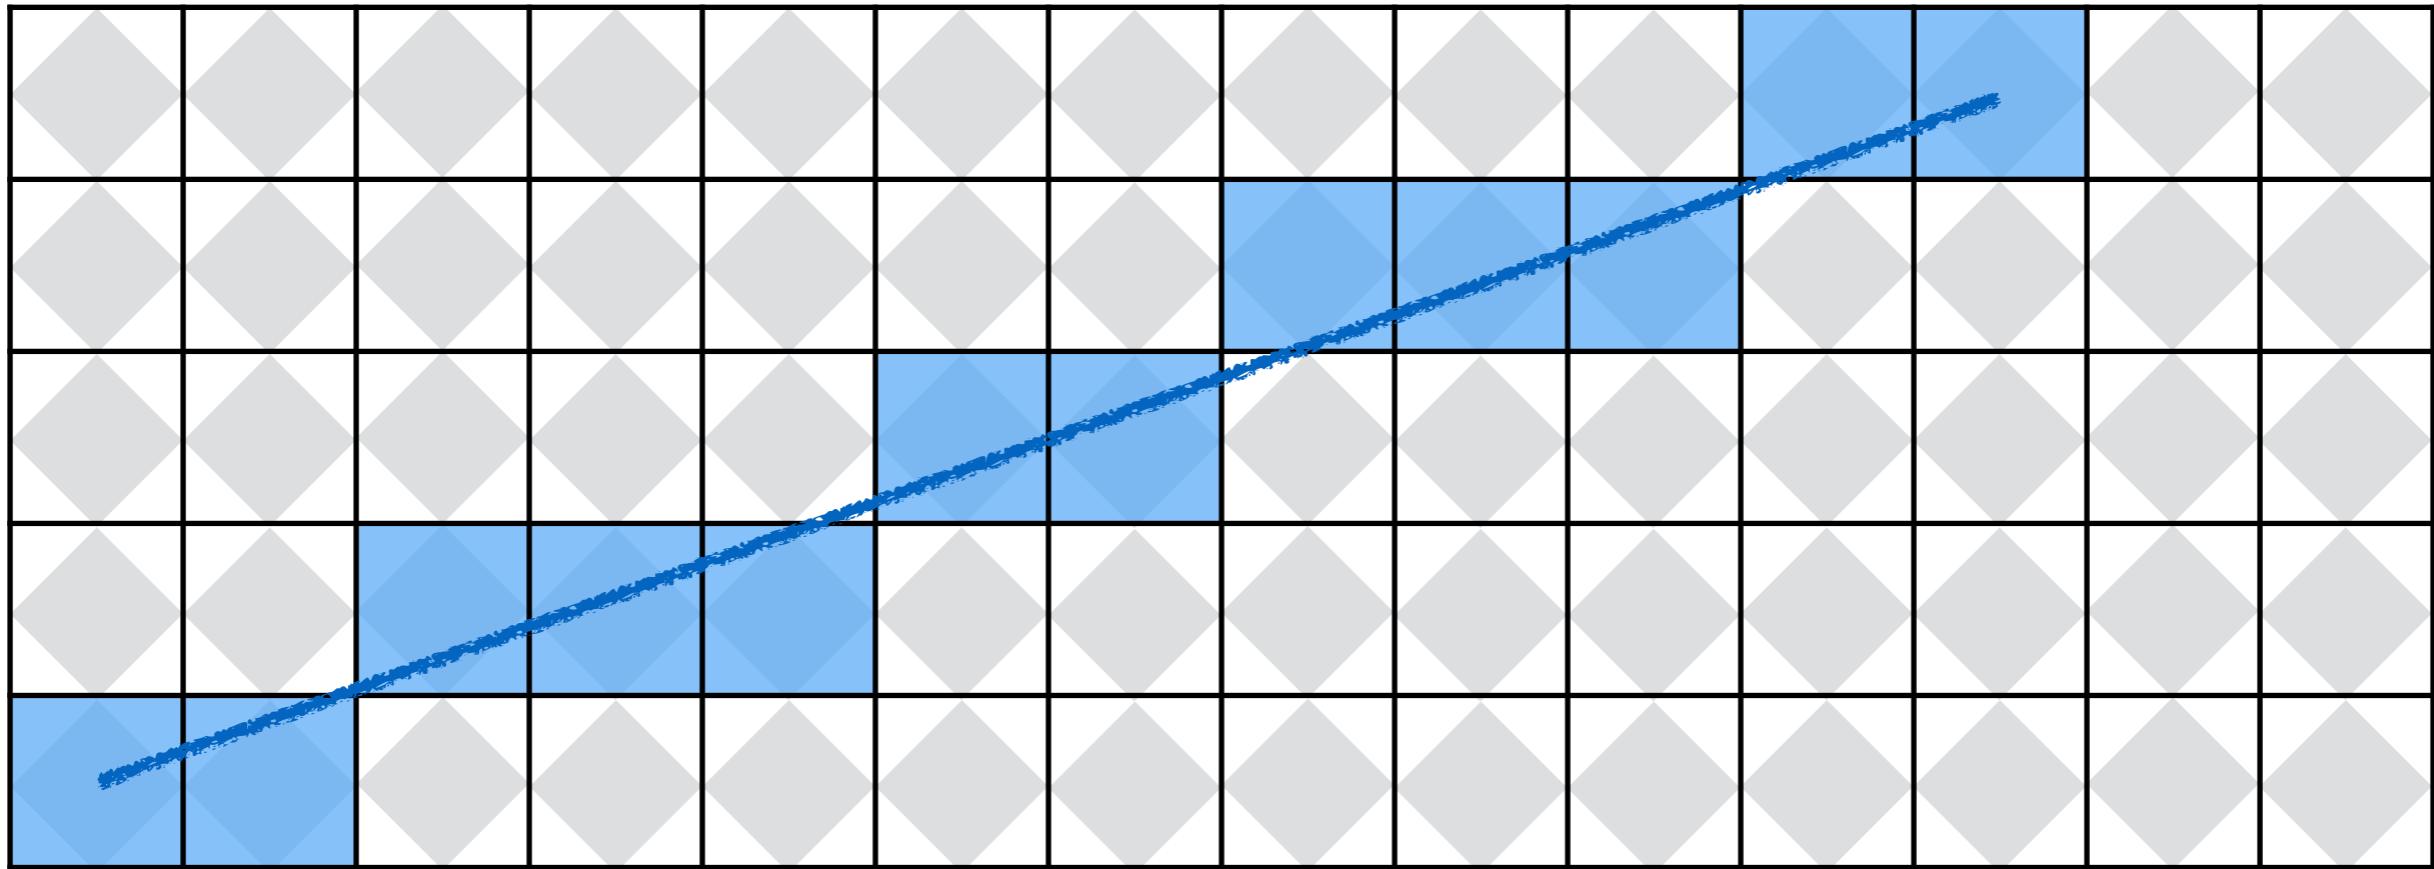
After a slide by Keenan Crane



# Rasterization Rules

Is there a right answer?

After a slide by Keenan Crane



# Line Rasterization

How do we find the pixels satisfying a chosen rasterization rule?

- Could check every single pixel in image to see if it meets the condition...
  - $O(n^2)$  image pixels vs. at most  $O(n)$  “lit up”
  - *must* be able to do better!
    - e.g., work proportional to number of pixels in the drawing of the line

# Incremental Line Rasterization

Let's say a line has integer endpoints:

$(u_1, v_1), (u_2, v_2)$

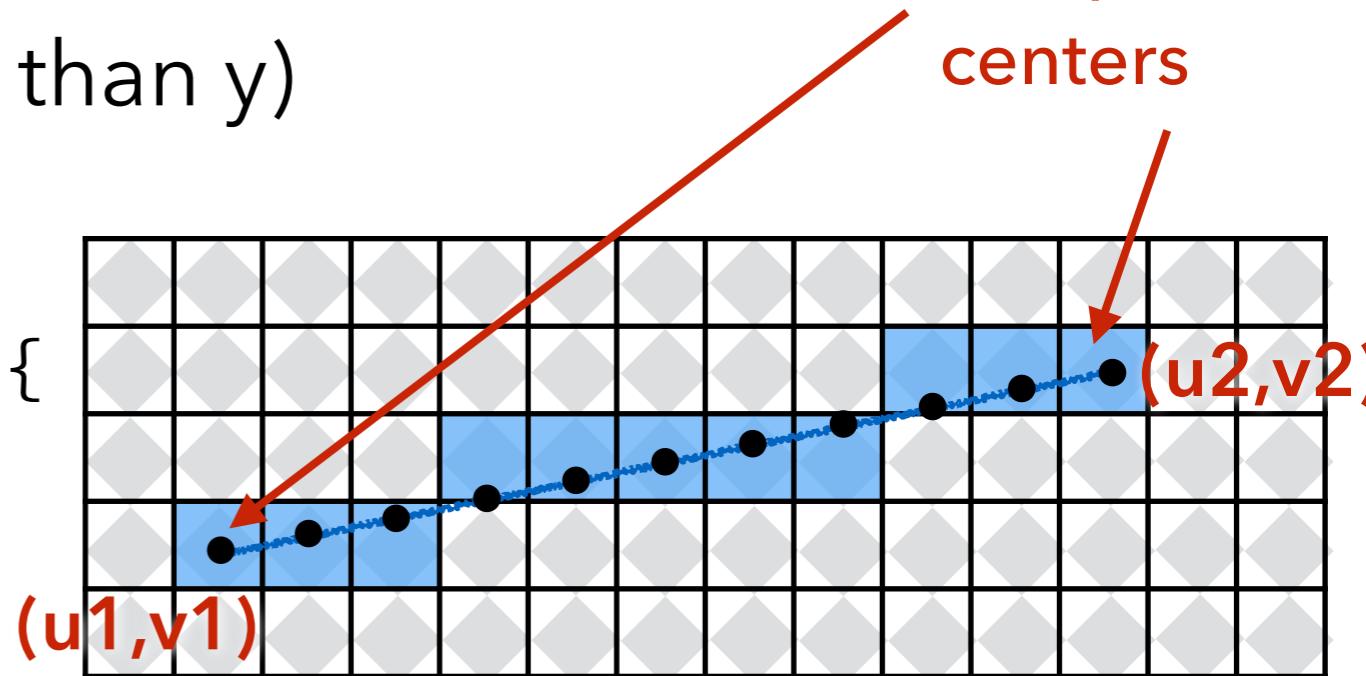
The line's slope is  $s = (v_2 - v_1)/(u_2 - u_1)$

Consider the case where:

- $u_1 < u_2, v_1 < v_2$  (line points toward upper-right)
- $0 < s < 1$  (more change in x than y)

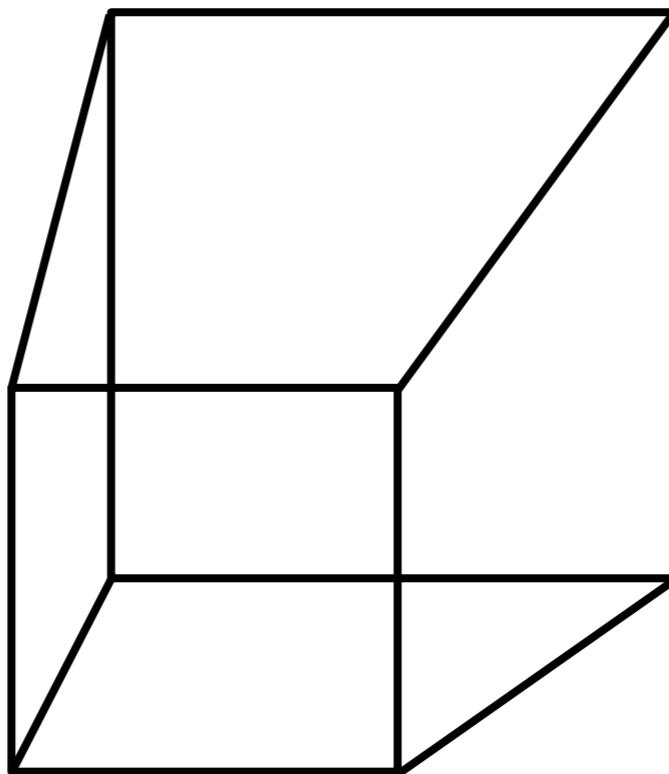
Assume integer coordinates are at pixel centers

```
v = v1;  
for (u = u1; u <= u2; ++u) {  
    v += s;  
    draw(u, round(v))  
}
```



# But wait...

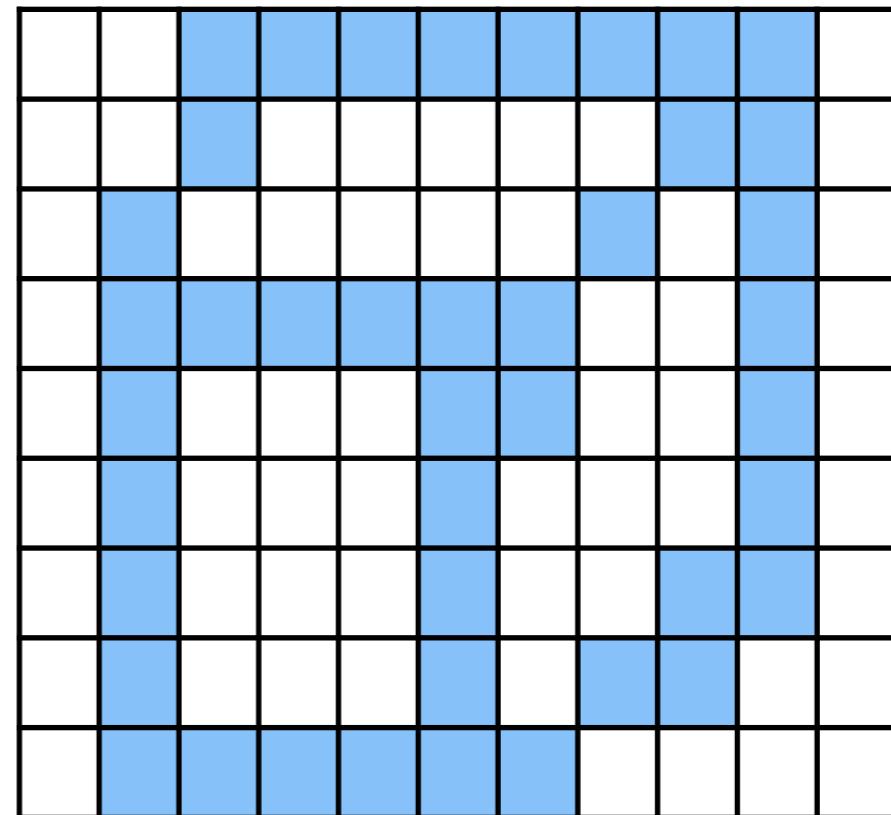
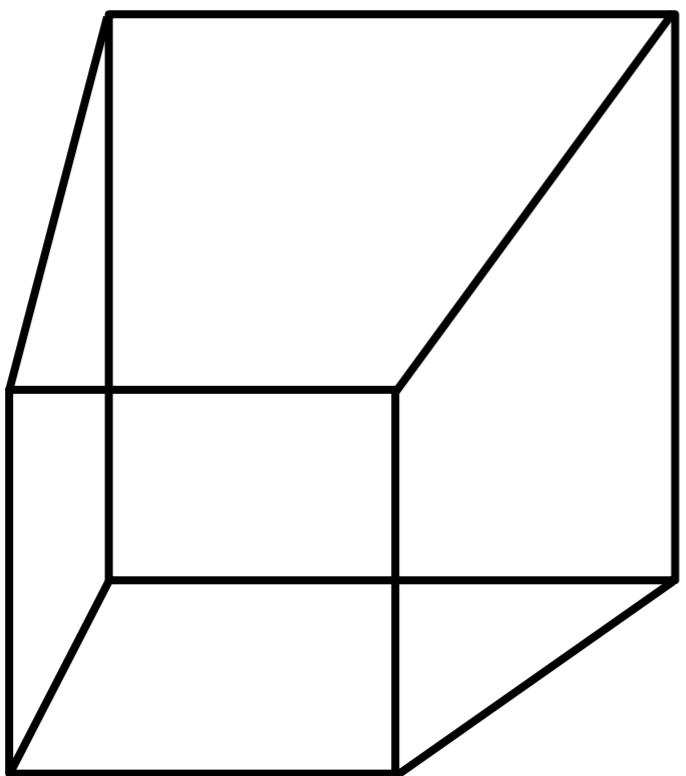
## How do we actually draw lines on a computer?



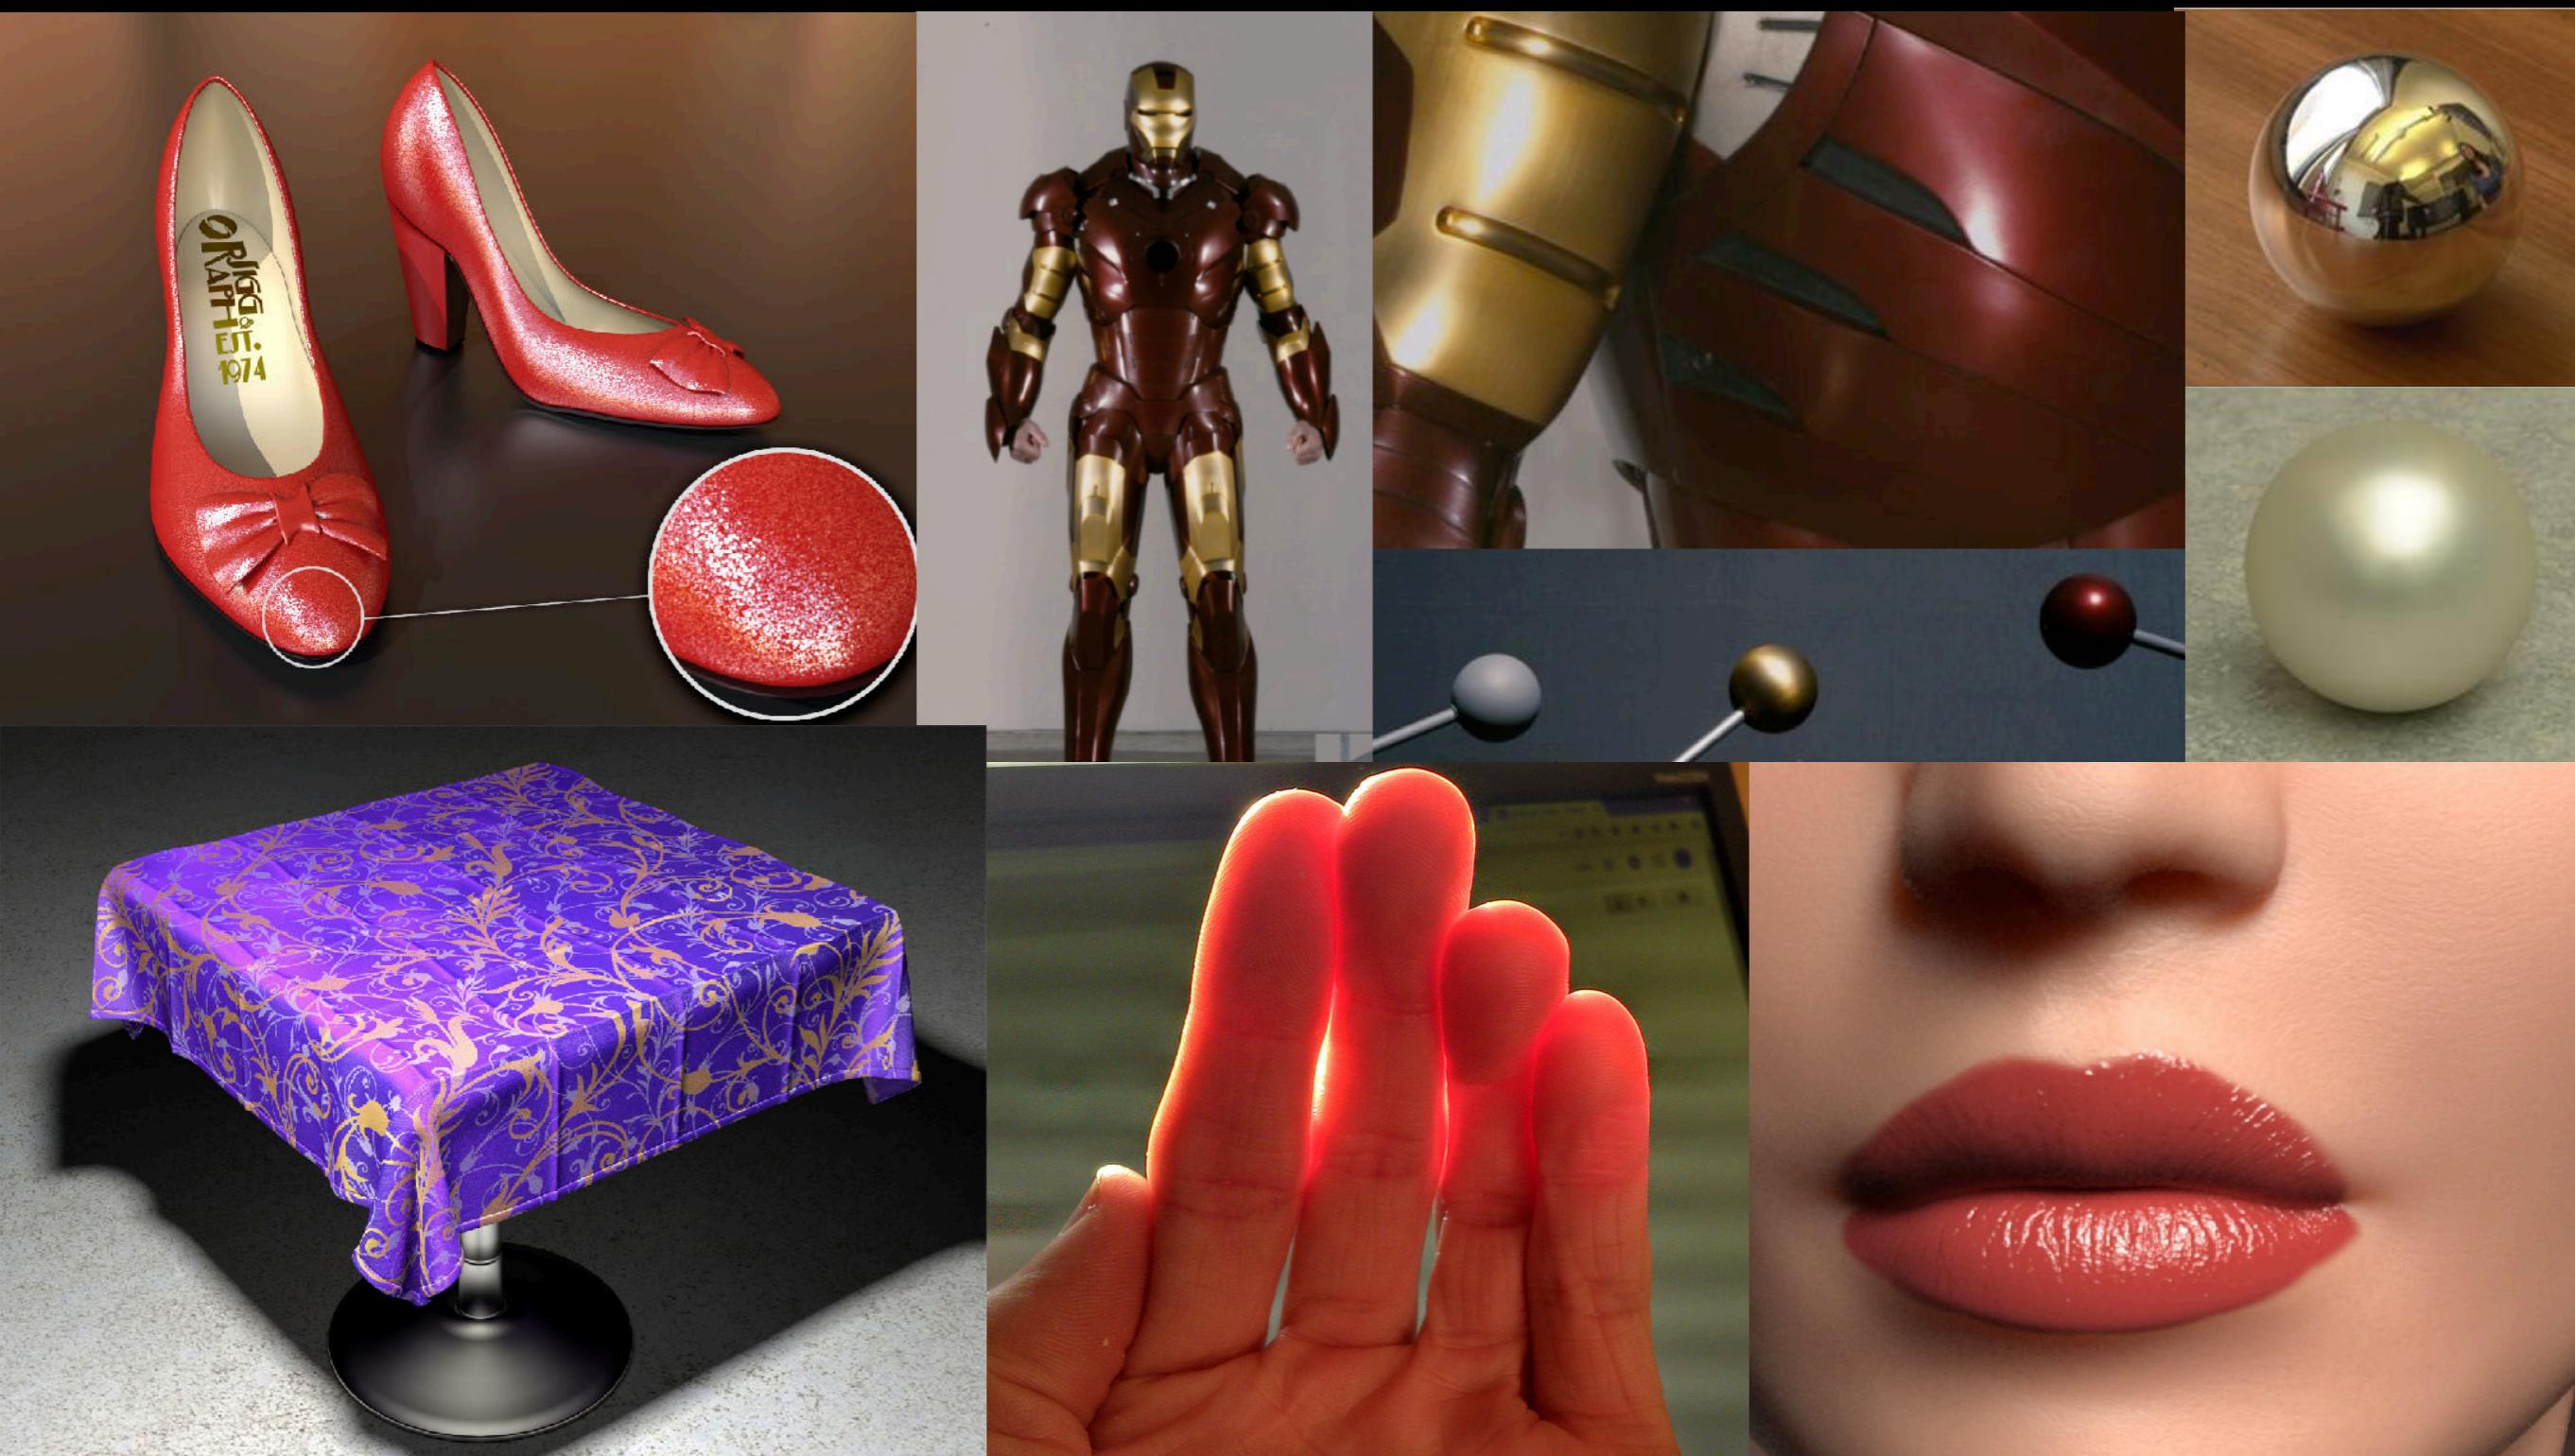
# But wait...

~~How do we actually draw lines on a computer?~~  
What is this object/picture missing?

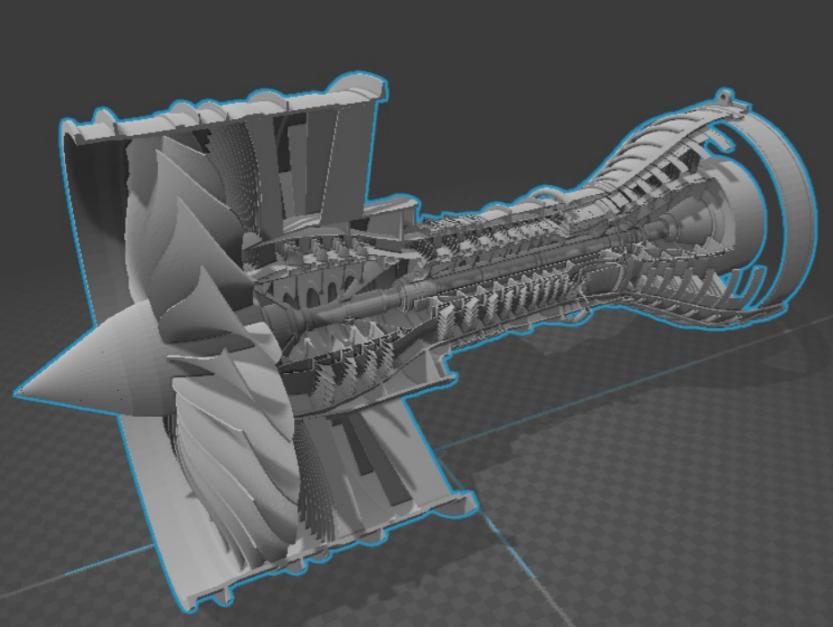
- surface, material, lighting/shading, motion...



# Modelling Material Properties



# Complex 3D Surfaces & Volumes



# Realistic Lighting Environments



WALL·E (Pixar, 2008)

# Realistic Lighting Environments



Industrial Light + Magic

# Realistic Lighting Environments

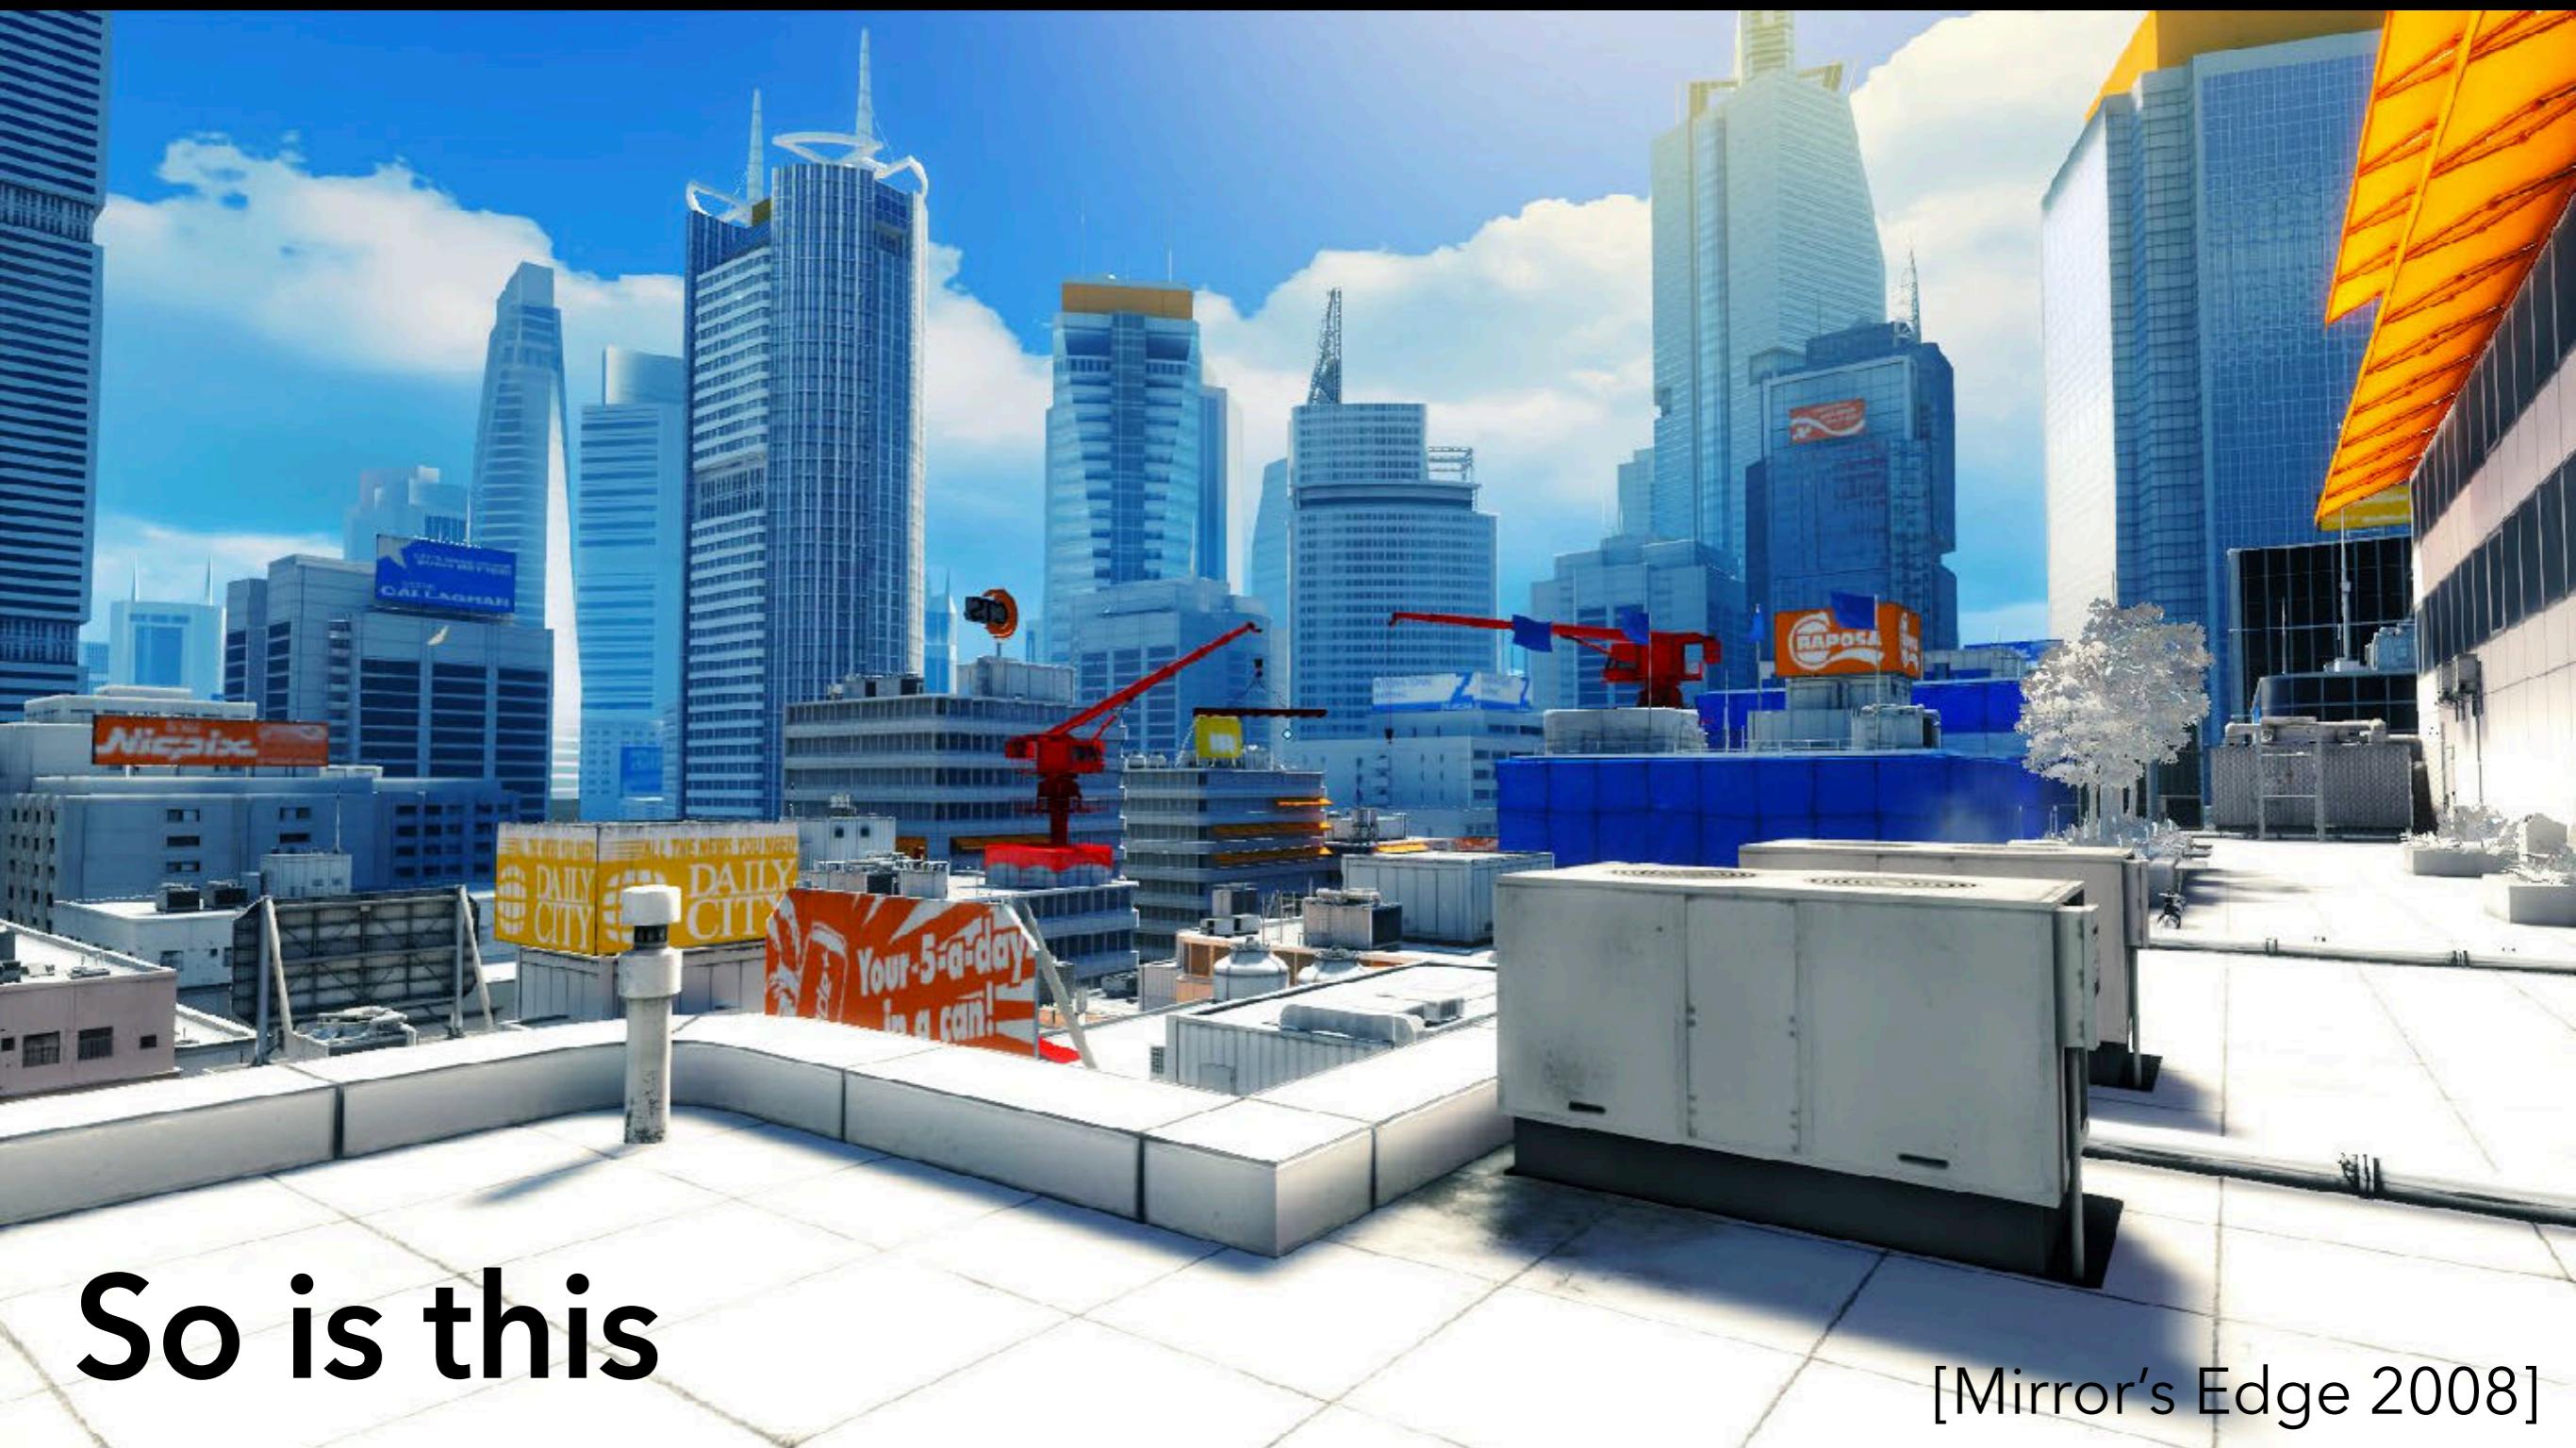


Big Hero 6 (Disney 2014)

**This image is rendered in realtime on a modern GPU**



Unreal Engine Kite Demo (Epic Games 2015)



So is this

[Mirror's Edge 2008]

# Rasterization Pipeline

The modern (programmable) rasterization **pipeline** allows us to scale simple line drawings to these complex examples

To do so, it works in “stages”

- turns out that **triangles** subsume lines as the baseline primitive in a rasterizer
- moving forward, you can assume that every rasterizer acts on a set of triangles

# Rasterization Pipeline – GPUs

You'll be exposed to some of these gritty API details in the tutorial sessions

- we abstract a good deal of this from you\*

At a high-level, the GPU rasterizer and these APIs allow you to:

- take triangles,
- figure out which pixels they cover, and
- perform some computation on these pixels

# these APIs allow you to:

- take triangles,
- figure out which **pixels** they cover, and
- perform some computation on these pixels

```
for(each triangle)
    for(each pixel)
        if(triangle covers pixel)
            perform some computation
```

We're not actually going to test all pixels for each triangle; that would be impractical!

# Programmable Rasterization Pipeline™

After a slide by Steve Marschner

