


ECSE 446/546

# IMAGE SYNTHESIS

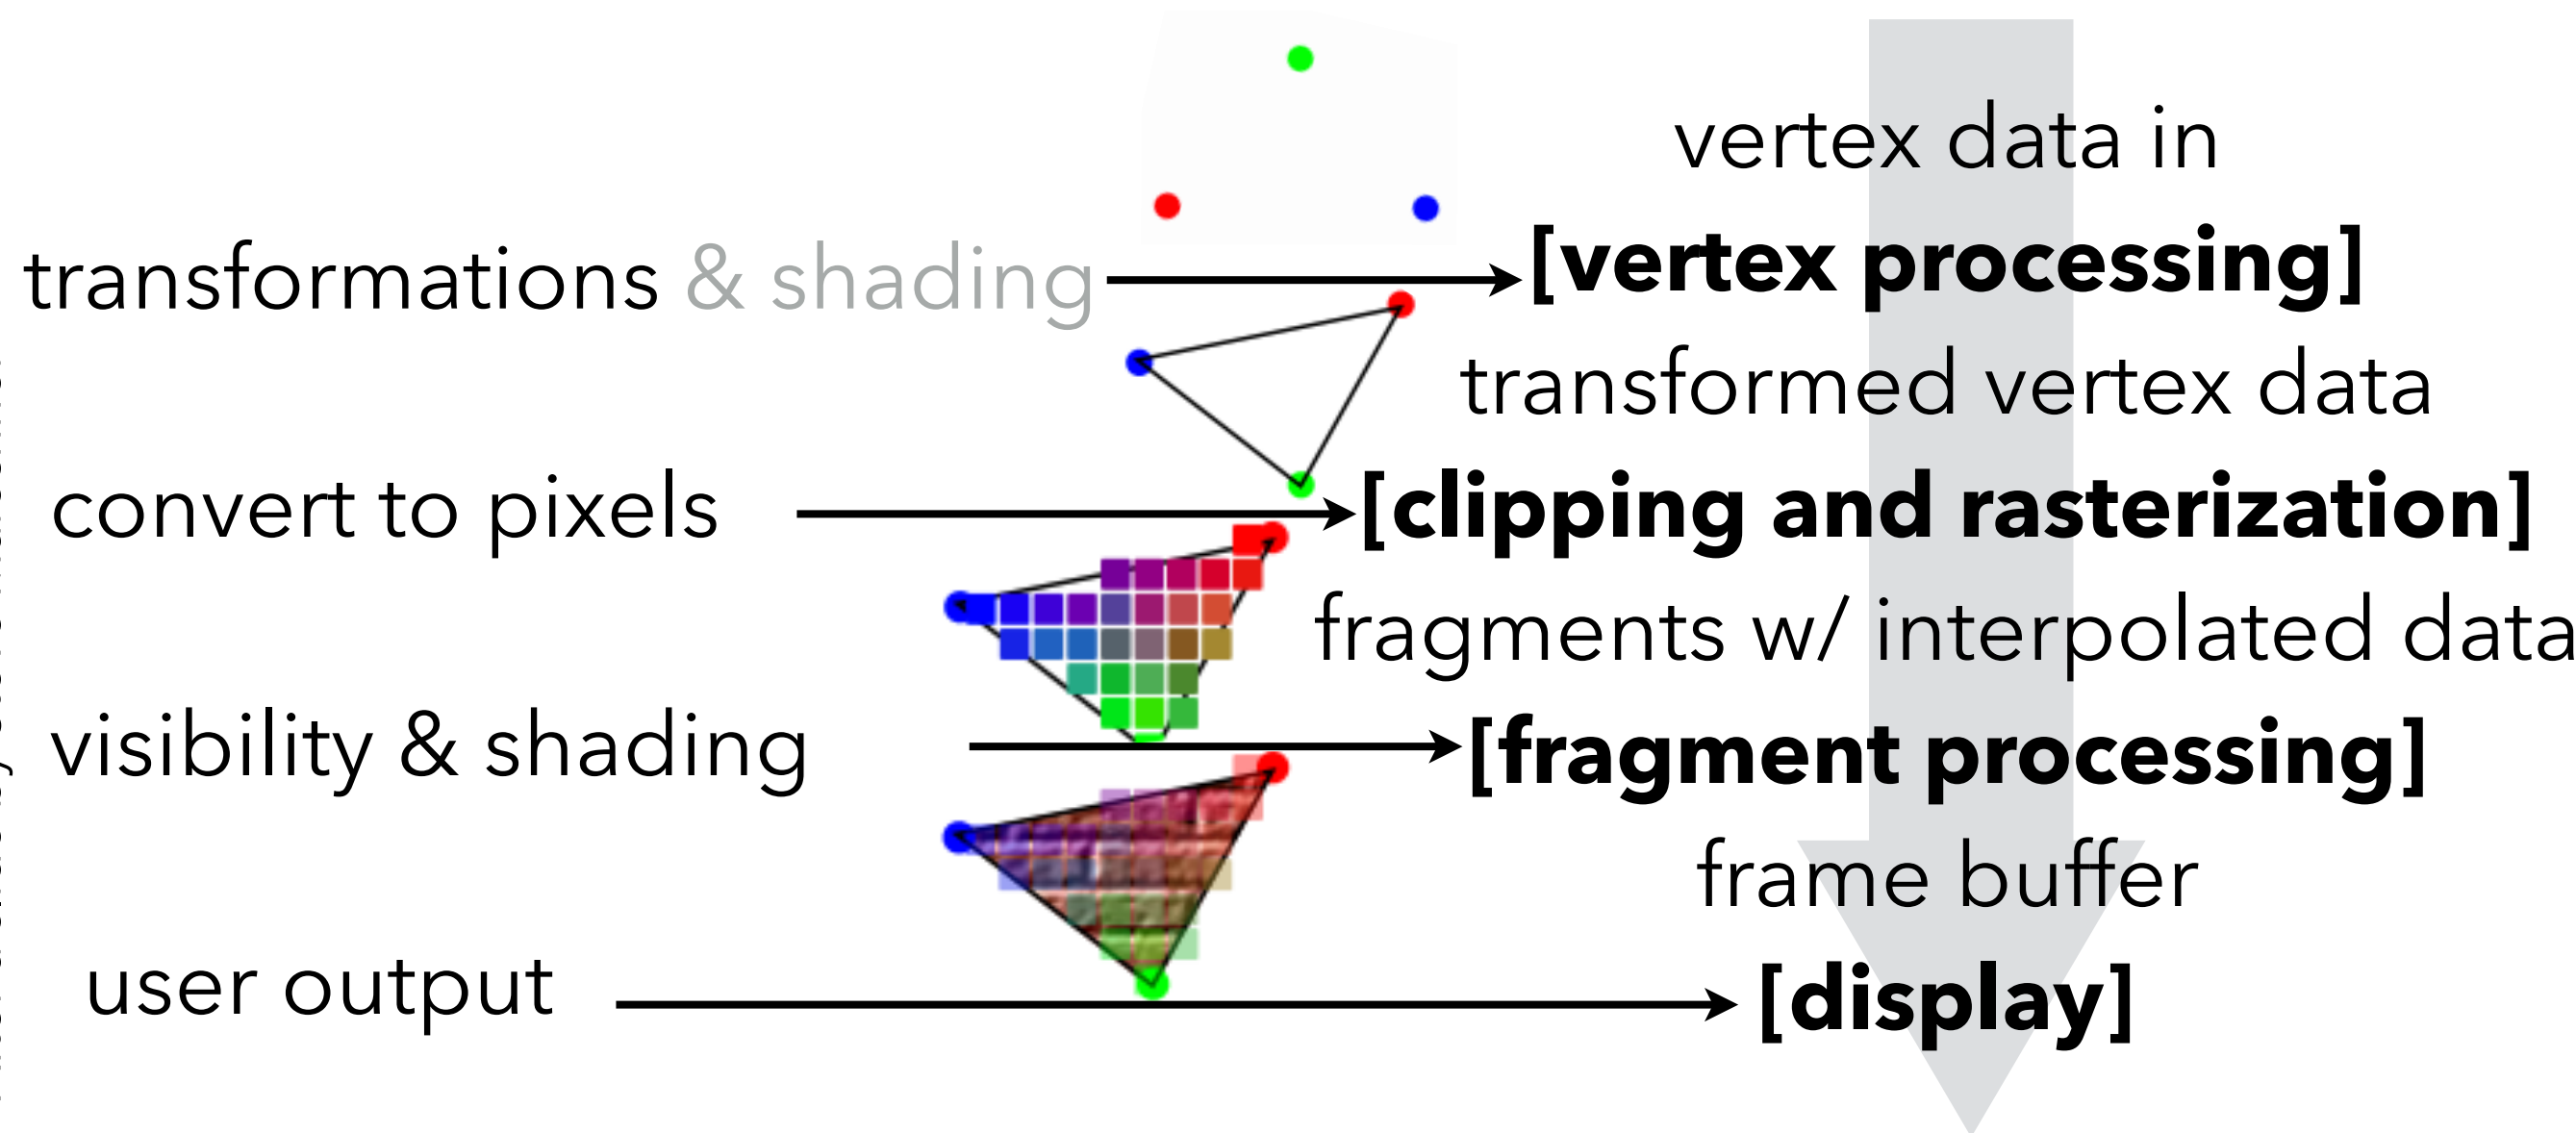


## SYSTEMS 1 – RASTERIZATION

Prof. Derek Nowrouzezahrai

[derek@cim.mcgill.ca](mailto:derek@cim.mcgill.ca)


# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

# Programmable Rasterization Pipeline™

vertex data in  
**[vertex processing]**  
transformed vertex data  
**[clipping and rasterization]**  
fragments w/ interpolated data  
**[fragment processing]**  
frame buffer  
**[display]**



```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
```

```
    if(x,y in triangle)
```


# Programmable Rasterization Pipeline™

```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```

# Programmable Rasterization Pipeline™

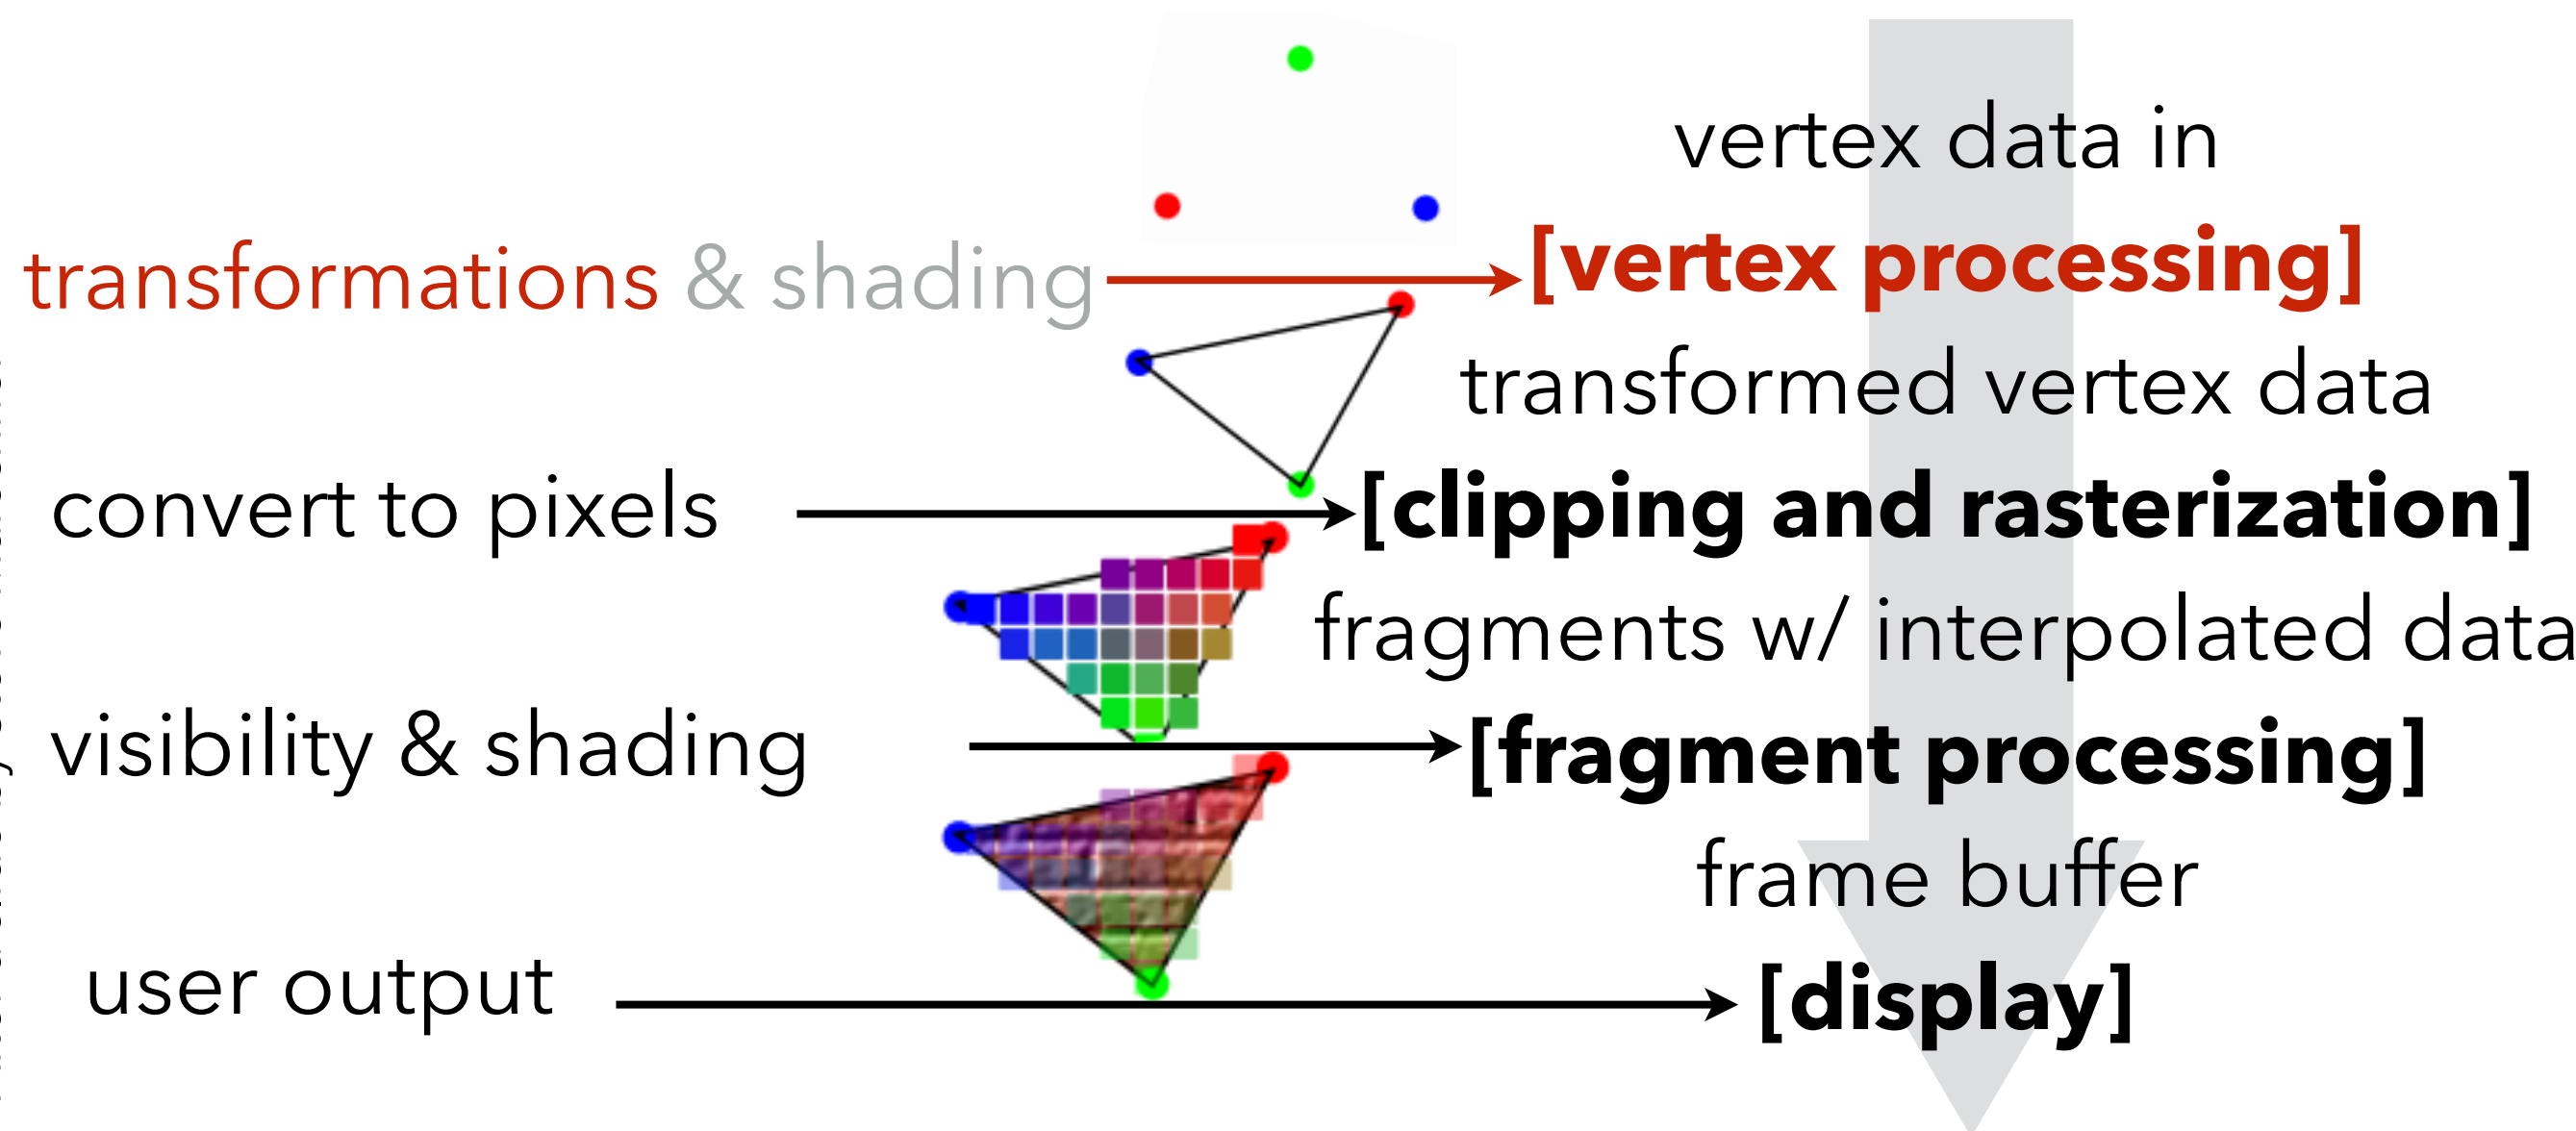
```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```

# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

# Programmable Rasterization Pipeline™

```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```



# Programmable Rasterization Pipeline™

```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```



# Model/View/Projection (MVP) Transformations

# Programmable Rasterization Pipeline™

```
for(each triangle)
    transform vertices into eye space
    project vertices to image space
    for(each pixel x,y)
        if(x,y in triangle)
            compute z
            if(z < zbuffer[x,y])
                zbuffer[x,y] = z
                framebuffer[x,y] = shade()
```


# Rasterization

---

**Goal:** convert from object point to image plane

- start with a 3D object point
- apply a sequence of transforms
  - transforms can be specified by a matrix
- determine the 2D image plane point it projects to

# Pipeline of transformations



## 1. Model

map local  
object  
coords to  
world  
coords

## 2. Viewing

map world  
coords to  
camera  
coords

## 3. Projection




map camera  
coords to  
canonical view  
volume

## 4. Viewport

map canonical  
view volume to  
screen space

These two stages perform the  
actual 3D-to-2D projection

## 2. Viewing transform



The view matrix transforms all coordinates into "eye space"

## 2. Viewing transform

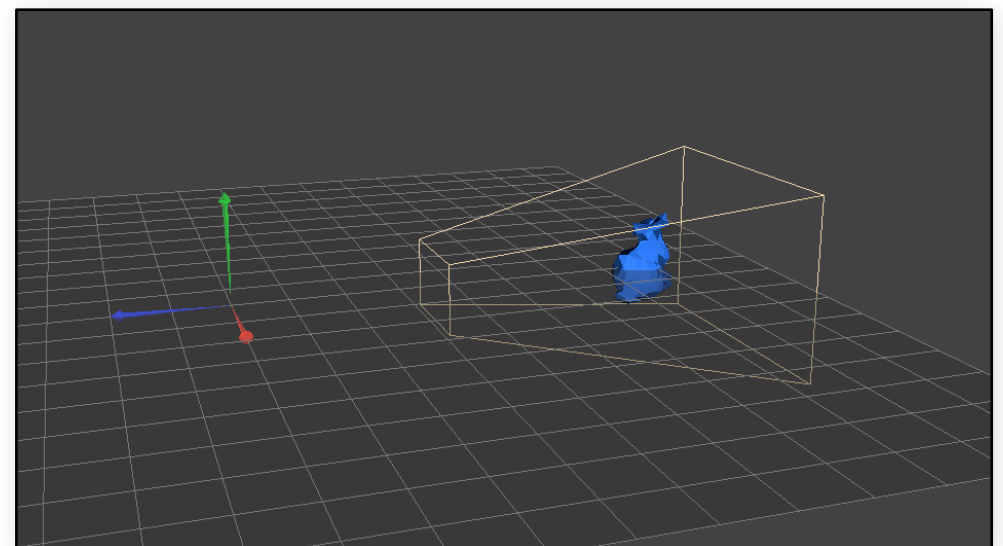


---

Many ways to construct a view transform (i.e., matrix)

For example, can specify an orthonormalized frame with the:

- **eye** point
- **up**-vector, and
- **look-at** point

# 2. Viewing transform





# Programmable Rasterization Pipeline™


```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```

# 3. Projection

Generally, a function that transforms points from  $m$ - to  $n$ -space where  $m > n$

In graphics, map 3D points to 2D image coordinates

- except we will keep around the third coordinate



# Mathematics of Projection

---

Always work in eye coords


- assume eye point at **0** and plane perpendicular to  $z$

Orthographic case: just toss out the  $z$ -coord

Perspective case: scale diminishes with  $z$

# Orthographic projection


## Box view volume



After a slide by Fabio Pellacini


# Orthographic projection

Viewing rays are parallel




# Perspective projection

## Truncated pyramid view volume




# Perspective projection

Viewing rays converge to a point



# Perspective projection

Use similar triangles:




$$\frac{y'}{-1} = \frac{y}{z}$$
$$y' = -\frac{y}{z}$$




# Remapping the view frustum

We also want the projection matrix to remap the space between **near** and **far** planes to a canonical view volume



# Pipeline of transformations



1. Model  
map local  
object  
coords to  
world  
coords


2. Viewing  
map world  
coords to  
camera  
coords

3. Projection  
map camera  
coords to  
canonical view  
volume

4. **Viewport**  
map canonical  
view volume to  
screen space

These two stages perform the  
actual 3D-to-2D projection

# 4. Viewport transformation



$$\begin{bmatrix} x_{\text{screen}} \\ y_{\text{screen}} \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{n_x}{2} & 0 & \frac{n_x - 1}{2} \\ 0 & \frac{n_y}{2} & \frac{n_y - 1}{2} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_{\text{canonical}} \\ y_{\text{canonical}} \\ 1 \end{bmatrix}$$

# Orthographic Transformation Chain

Start with coordinates in object's local coordinates

Transform into world coords (modeling transform,  $\mathbf{M}_m$ )

Transform into eye coords (camera xf.,  $\mathbf{M}_{cam} = \mathbf{F}_c^{-1}$ )

Orthographic projection,  $\mathbf{M}_{orth}$

Viewport transform,  $\mathbf{M}_{vp}$

$$\mathbf{p}_s = \mathbf{M}_{vp} \mathbf{M}_{orth} \mathbf{M}_{cam} \mathbf{M}_m \mathbf{p}_o$$

$$\begin{bmatrix} x_s \\ y_s \\ z_c \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{n_x}{2} & 0 & 0 & \frac{n_x-1}{2} \\ 0 & \frac{n_y}{2} & 0 & \frac{n_y-1}{2} \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{2}{n-f} & -\frac{n+f}{n-f} \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \mathbf{u} & \mathbf{v} & \mathbf{w} & \mathbf{e} \\ 0 & 0 & 0 & 1 \end{bmatrix}^{-1} \mathbf{M}_m \begin{bmatrix} x_o \\ y_o \\ z_o \\ 1 \end{bmatrix}$$

# Perspective Transformation Chain

Start with coordinates in object's local coordinates

Transform into world coords (modelling transform,  $\mathbf{M}_m$ )

Transform into eye coords (camera xf.,  $\mathbf{M}_{cam} = \mathbf{F}_c^{-1}$ )

Perspective projection,  $\mathbf{M}_{persp}$

Viewport transform,  $\mathbf{M}_{vp}$

$$\mathbf{p}_s = \mathbf{M}_{vp} \mathbf{M}_{persp} \mathbf{M}_{cam} \mathbf{M}_m \mathbf{p}_o$$

$$\begin{bmatrix} x_s \\ y_s \\ z_c \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{n_x}{2} & 0 & 0 & \frac{n_x-1}{2} \\ 0 & \frac{n_y}{2} & 0 & \frac{n_y-1}{2} \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \frac{2n}{r-l} & 0 & \frac{l+r}{l-r} & 0 \\ 0 & \frac{2n}{t-b} & \frac{b+t}{b-t} & 0 \\ 0 & 0 & \frac{f+n}{n-f} & \frac{2fn}{f-n} \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} \mathbf{u} & \mathbf{v} & \mathbf{w} & \mathbf{e} \end{bmatrix}^{-1} \mathbf{M}_m \begin{bmatrix} x_o \\ y_o \\ z_o \\ 1 \end{bmatrix}$$

# Recap: Transformation Pipeline

Perform rotation/translation/other transforms to put viewpoint at origin and view direction along z axis

- combination of model and view matrix called "modelview" matrix (e.g., in OpenGL)

Combine with projection matrix (perspective or orthographic)

- this is the OpenGL "projection" matrix

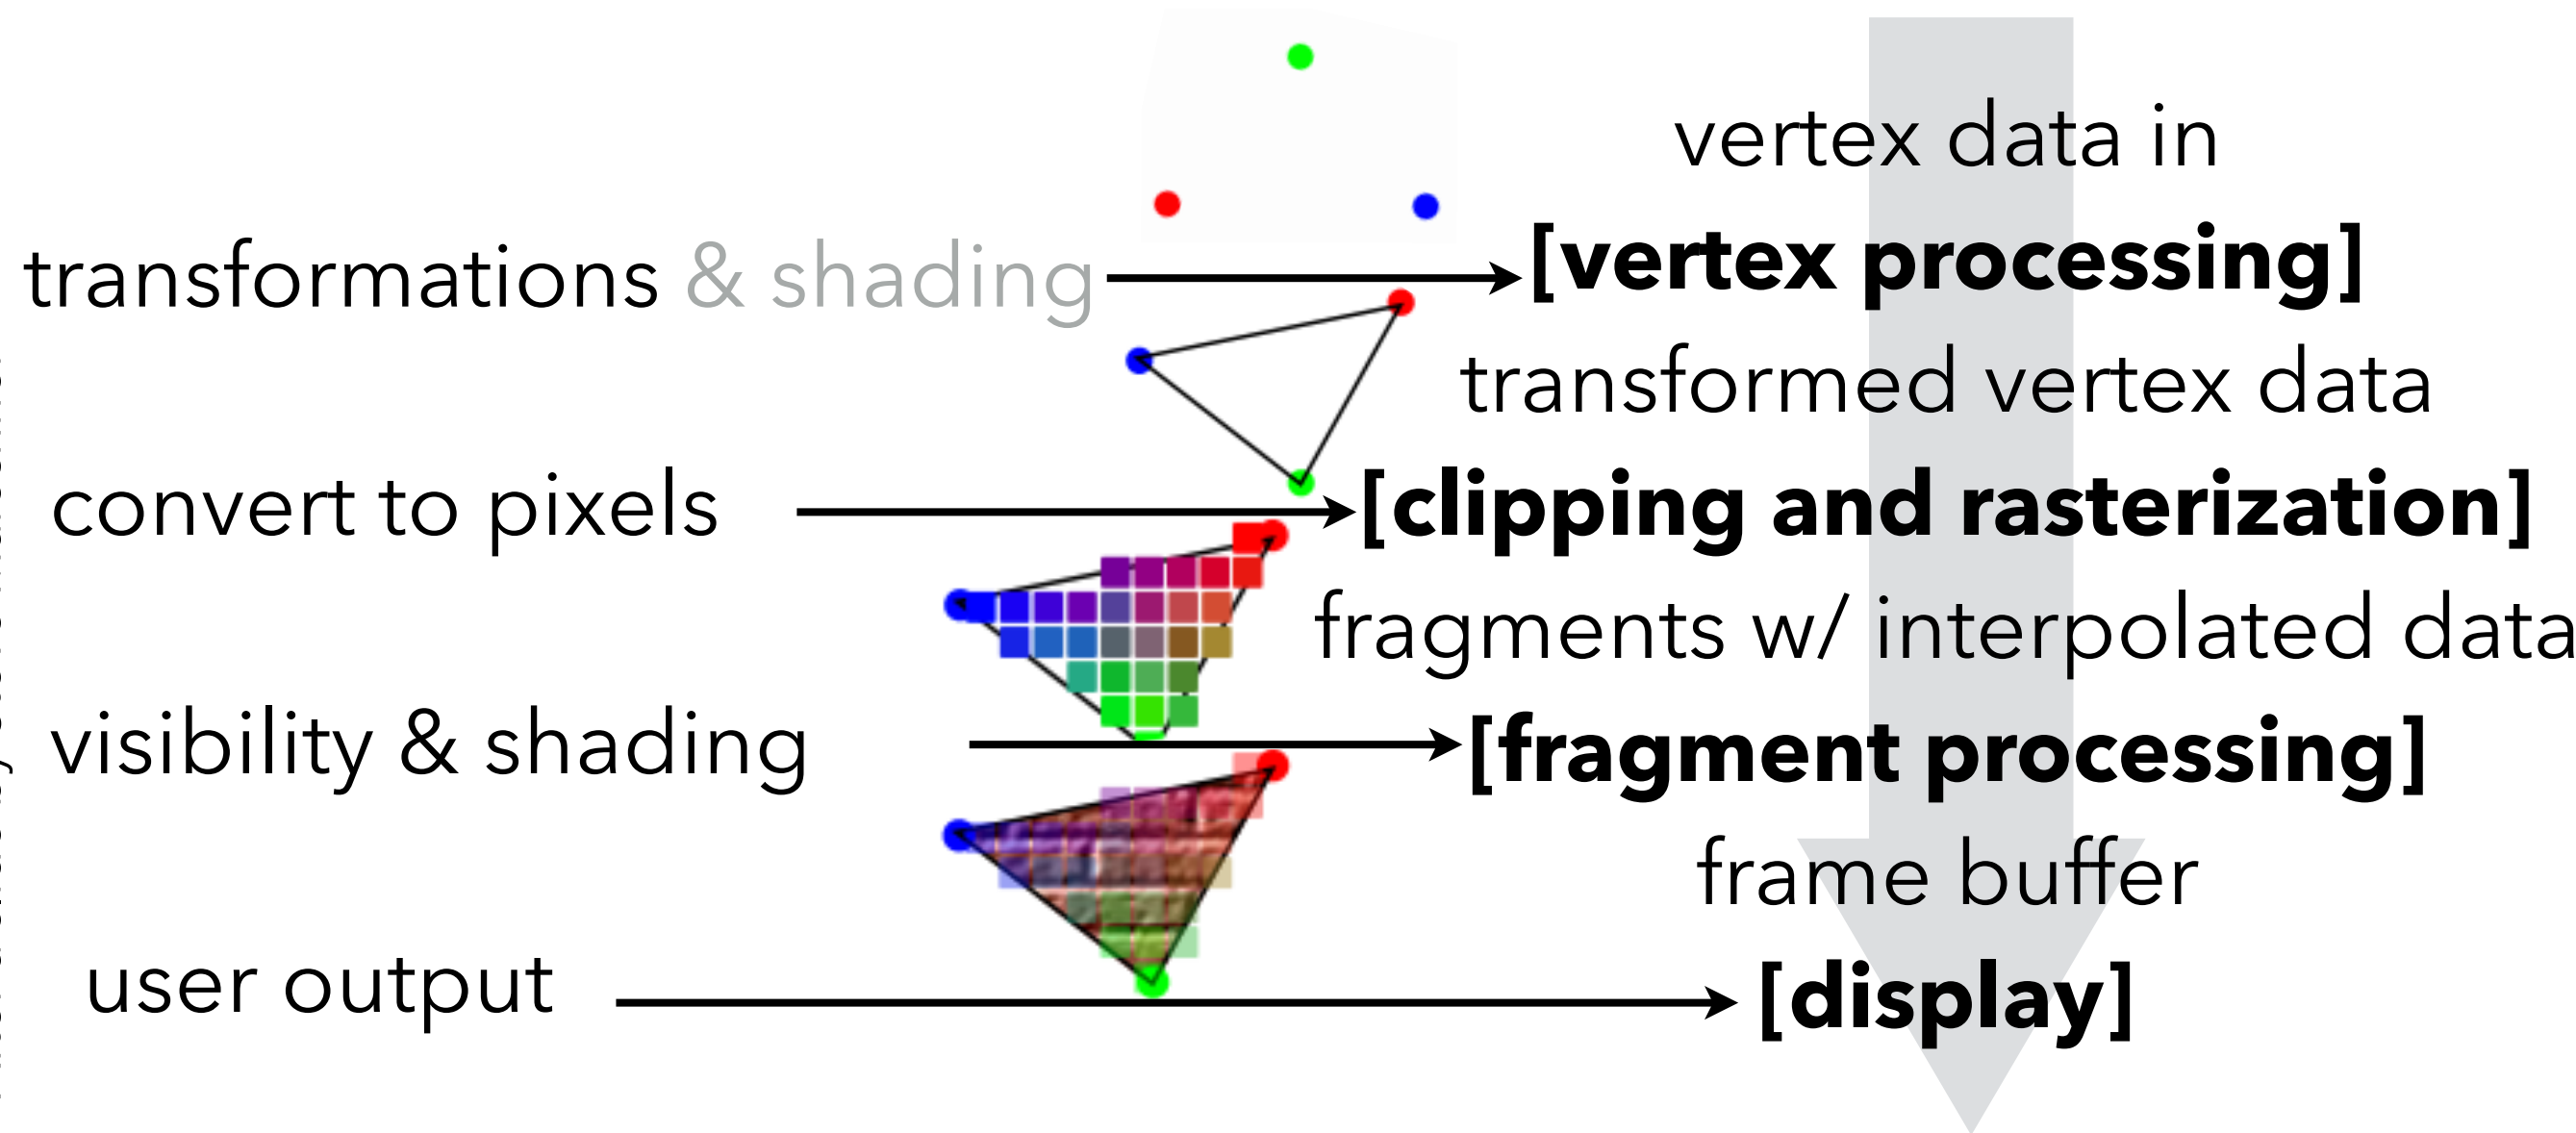
Convert canonical view volume  $[-1, 1]^3$  to screen space

**Corollary:** The entire transformation from local object space to screen space is a single 4x4 matrix!

# Questions?

---

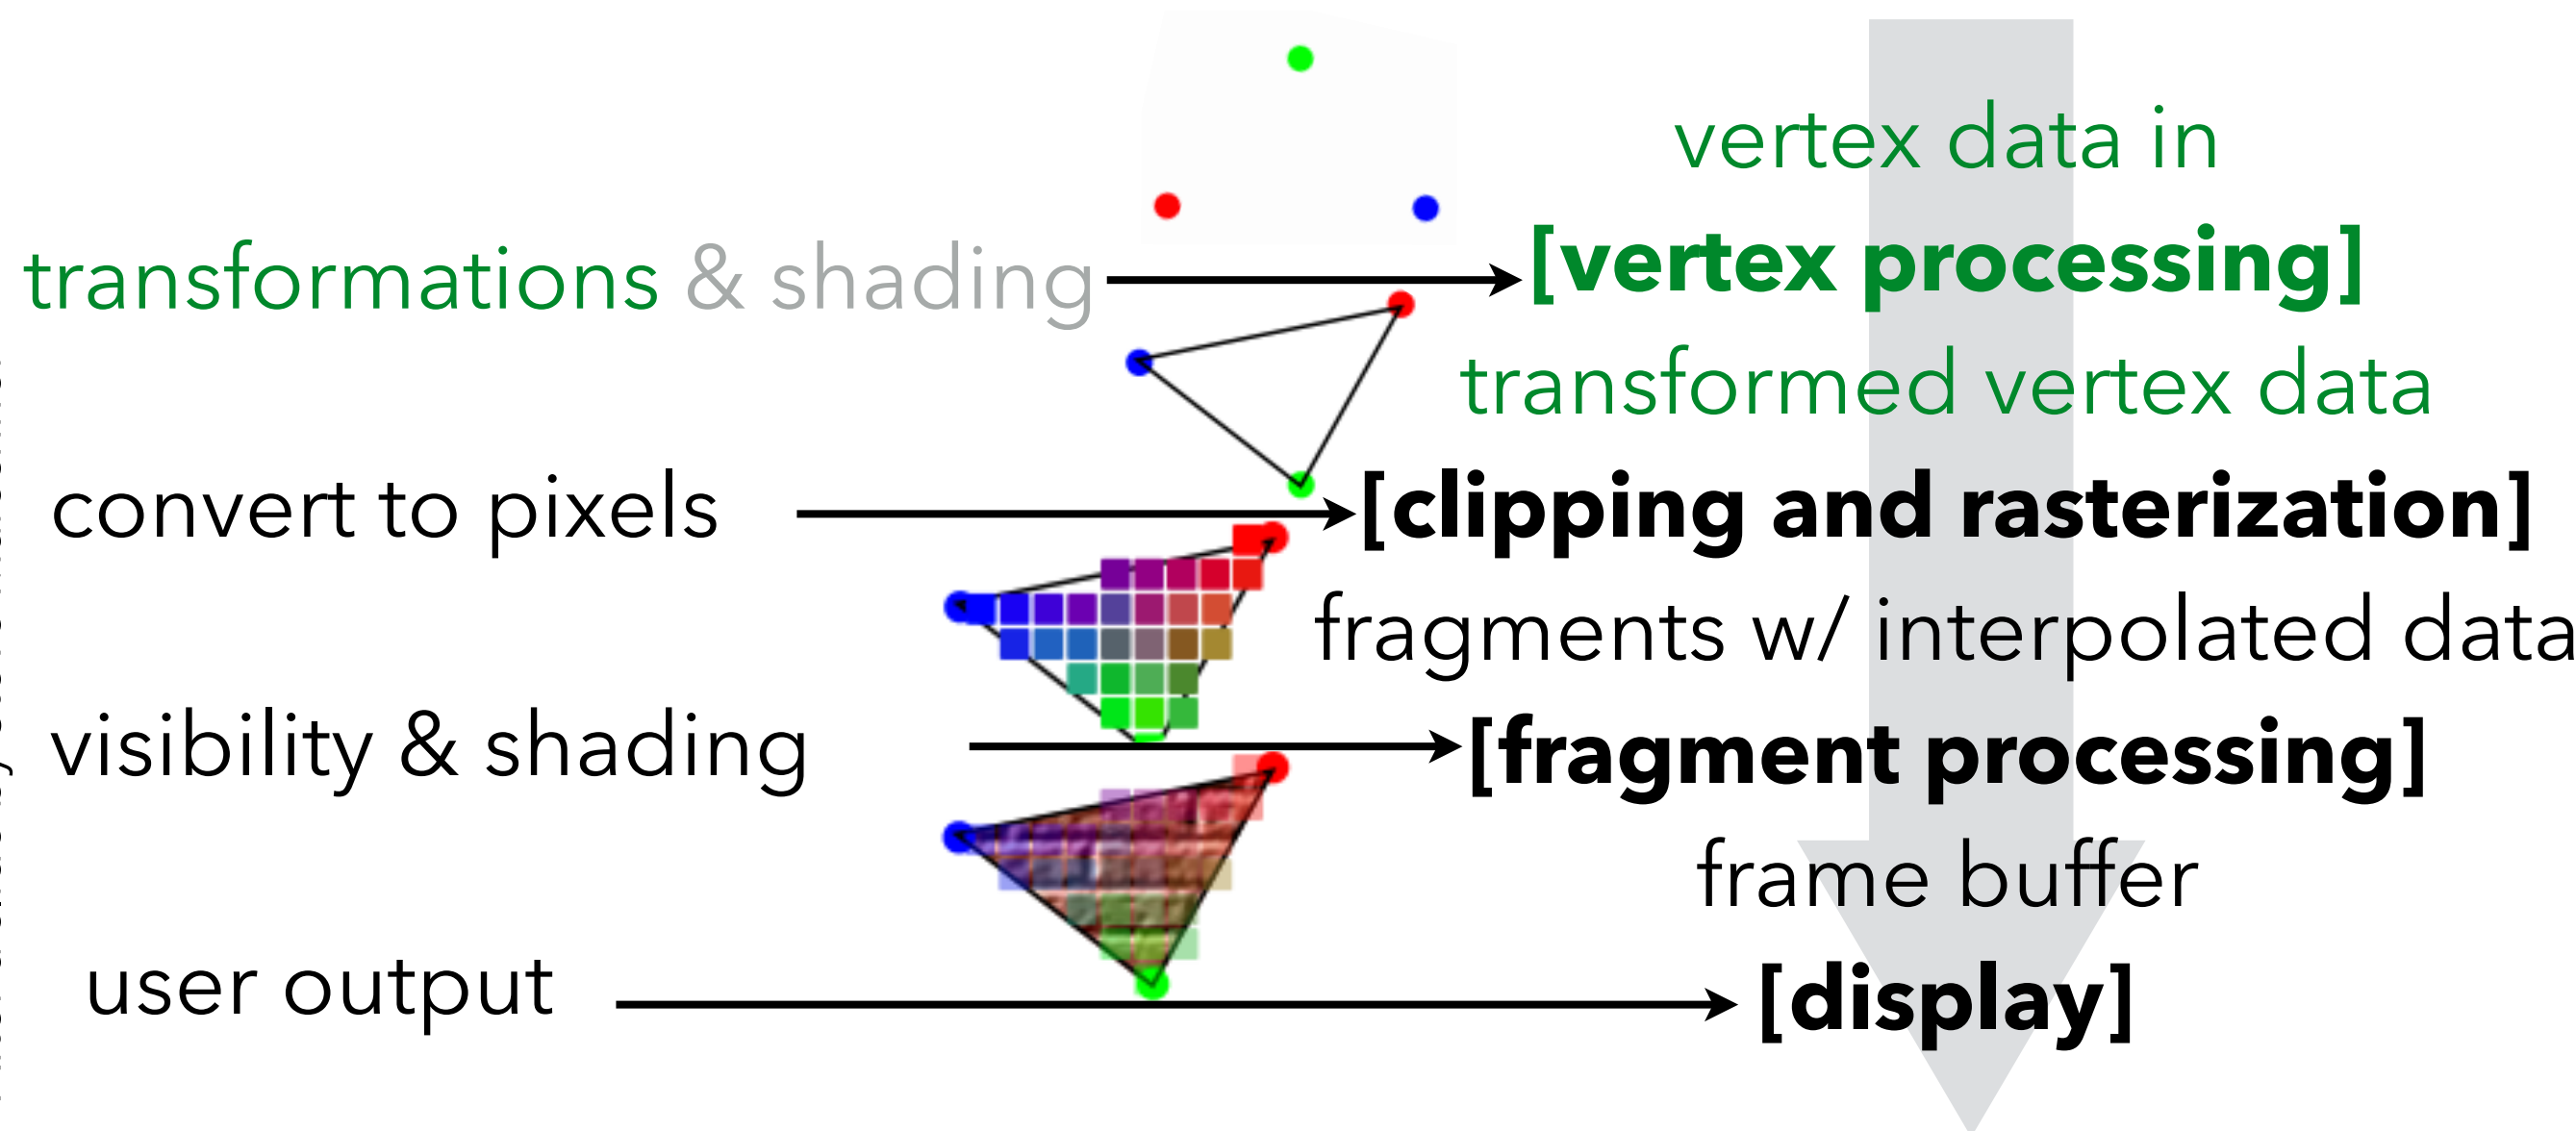
# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

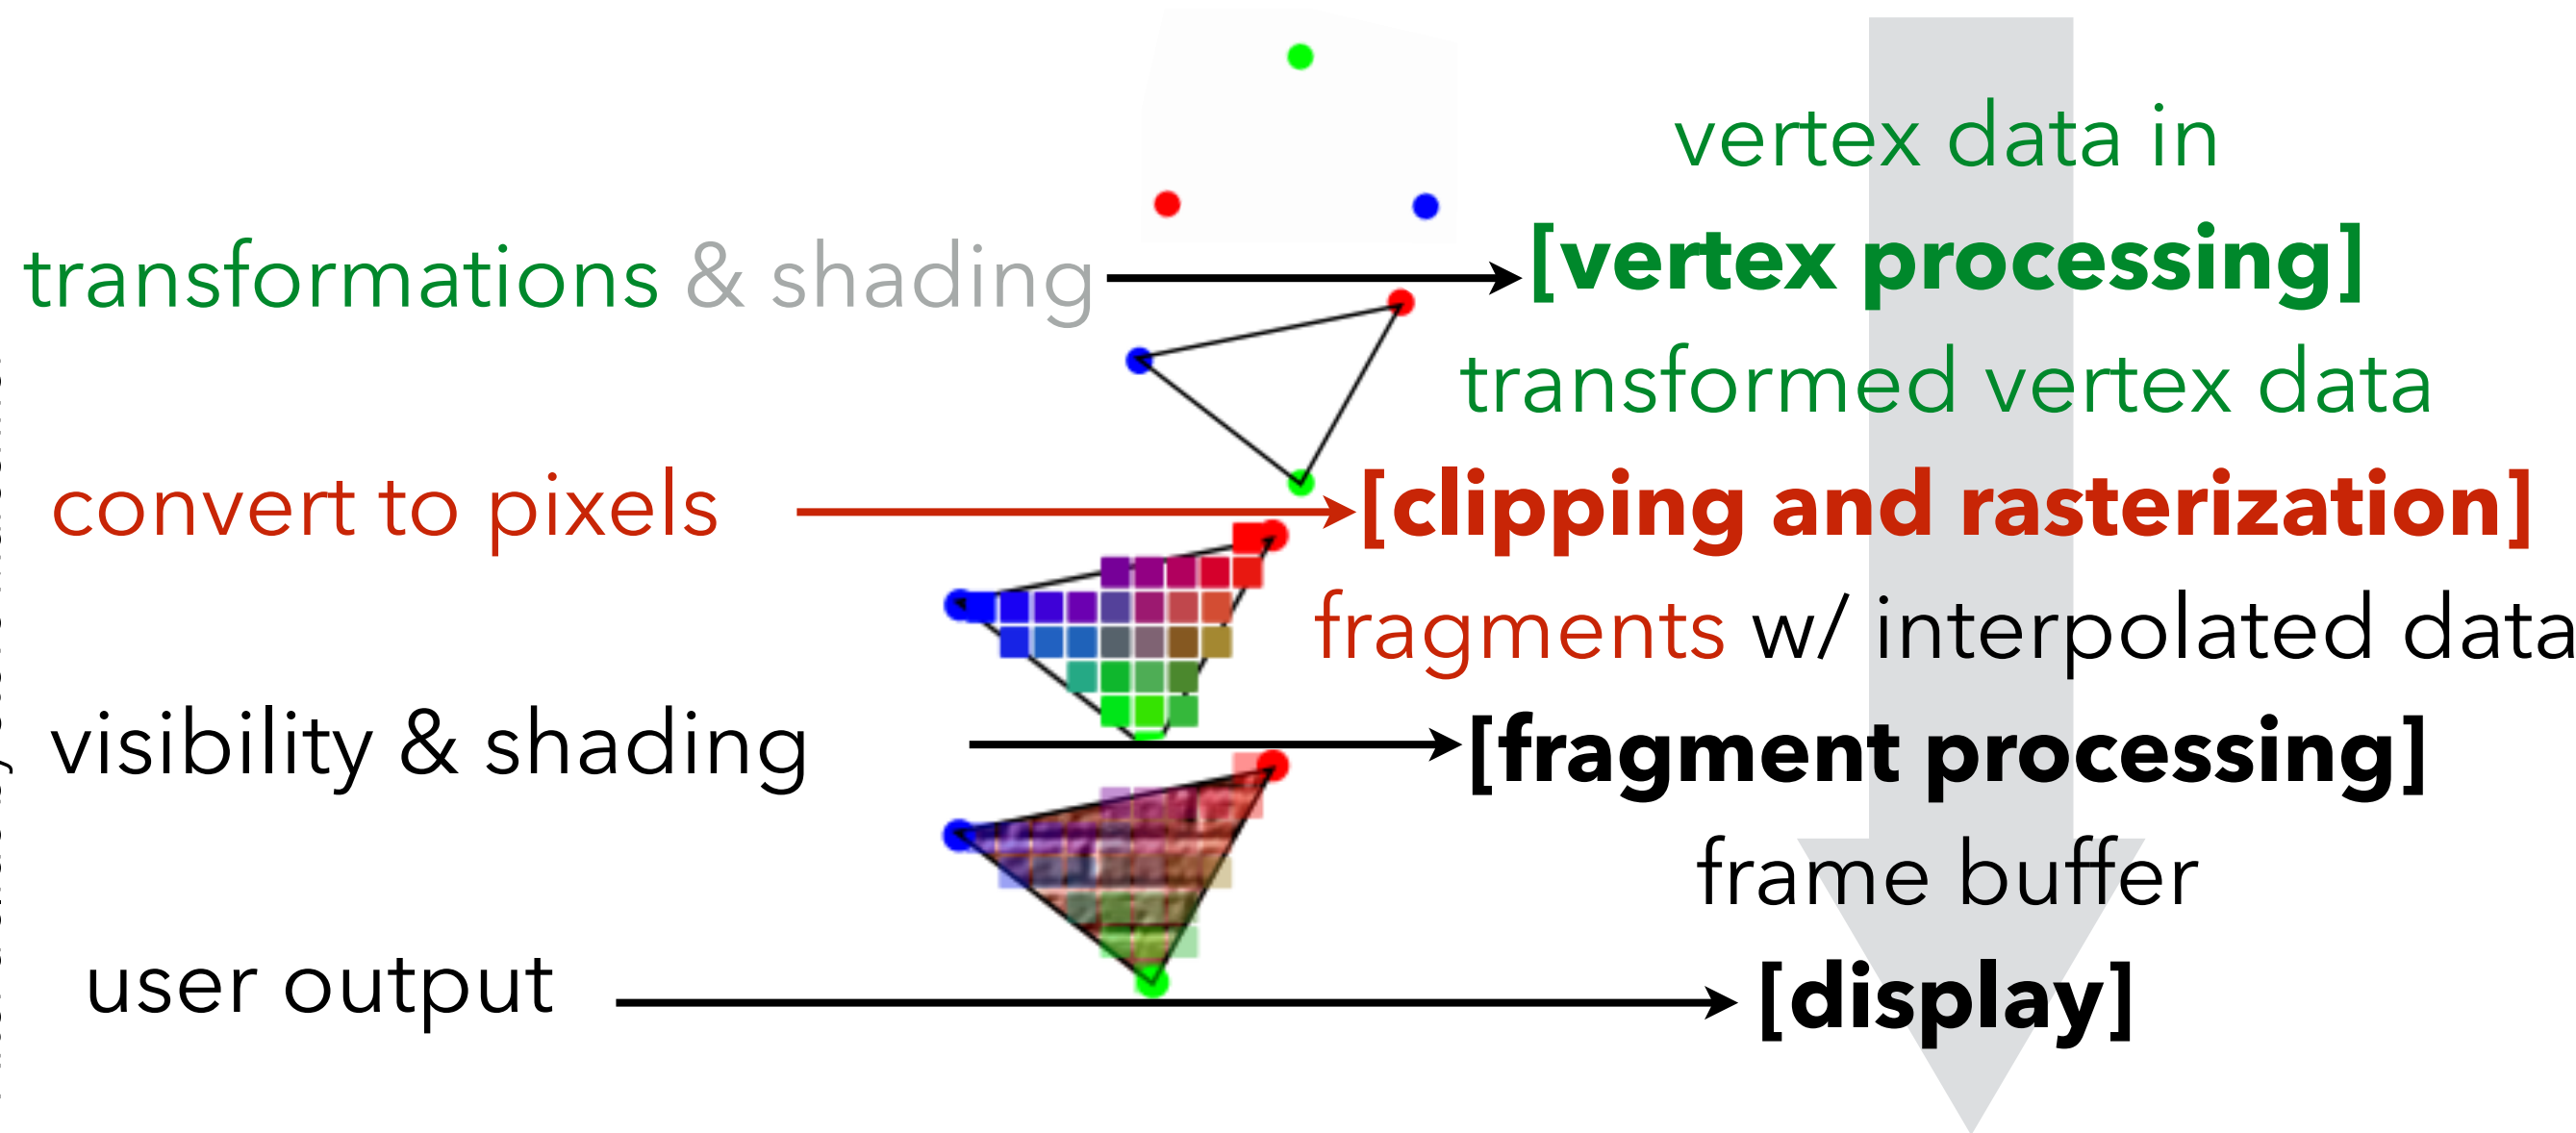


# Programmable Rasterization Pipeline™



After a slide by Steve Marschner


# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

# Programmable Rasterization Pipeline™


```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```



# Triangle Rasterization

# Triangle Rasterization


Primitives are "continuous" geometric objects;  
screen is discrete (pixels)



# Rasterization

Primitives are "continuous" geometric objects;  
screen is discrete (pixels)


Rasterization computes a discrete  
approximation in terms of pixels



# Edge Functions

A triangle's 3D edges project to line segments in the image (thanks to planar perspective)

- lines map to lines, not curves




# Edge Functions

A triangle's 3D edges project to line segments in the image (thanks to planar perspective)

- interior of the triangle is the set of points lie inside **all** 3 half-spaces defined by these lines

$$E_i(x, y) = a_i x + b_i y + c_i$$

$(x, y)$  within triangle iff  
 $E_i(x, y) \geq 0, \forall i = 1, 2, 3$






# Brute Force Triangle Rasterizer

Compute  $E_1, E_2, E_3$  coefficients from projected vertices


- called "triangle setup": yields  $a_i, b_i, c_i$  for  $i = 1, 2, 3$



# Brute Force Triangle Rasterizer

for each pixel  $(x,y)$   
evaluate edge functions at pixel center  
if all non-negative, pixel is in


**Problem?**



# Brute Force Triangle Rasterizer

for each pixel  $(x,y)$   
    evaluate edge functions at pixel center  
    if all non-negative, pixel is in


If the triangle is small,  
lots of useless  
computation (if we  
really test all pixels)



# Easy Optimization

Improvement: only scan over pixels overlapping the **screen bounding box** (BBox) of the triangle


- how do we get such a bounding box?
  - $x_{\min}$ ,  $x_{\max}$ ,  $y_{\min}$ ,  $y_{\max}$  of projected triangle vertices



# Triangle Rasterization Pseudocode

```
for every triangle
  project vertices, compute the  $E_i$ 
  compute bbox, clip box to screen limits
  for all pixels in box
    evaluate  $E_i$  functions
    if all  $> 0$ 
      framebuffer[x,y] = c;
```

Bounding box clipping  
is easy, just clamp the  
coordinates to the  
screen rectangle



Note: no visibility!

# Triangle Rasterization Pseudocode



**Jaakko Lehtinen**

@jaakkolehtinen

Follow

My most succinct rasterizer yet. Still kind of readable.

```
% triangle vertices in [-1,1]^2
P = [-1 -0.51; -0.3 0.77; -0.1 -0.32];
% pixel samples 0.1 units apart
[px,py] = meshgrid(-1+0.1/2:0.1:1, -1+0.1/2:0.1:1);
% form edge functions
E = [P([2 3 1],2), P([1 2 3],1)] - [P([1 2 3],2), P([2 3 1],1)];
E(:,3) = -diag(E*P');
% evaluate edge functions at sample points
res = E*[px(:)'; py(:)'; ones(1,length(px(:)))];
% which ones are in?
inside = min(res,[],1) > 0;
```

7:32 AM - 13 Nov 2015


7 Retweets 23 Likes



7



23



# Options, Optimizations and Caveats

---

Could use barycentric coordinates instead of edge functions

Can operate on blocks of pixels before going per-pixel


Incremental computation possible, but...

- ... parallelism most important for hardware

Be careful about pixels on triangle edges!

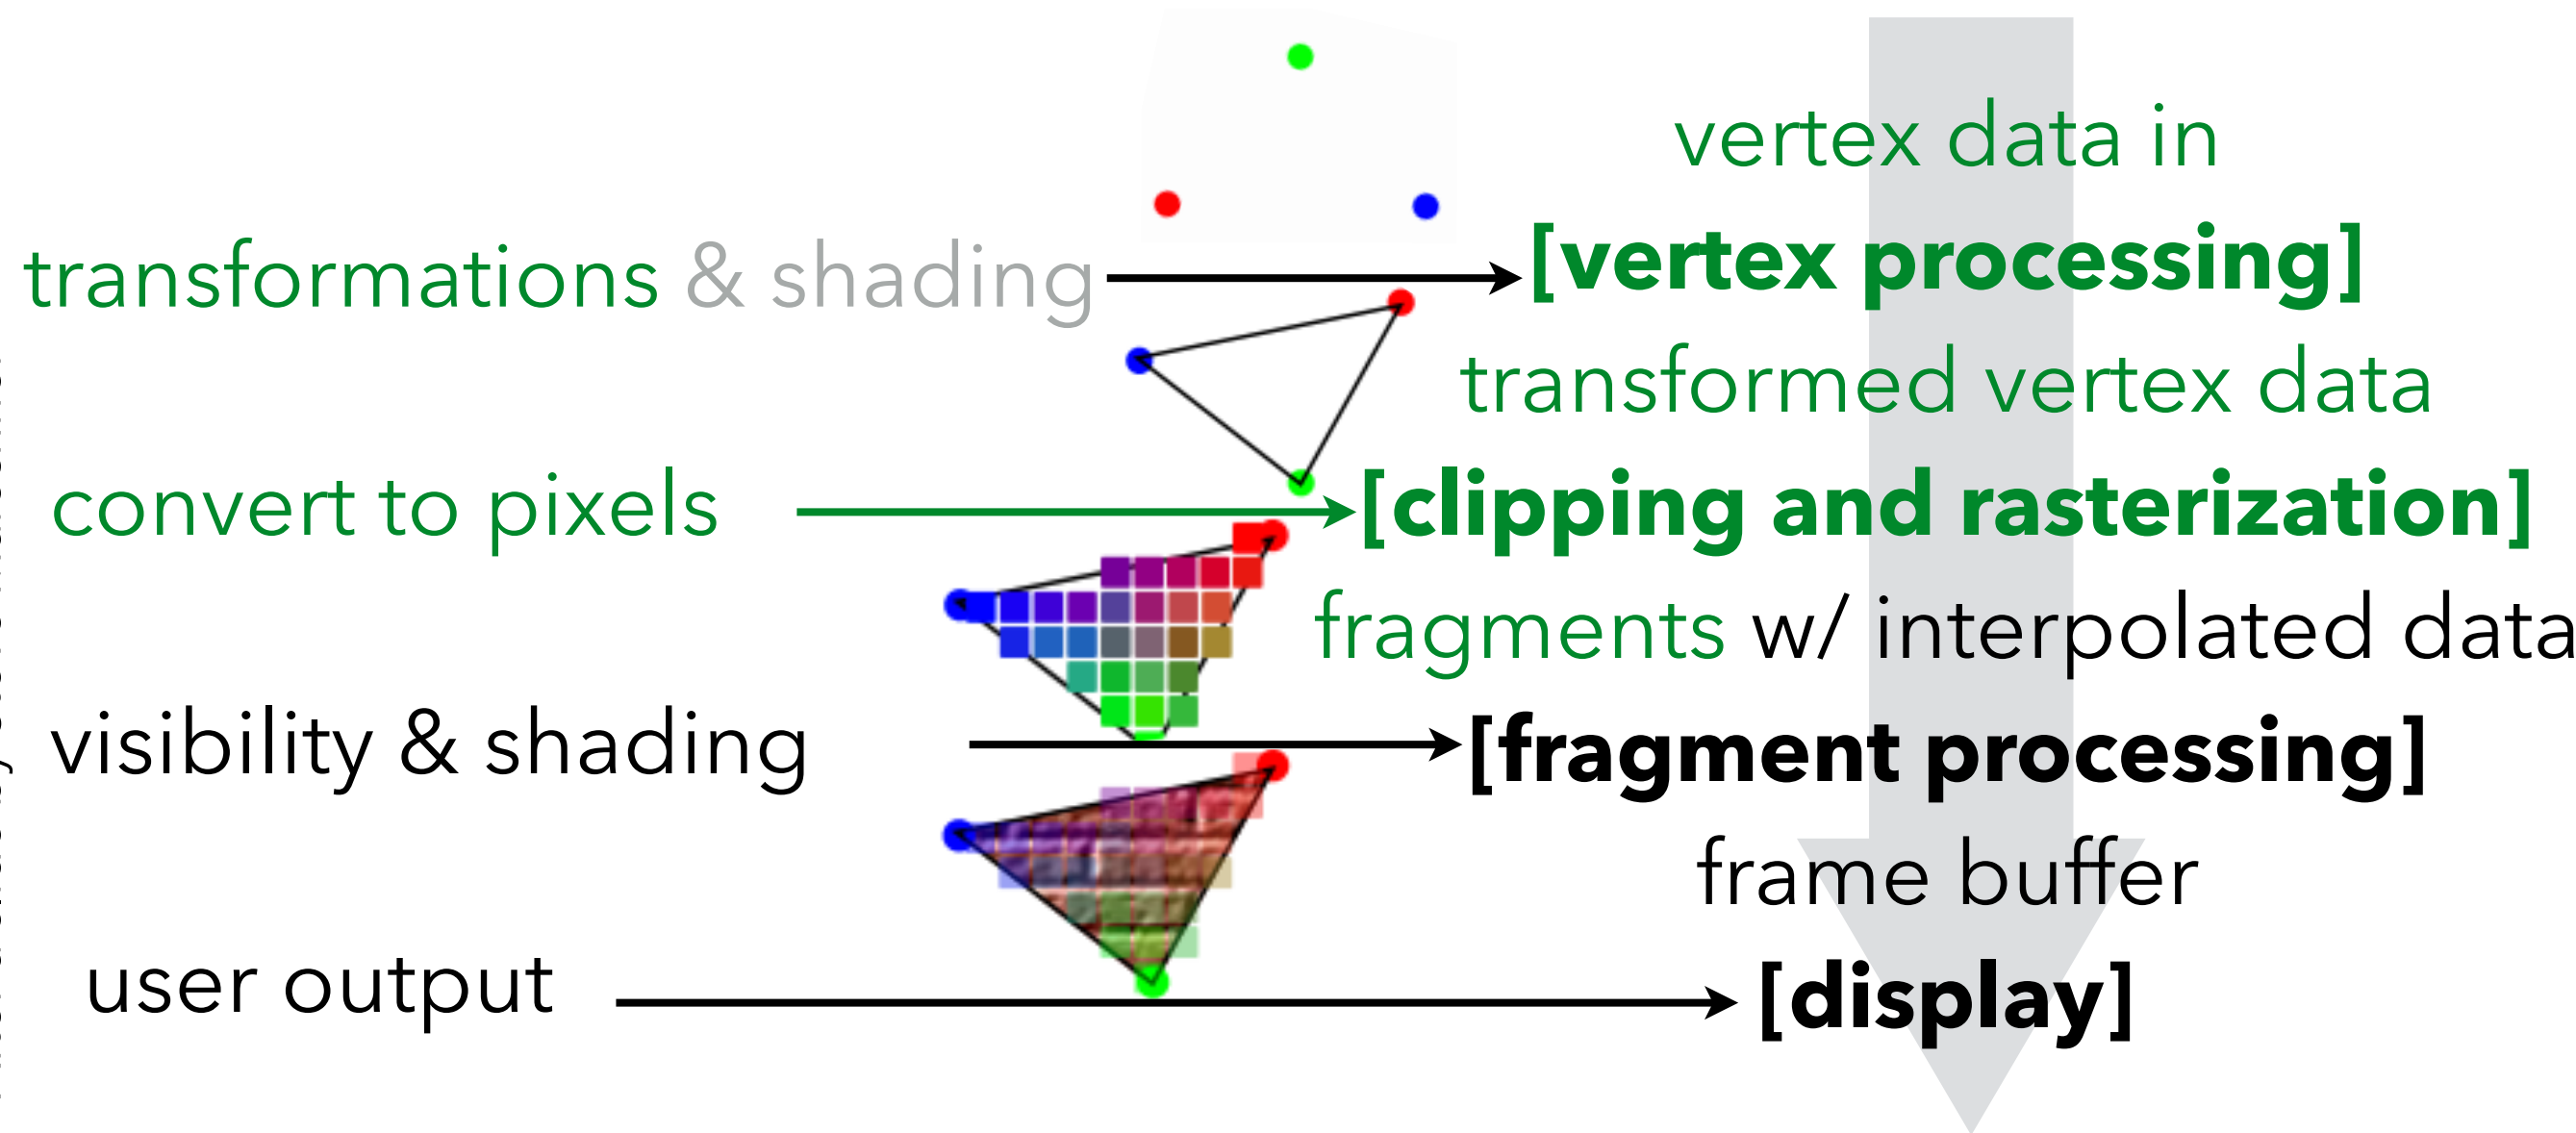
- pixels should only be drawn once, but no holes between adjacent triangles

# Programmable Rasterization Pipeline™



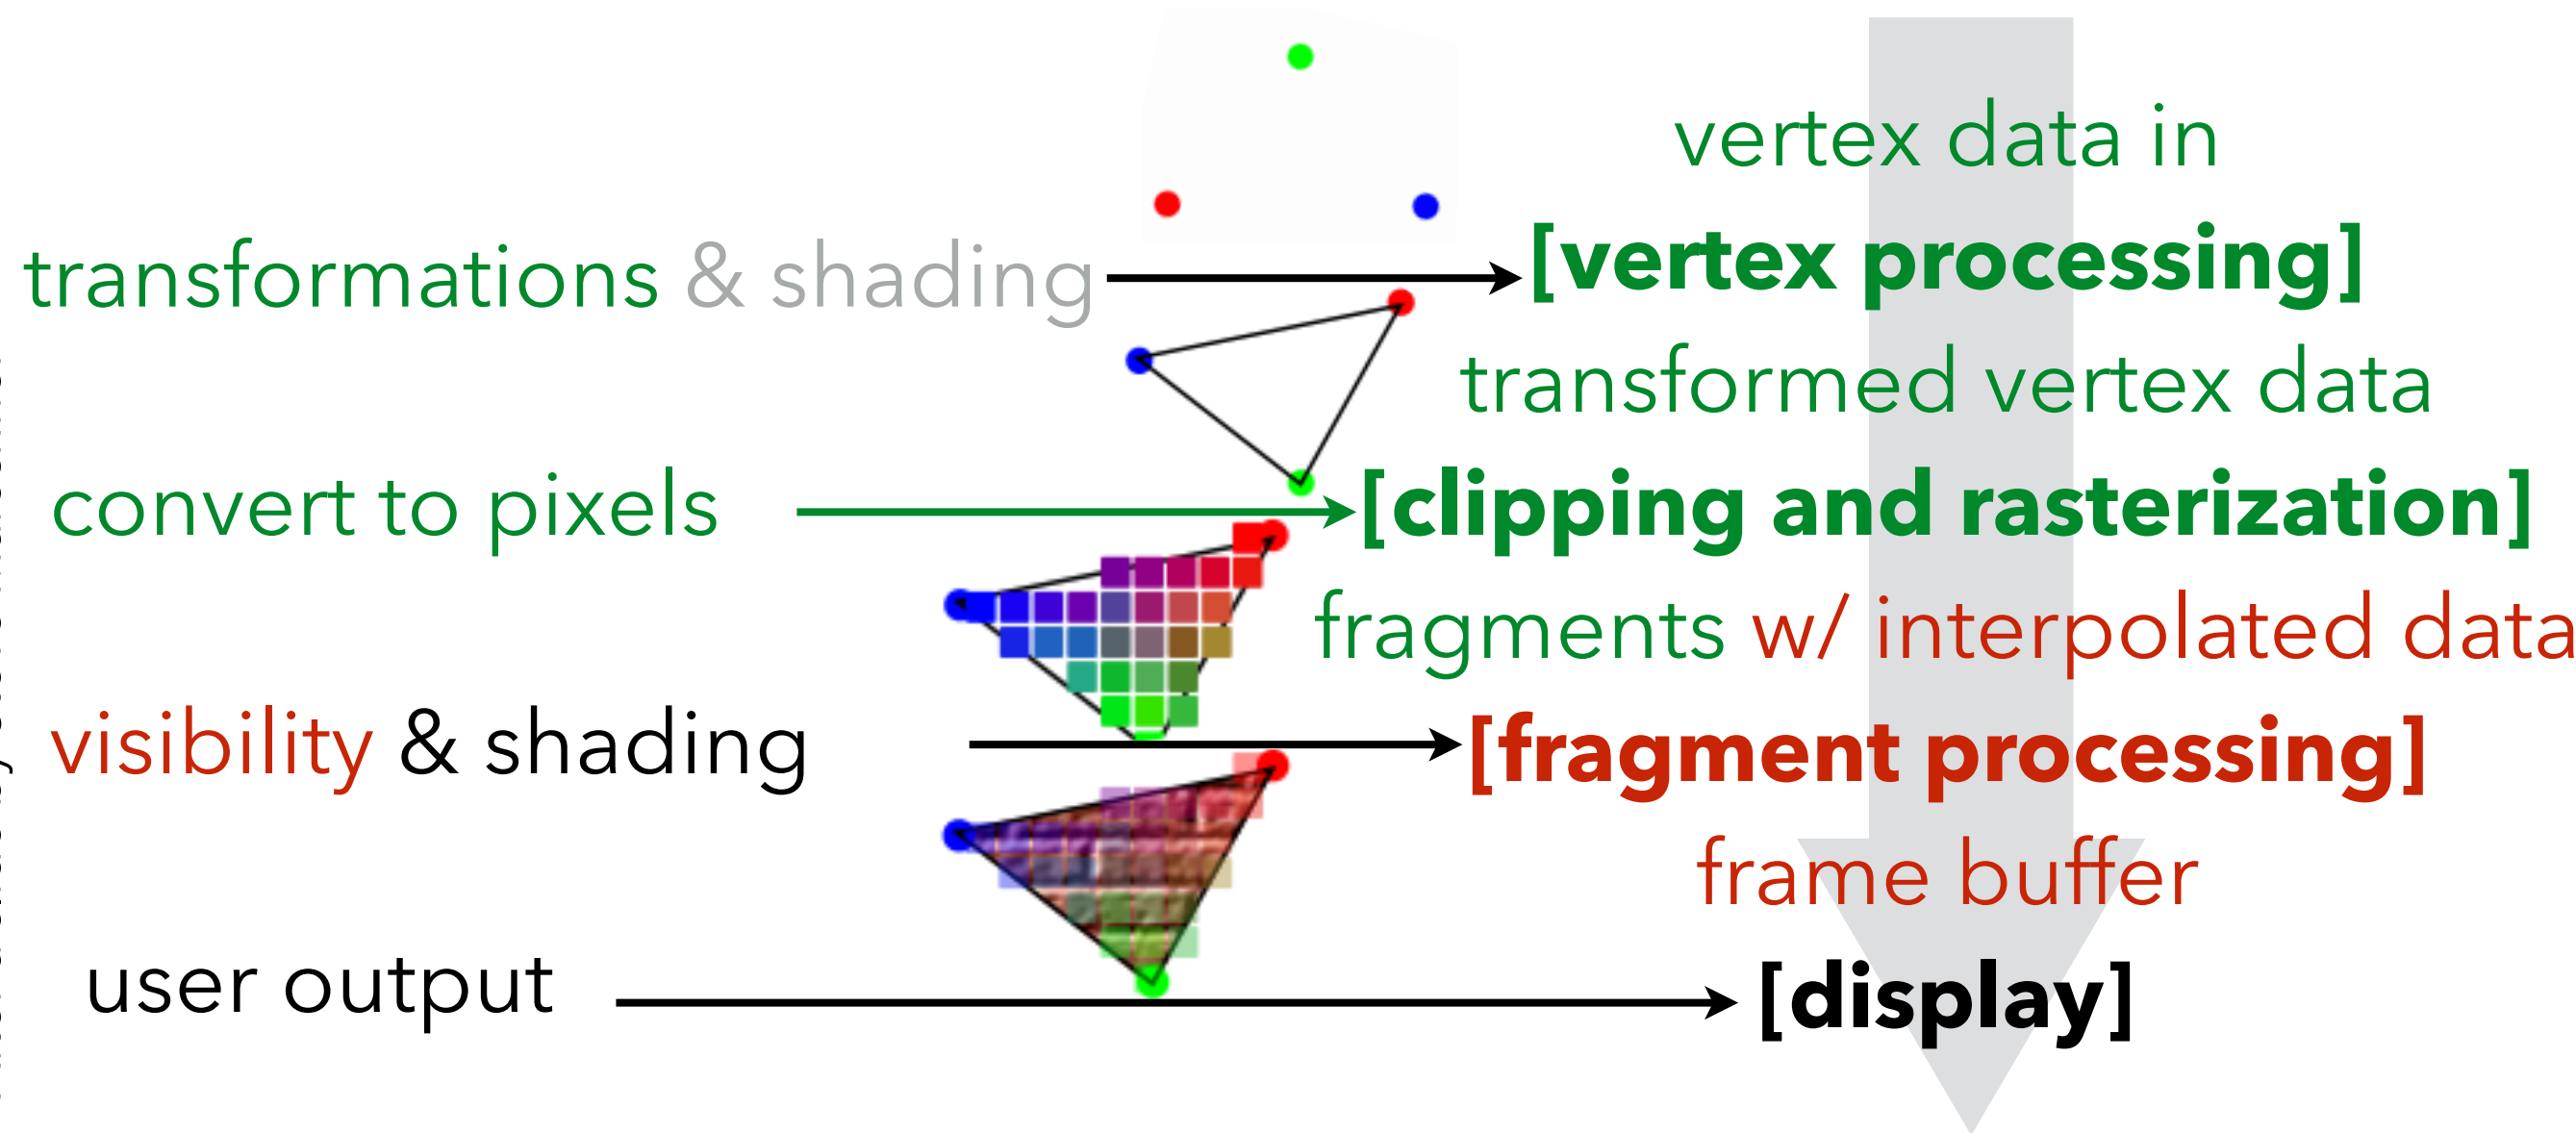


# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

# Programmable Rasterization Pipeline™



# Programmable Rasterization Pipeline™

```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```



# Hidden Surface Elimination




# Visible Surface Determination

# Hidden Surface Elimination

We have discussed how to map primitives to image space


- projection and perspective are depth cues
- occlusion is another very important cue



# Backface Culling

For closed shapes you will never see the inside of the shape


- therefore, only draw surfaces that face the camera
- implement by checking  $\mathbf{n} \cdot \mathbf{v}$



# Painter's Algorithm

Simplest way to treat hidden surfaces:




- draw from back to front, overwriting the framebuffer along the way





# Painter's Algorithm

Useful when a valid order is easy to come by



[Foley et al.]

# The z-buffer (a.k.a. depth buffer)

In many/most applications, maintaining a z sort is too expensive


- changes all the time as the view changes
- many data structures exist, but complex

Solution: draw in any order, keep track of closest

- allocate extra channel per pixel to keep track of **closest depth so far**
- when drawing, compare object's depth to current closest depth, and discard if greater\*
- works just like any other compositing operation

# The z-buffer

Another example of a memory-intensive, brute force approach that works and has become the standard




[Foley et al.]


# Precision in the z-buffer

The precision is distributed between the near and far clipping planes

- this is why these planes have to exist
- also why you can't always just set them to very small and very large distances
- generally use  $z'$  (not world  $z$ ) in z buffer



# Interpolating in Projection



linear interp. in screen space  $\neq$  linear interp. in world (eye) space

# Pipeline for Basic z-buffer

Vertex stage (input: position / vtx; color / tri)

- transform position (object to screen space)
- pass through color


Rasterizer

- interpolated parameter:  $z'$  (screen  $z$ )
- pass through color

Fragment stage (output: color,  $z'$ )

- write to frame buffer only if interpolated  $z' <$  current  $z'$

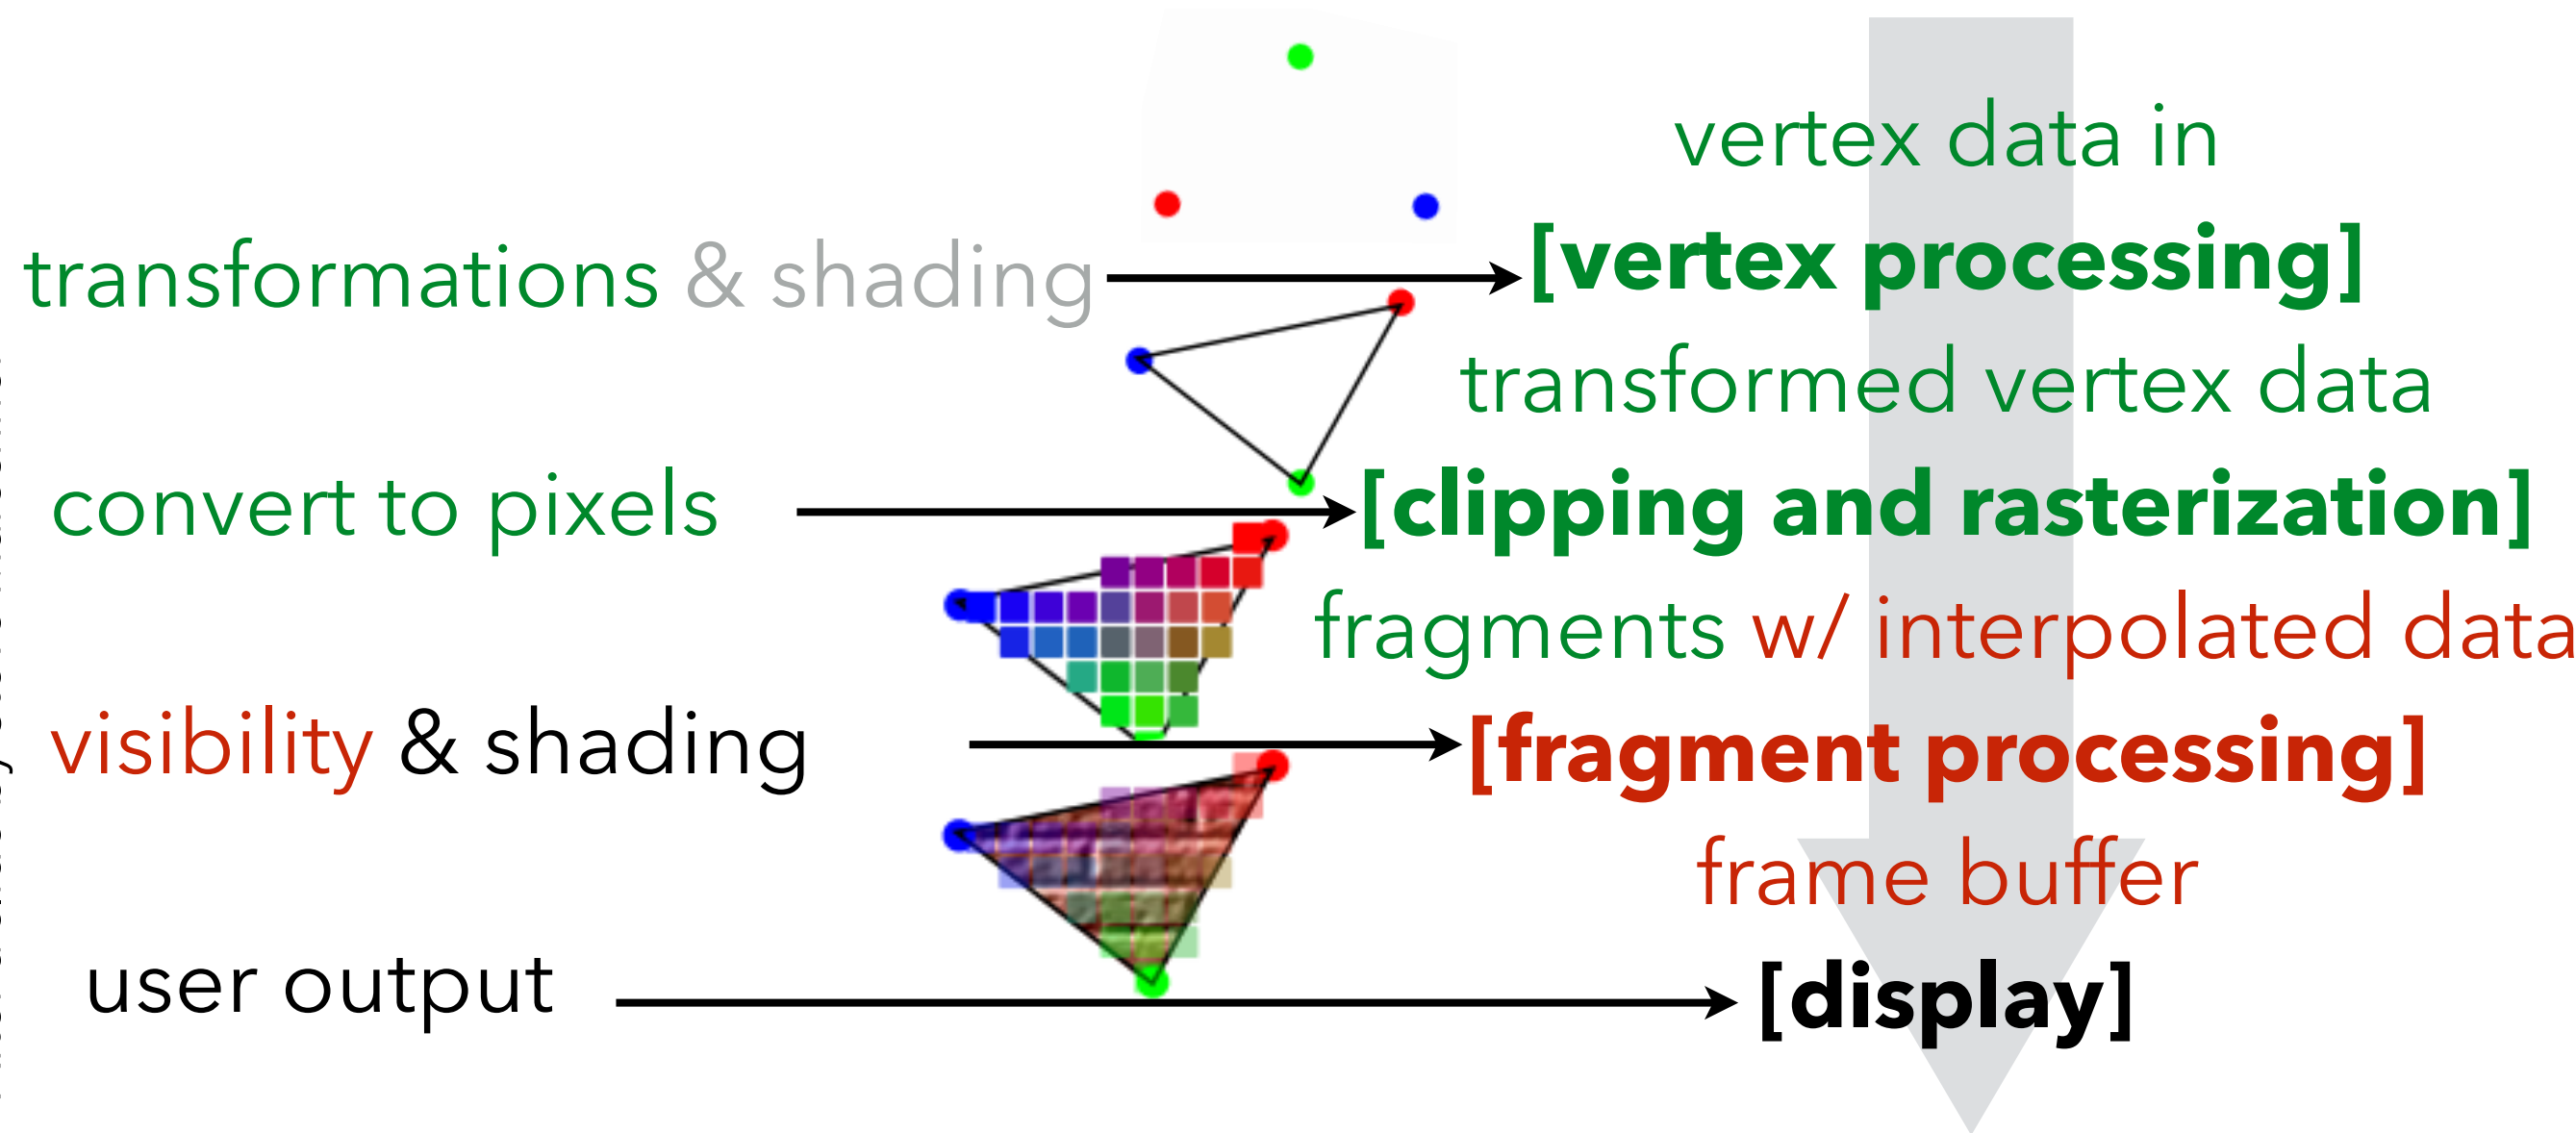
# Result of Basic z-buffer Pipeline



After a slide by Steve Marschner

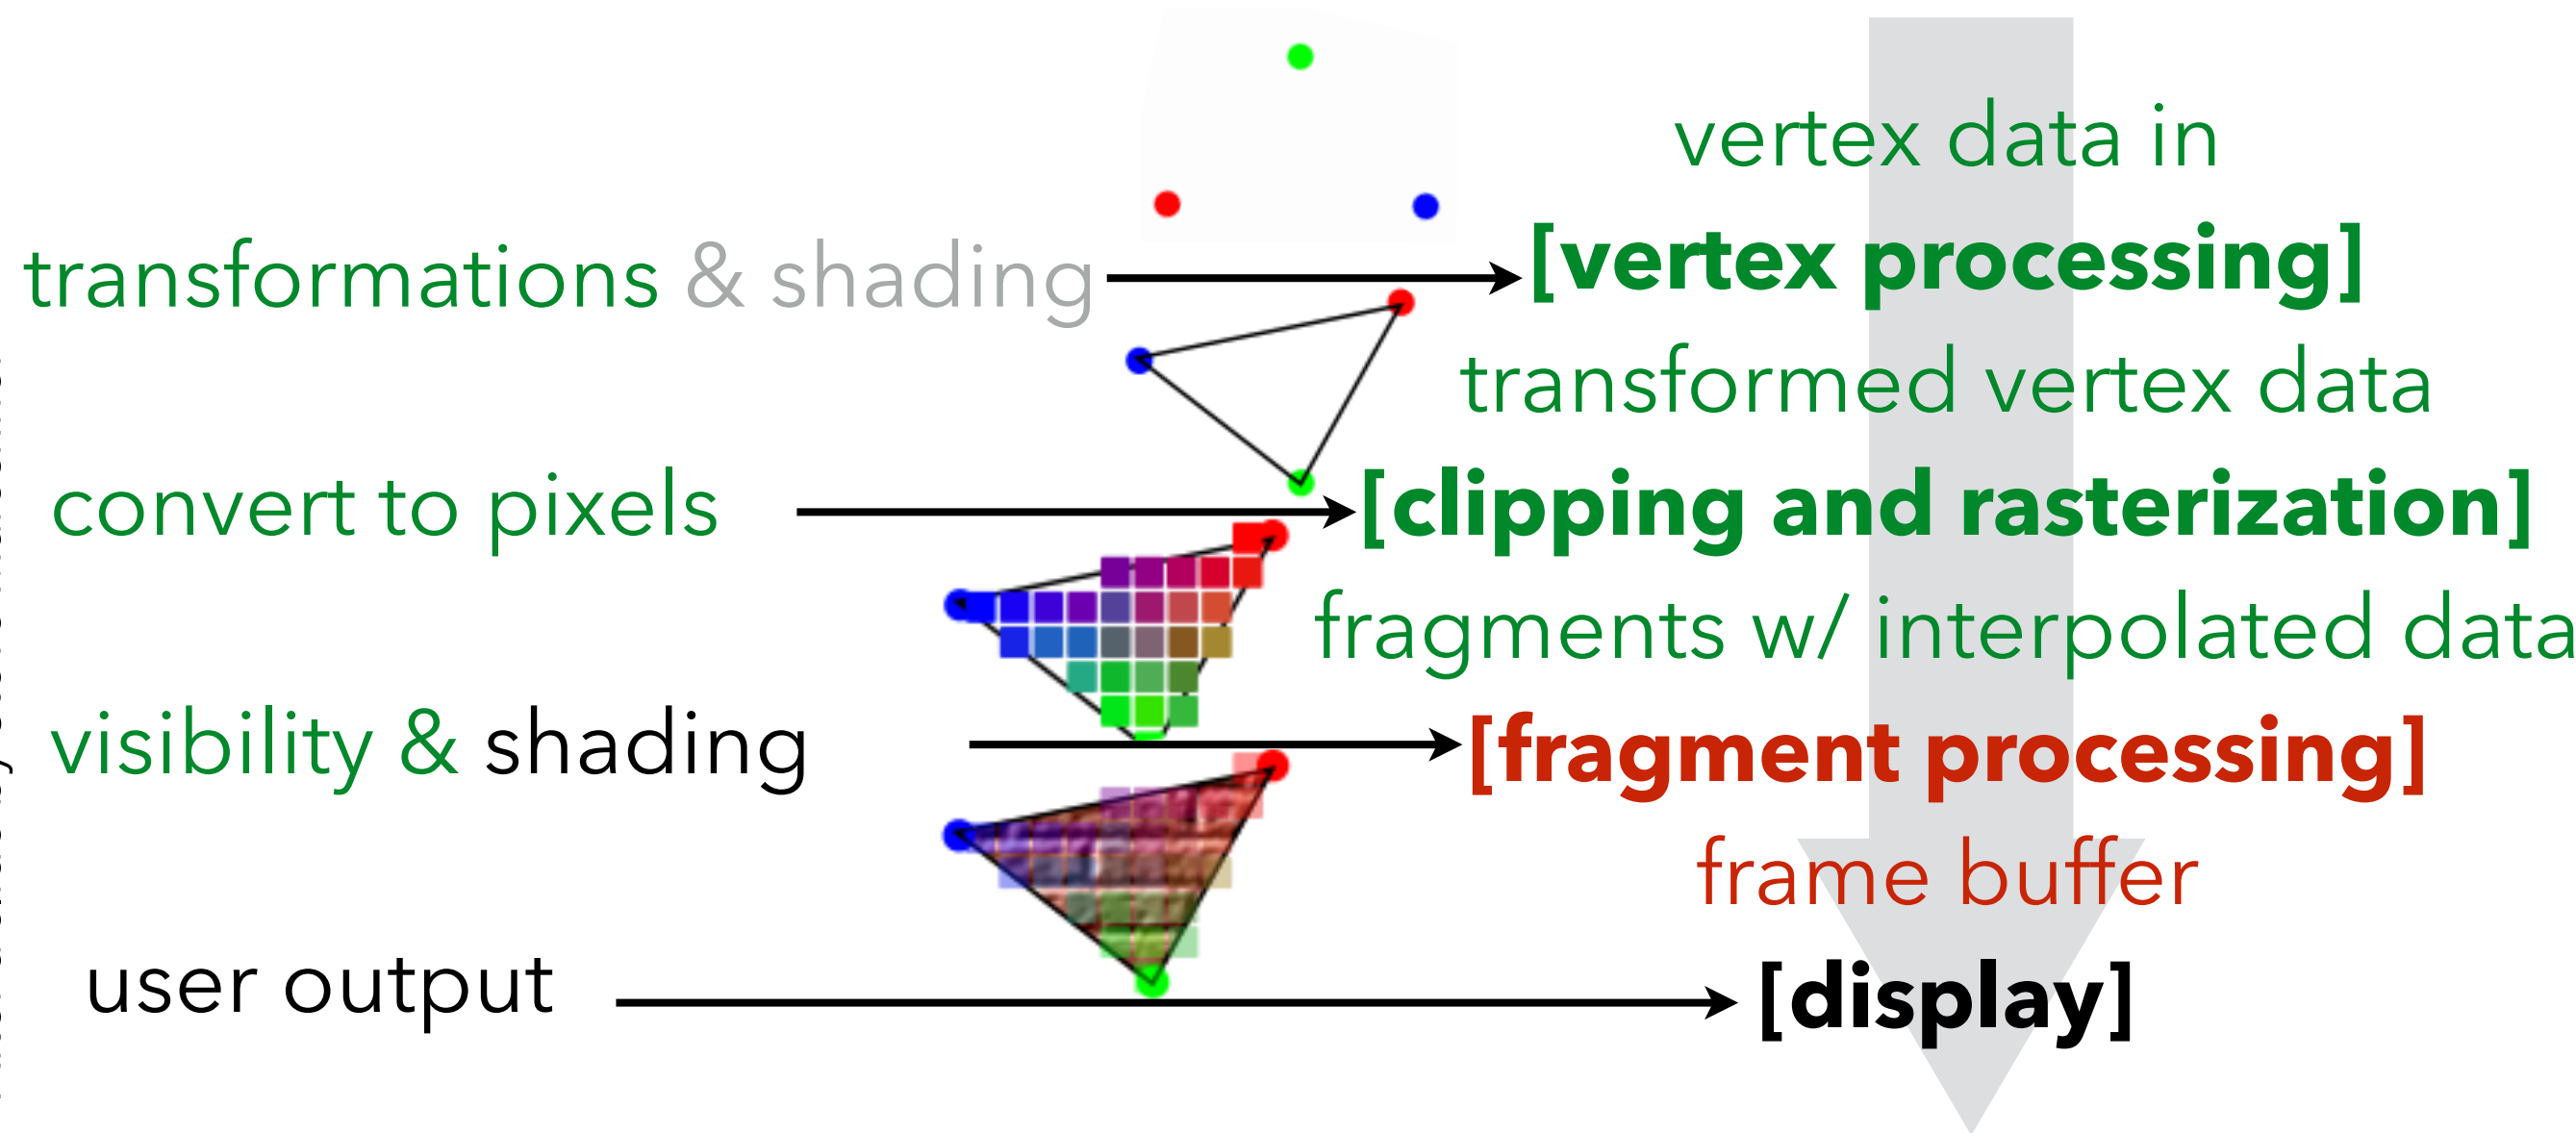


# Programmable Rasterization Pipeline™



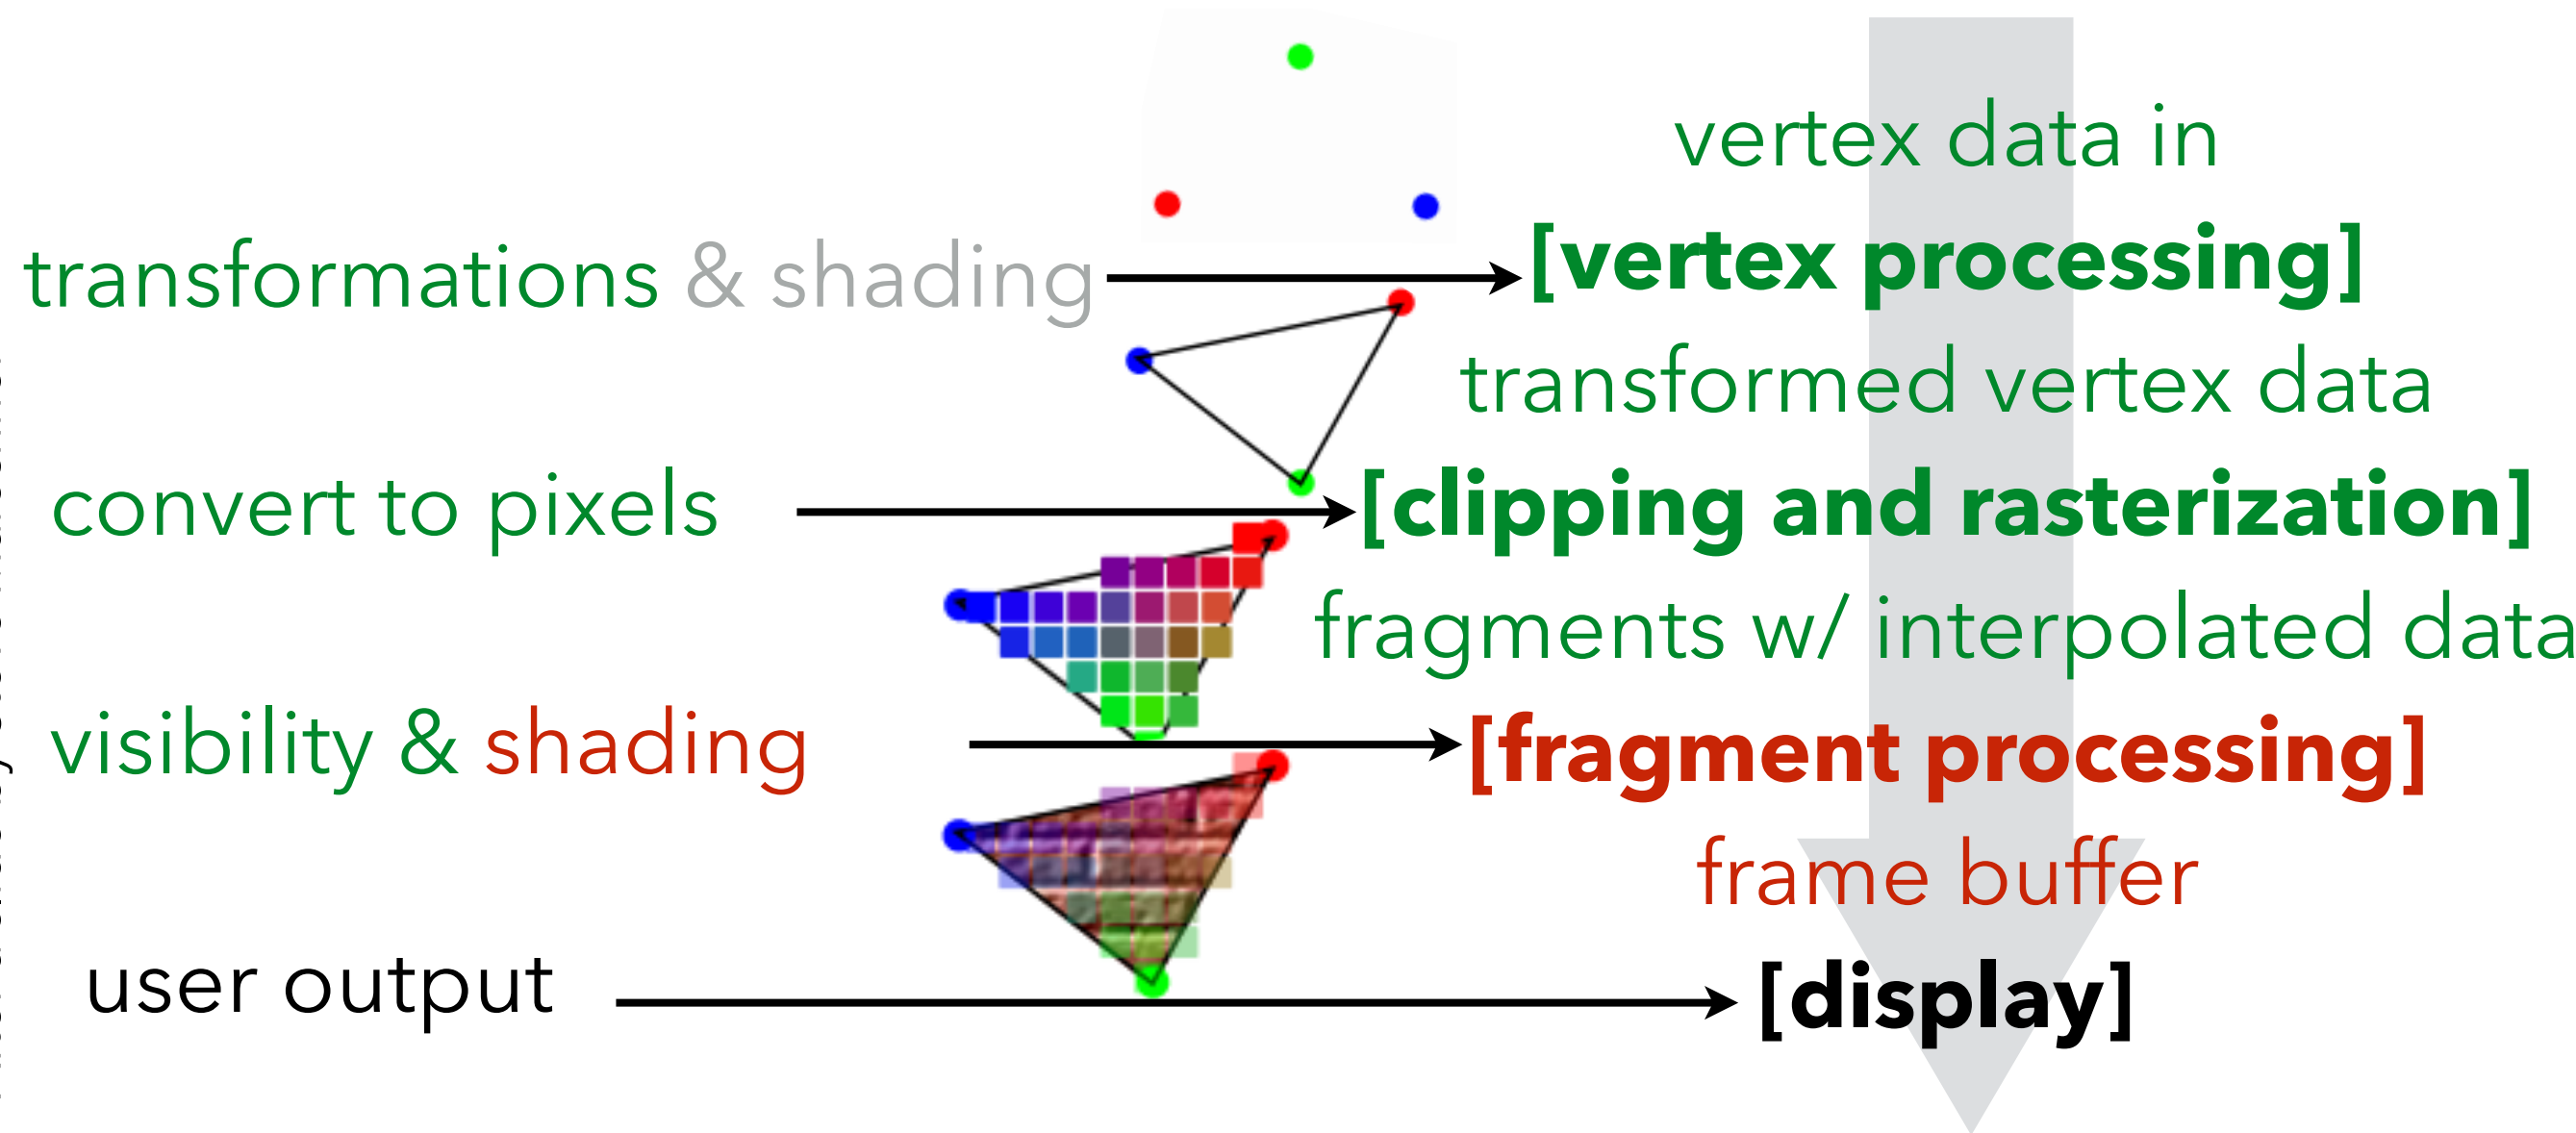


# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

# Programmable Rasterization Pipeline™



After a slide by Steve Marschner

# Programmable Rasterization Pipeline™

```
for(each triangle)
  transform vertices into eye space
  project vertices to image space
  for(each pixel x,y)
    if(x,y in triangle)
      compute z
      if(z < zbuffer[x,y])
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```